



SAA 03 Practice Test 3

Question 1 of 65

Exit Quiz

Domain: Design Cost-Optimized Architectures

A start-up firm has created a cloud storage application that gives users the ability to store any amount of personal data & share them with their connections. For this, they are using Amazon S3 buckets to store user data. The firm has used Amazon S3 multipart upload to upload large objects in parts. During the last quarter, the finance team has observed a surge in storage costs for the S3 bucket. On further checking, the firm observed that many 100 GB files are uploaded by users & are in a partially completed state. As an AWS consultant, the IT Team requests you prevent this from happening again. Which of the following actions can be taken to meet this requirement cost-effectively with the least effort?

- A. Create an S3 lifecycle Configuration to abort incomplete multipart uploads.
- B. Manually delete incomplete multipart uploads from the S3 bucket.
- C. Use Cron tool to identify incomplete uploads & delete those files.
- D. No action is required. All Incomplete uploads are automatically deleted every three months by Amazon S3.

Cevap : A

Lifecycle rule actions

Choose the actions you want this rule to perform. Per-request fees apply. [Learn more](#) or see [Amazon S3 pricing](#)

- Transition current versions of objects between storage classes
- Transition previous versions of objects between storage classes
- Expire current versions of objects
- Permanently delete previous versions of objects
- Delete expired delete markers or incomplete multipart uploads

When a lifecycle rule is scoped with tags, these actions are unavailable.



Question 2 of 65

[Exit Quiz](#)

Domain: Design High-Performing Architectures

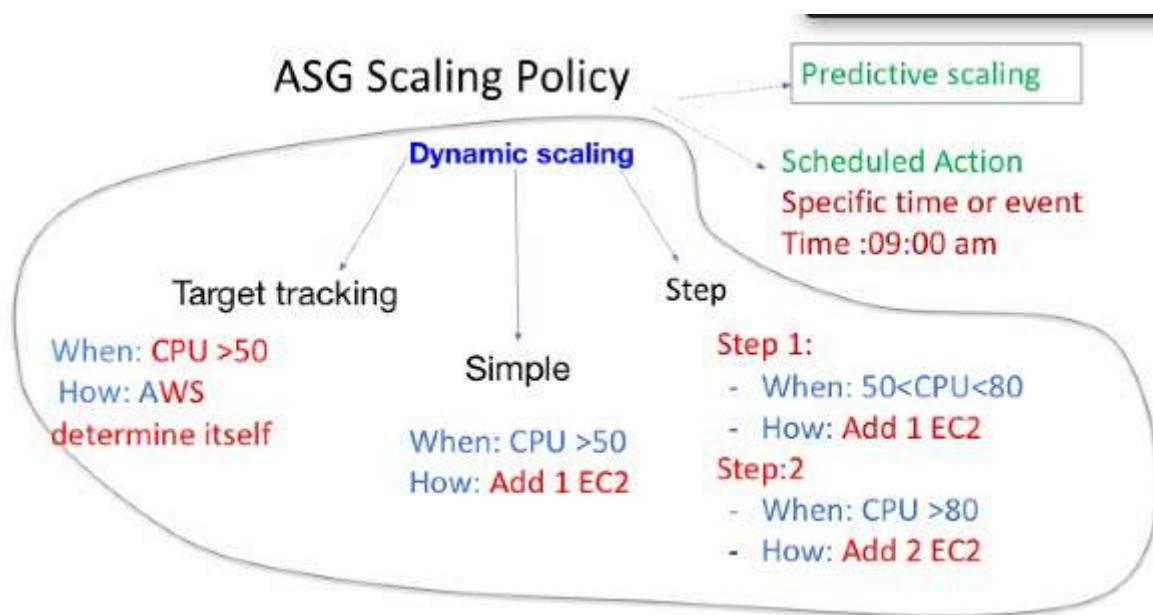
Your Operations department is using an incident-based application hosted on a set of EC2 Instances. These instances are placed behind an Auto Scaling Group to ensure that the right number of instances are in place to support the application. The Operations department has expressed dissatisfaction concerning poor application performance every day at 9:00 AM. However, it is also noted that the system performance returns to optimal at 9:45 AM.

What could be done to fix this issue?

- A. Create another Dynamic Scaling Policy to ensure that the scaling happens at 9:00 AM.
- B. Add another Auto Scaling group to support the current one.
- C. Change the Cool Down Timers for the existing Auto Scaling Group.
- D. Add a Scheduled Scaling Policy at 8:30 AM.

Cevap : D

Sınavda eger dinamik veya CPU geçiyorsa cevap Target tır



Question 3 of 65

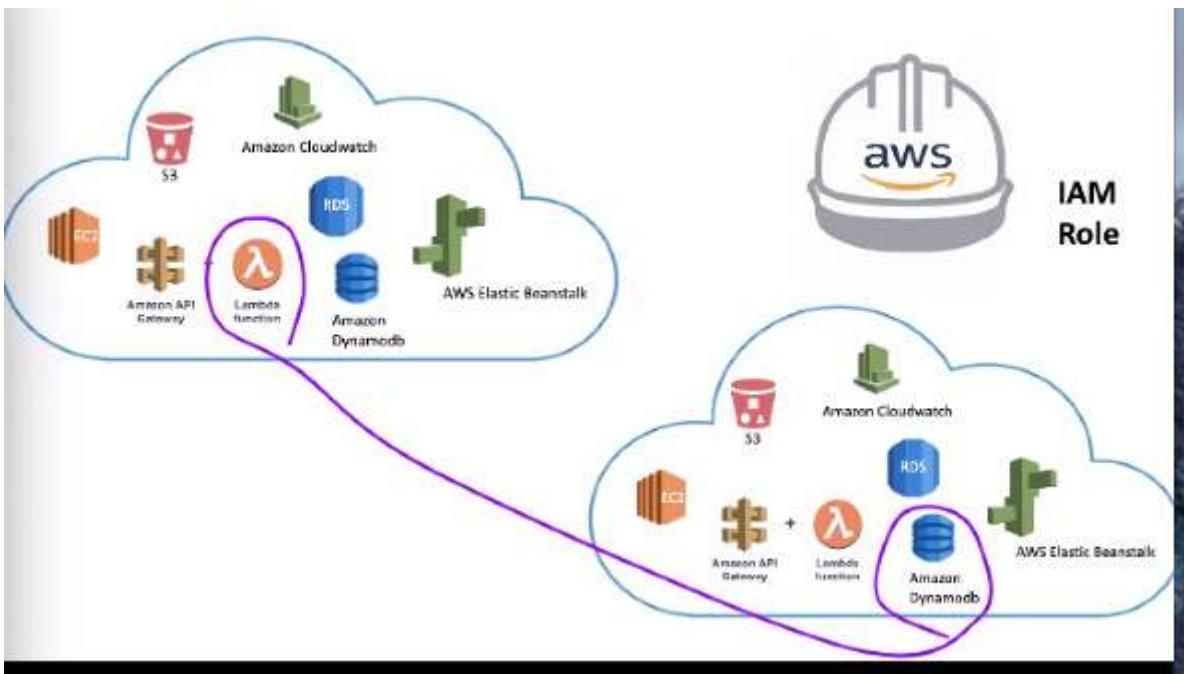
[Exit Quiz](#)

Domain: Design Secure Architectures

You have created an AWS Lambda function that will write data to a DynamoDB table. Which of the following must be in place to ensure that the Lambda function can interact with the DynamoDB table?

- A. Ensure an IAM Role is attached to the Lambda function which has the required DynamoDB privileges.
- B. Ensure an IAM User is attached to the Lambda function which has the required DynamoDB privileges.
- C. Ensure the Access keys are embedded in the AWS Lambda function.
- D. Ensure the IAM user password is embedded in the AWS Lambda function.

Cevap A



Domain: Design Cost-Optimized Architectures

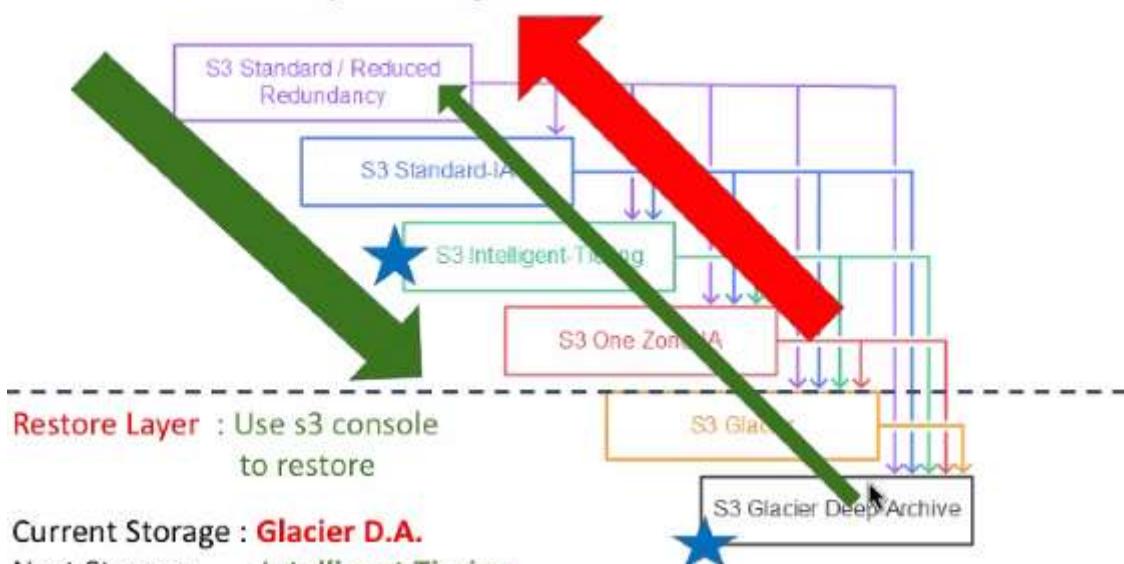
A Media firm is saving all its old videos in S3 Glacier Deep Archive. Due to the shortage of new video footage, the channel has decided to reuse all these old videos. Since these are old videos, the channel is not sure of their popularity & response from users. Channel Head wants to make sure that these huge size files do not shoot up their budget. For this, as an AWS consultant, you advise them to use the S3 intelligent storage class. The Operations Team is concerned about moving these files to the S3 Intelligent-Tiering storage class. Which of the following actions can be taken to move objects in Amazon S3 Glacier Deep Archive to the S3 Intelligent-Tiering storage class?

- A. Use Amazon S3 Console to copy these objects from S3 Glacier Deep Archive to the required S3 Intelligent-Tiering storage class.
- B. Use Amazon S3 Glacier Console to restore objects from S3 Glacier Deep Archive & then copy these objects to the required S3 Intelligent-Tiering storage class.
- C. Use Amazon S3 console to restore objects from S3 Glacier Deep Archive & then copy these objects to the required S3 Intelligent-Tiering storage class.
- D. Use the Amazon S3 Glacier console to copy these objects to the required S3 Intelligent-Tiering storage class.

Cevap C

Glacierden alırken kopyasını alıyoruz. Çekilirken aslında kopyası çekilir.

S3 Lifecycle Configuration-AWS console or CLI



Question 5 of 65

Exit Quiz

Domain: Design Secure Architectures

You have a cluster of Windows instances joined to an AWS Managed Active Directory. You want to have a shared storage for all these instances and control this storage access with the Managed Active Directory. Which of the following services allows you to achieve this?

- A. Amazon FSx for Lustre
- B. Amazon FSx for Windows File Server
- C. Amazon EFS
- D. Use S3 and AD Connector

Cevap B

Amazon FSx file systems

FSx _N	FSx _Z	FSx _W	FSx _L
Amazon FSx for NetApp ONTAP Fully managed shared storage built on NetApp's popular ONTAP file system.	Amazon FSx for OpenZFS Fully managed shared storage built on the popular OpenZFS file system.	Amazon FSx for Windows File Server Fully managed shared storage built on Windows Server.	Amazon FSx for Lustre Fully managed shared storage built on the world's most popular high-performance file system.



- For Linux Instance
- Can't write/read S3

- Compatible with Windows Active Directory

- HPC
- Compatible with S3

- Compatible with FSx File Gateway

Question 6 of 65

Exit Quiz

Domain: Design Cost-Optimized Architectures

You are part of the IT team of an insurance company. You have 4 M5.large EC2 instances used to compute some data of your core services. The amount of usage of these instances has been very consistent. So you predict that it will not increase in the next two or three years. However, your CFO is asking if there is a way to reduce costs in the EC2 instances. What do you suggest to get the maximum cost reduction?

- A. Use a Compute Savings Plan.
- B. Use an EC2 instance Savings Plan.
- C. Use a Convertible Reserved Instance.
- D. Use a Dedicated Instance.

Cevap B

Convertible da cpu artırımı olabiliyor.

Cost Savings Plan		Reserved Instances	
1 year/1000 dollars/limit		1 year/1500 dollars/limitless	
Compute Saving Plans	EC2 Fargate Lambda	Convertible RI (Reserved Instance)	EC2
%66 Cost Saving		%66 Cost Saving	
EC2 Instance Saving Plans	EC2	Standard RI (Reserved Instance)	EC2
%72 Cost Saving		%72 Cost Saving	
convertible to the other size of instance			

Question 7 of 65

[Exit Quiz](#)

Domain: Design Cost-Optimized Architectures

You are building an automated transcription service where Amazon EC2 worker instances process an uploaded audio file and generate a text file. You must store both of these files in the same durable storage until the text file is retrieved. Customers fetch the text files frequently. You do not know about the storage capacity requirements. Which storage option would be both cost-efficient and highly available in this situation?

- A. Multiple Amazon EBS Volume with snapshots
- B. A single Amazon Glacier Vault
- C. A single Amazon S3 bucket
- D. Multiple instance stores

Cevap : C

S3 de limit yok. Kapasiteyi bilmiyoruz çünkü

You are building an automated transcription service where Amazon EC2 worker instances process an uploaded audio file and generate a text file. You must store both of these files in the same durable storage until the text file is retrieved.

Customers fetch the text files frequently. You do not know about the storage capacity requirements. Which storage option would be both cost-efficient and highly available in this situation?



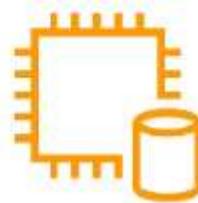
Amazon EBS



Amazon Glacier



amazon
S3



Instance Store

Question 8 of 65

[Exit Quiz](#)

Domain: Design High-Performing Architectures

A customer has an instance hosted in the public subnet of the default VPC. The subnet has the default settings for the Network Access Control List. An IT Administrator needs to be provided SSH access to the underlying instance. How could this be accomplished?

- A. Ensure the Network Access Control Lists allow Inbound SSH traffic from the IT Administrator's Workstation.
- B. Ensure the Network Access Control Lists allow Outbound SSH traffic from the IT Administrator's Workstation.
- C. Ensure that the Security group allows Inbound SSH traffic from the IT Administrator's Workstation.
- D. Ensure that the Security group allows Outbound SSH traffic from the IT Administrator's Workstation.

Cevap : C

Network ACL: acl-021ec68e42e806c8

Default NACL

Details Inbound Rules Outbound Rules Subnet associations Tags

Edit inbound rules

View	All rules				
Rule #	Type	Protocol	Port Range	Source	Allow / Deny
100	All Traffic	All	All	0.0.0.0/0	ALLOW
*	All Traffic	All	All	0.0.0.0/0	DENY

Edit outbound rules

View	All rules				
Rule #	Type	Protocol	Port Range	Destination	Allow / Deny
100	All Traffic	All	All	0.0.0.0/0	ALLOW
*	All Traffic	All	All	0.0.0.0/0	DENY

Question 9 of 65

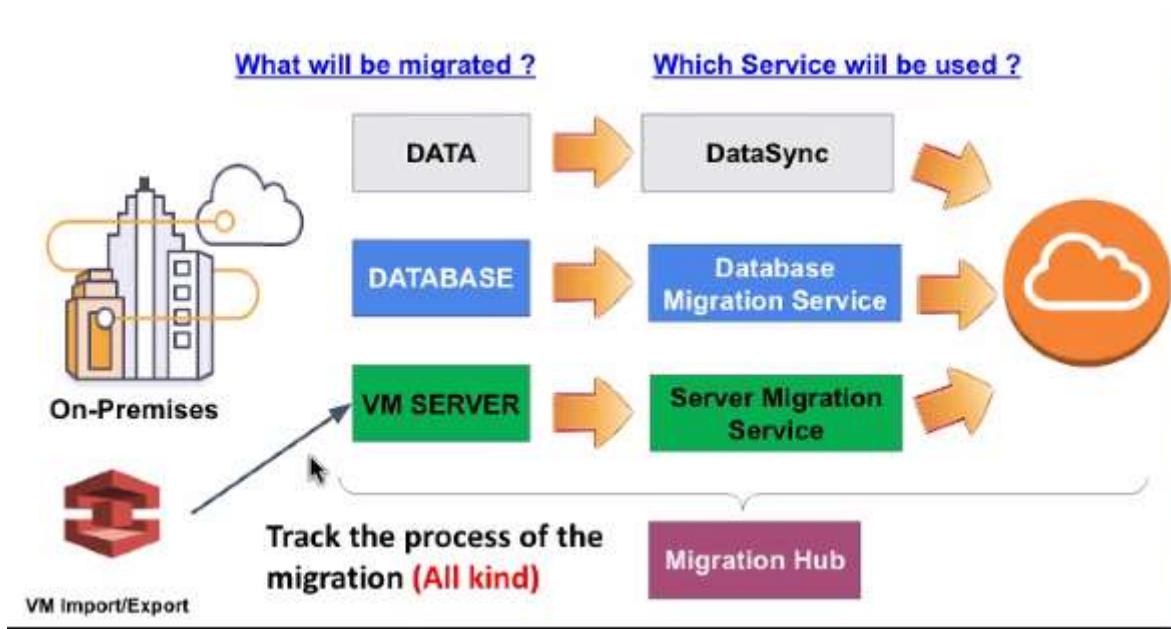
Exit Quiz

Domain: Design Secure Architectures

Your Organization is planning to move its on-premise databases to the AWS Cloud. You have been selected to migrate the main production database, and there are some requirements. The production database should remain active during the migration. You need to monitor the progress of the migration. The database is an SQL Server Database. You need to find an easy way to convert the actual schemas to MySQL schemas. What services could help to achieve this? (Select two)

- A. AWS DataSync.
- B. AWS Server Migration Service.
- C. AWS Database Migration Service.
- D. AWS Migration Hub.
- E. AWS Server Migration Service Connector.

Cevap C ve D tüm taşımaları izlemek için kullanılır.



Question 10 of 65

[Exit Quiz](#)

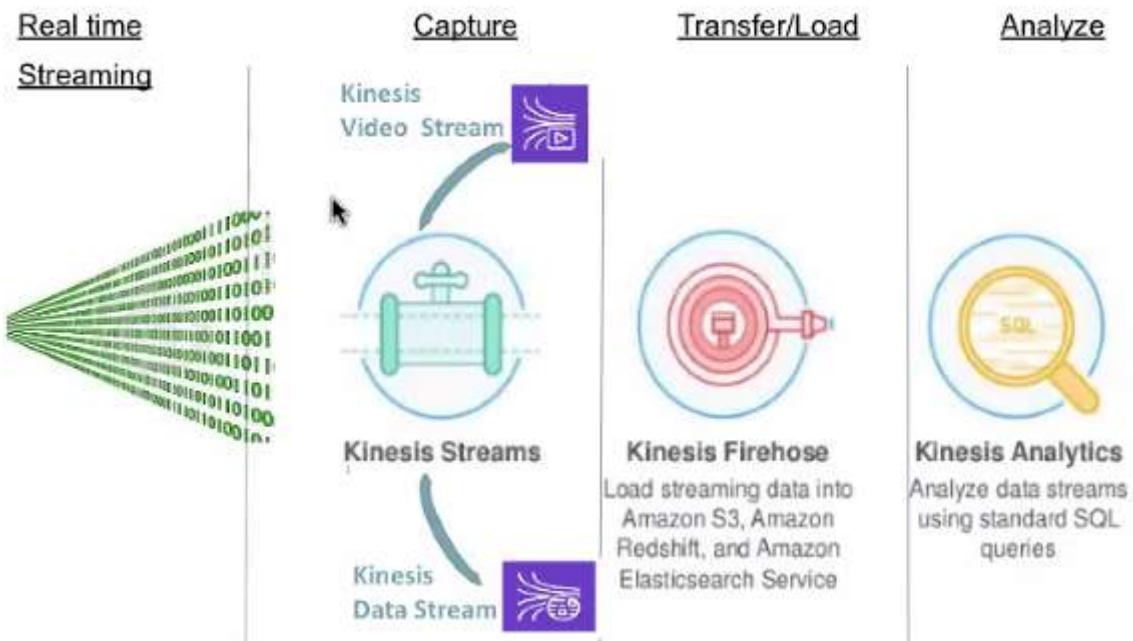
Domain: Design High-Performing Architectures

You are deploying an application to track the GPS coordinates of delivery trucks in the United States. Coordinates are transmitted from each delivery truck once every three seconds. You need to design an architecture that will enable real-time processing of these coordinates from multiple consumers. Which of the following services would you use to implement data ingestion?

- A. Amazon Kinesis
- B. AWS Data Pipeline
- C. Amazon Elastic Transcoder
- D. Amazon Simple Queue Service

Cevap A

GERcek zamanlı,anlık



Question 11 of 65

[Exit Quiz](#)

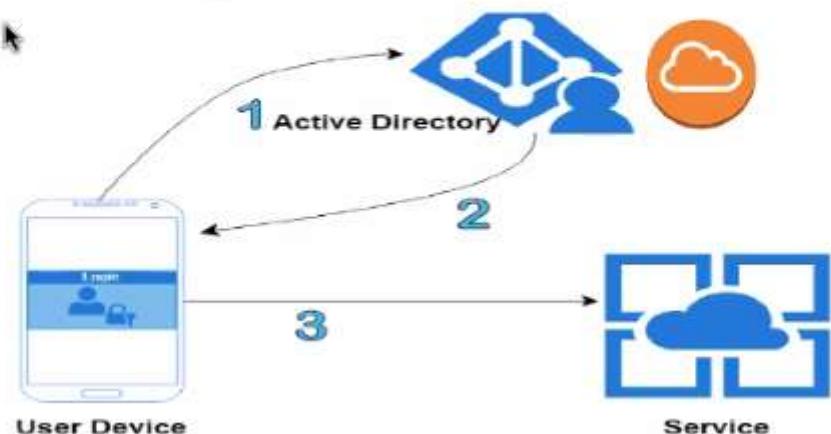
Domain: Design High-Performing Architectures

With its own Active Directory, your company authenticates users for different applications. You have been assigned the task of consolidating and migrating services to the cloud and using the same credentials if possible. What would you recommend?

- A. Use AWS Directory Service that allows users to sign in with their existing corporate credentials.
- B. Create two Active Directories – one for the cloud and one for on-premises – reducing username/password combinations to two.
- C. Require users to use third-party identity providers to log-in for all services.
- D. Build out Active Directory on EC2 instances to gain more control over user profiles.

Cevap : A

Your company authenticates users in a very disconnected network requiring each user to have several username/password combinations for different applications. You have been assigned a task of consolidating and migrating services to the cloud and reducing the number of usernames and passwords, employees need to use. What would you recommend?



Question 12 of 65

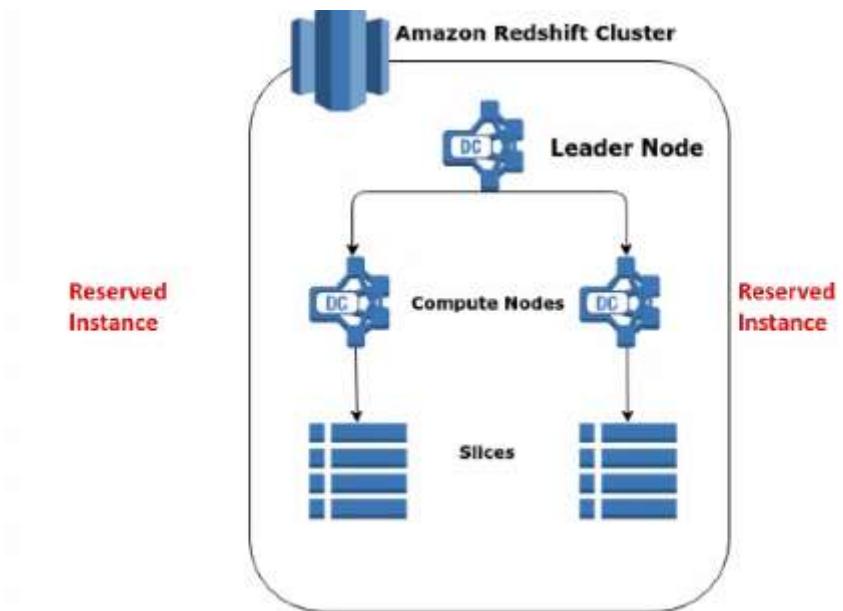
[Exit Quiz](#)

Domain: Design Cost-Optimized Architectures

A company is planning to use the AWS Redshift service. The Redshift service and data on it would be used continuously for the next 3 years as per the current business plan. What would be the most cost-effective solution in this scenario?

- A. Consider using On-demand instances for the Redshift Cluster.
- B. Enable Automated backup.
- C. Consider using Reserved Instances for the Redshift Cluster.
- D. Consider not using a cluster for the Redshift nodes.

Cevap :C



Question 13 of 65

[Exit Quiz](#)

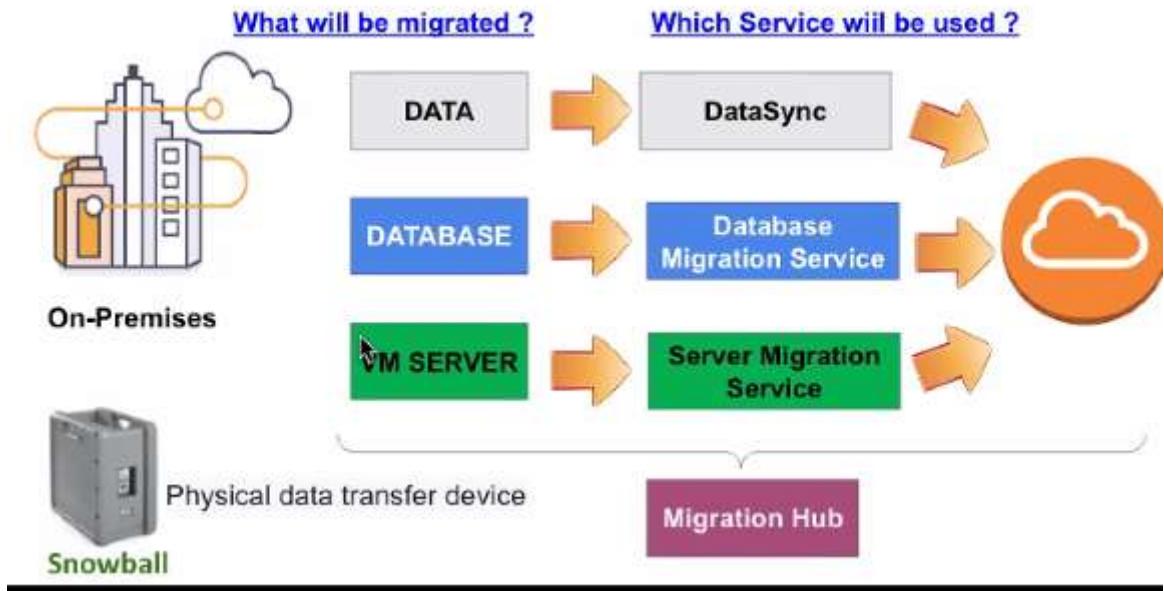
Domain: Design High-Performing Architectures

A company has recently started using AWS Cloud services and needs to transfer a large set of data online from on-prem Windows servers to AWS Storage Services S3, EFS, and FSx. The data can be transferred in opposite directions periodically and should be incremental based on schedules.

How would a Solution Architect design this solution?

- A. Use Snowball devices to transfer data to S3, EFS and FSx.
- B. Use AWS DataSync service to transfer data to AWS Services.
- C. AWS Database Migration Service to transfer data to AWS services.
- D. Use AWS S3 Transfer acceleration to transfer a large set of data.

Cevap B



Question 14 of 65

[Exit Quiz](#)

Domain: Design Secure Architectures

You have a PHP application deployed in an Auto Scaling group. In production, you want to use AWS WAF to block requests associated with exploiting vulnerabilities specific to the PHP use, including injection of unsafe PHP functions. Which method is appropriate?

- A. Add the AWS managed PHP application rule to AWS Shield.
- B. Add the AWS managed PHP application rule in the web ACL of AWS WAF.
- C. Add a PHP protection rule from AWS Marketplace to the WAF web ACL.
- D. Override the PHP rule's actions under the 'ExcludedRules' specification inside a rule group of a web ACL.

Cevap : B

Linux operating system	Contains rules that block request patterns associated with exploitation of vulnerabilities specific to Linux, including LFI attacks. This can help prevent attacks that expose file contents or execute code for which the attacker should not have been granted.	200	<input type="radio"/> Add to web ACL
PHP application	Contains rules that block request patterns associated with exploiting vulnerabilities specific to the use of PHP, including injection of unsafe PHP functions. This can help prevent exploits that allow an attacker to remotely execute code or run commands.	100	<input type="radio"/> Add to web ACL
POSIX operating system	Contains rules that block request patterns associated with exploiting vulnerabilities specific to POSIX/POSIX-like OS, including LFI attacks. This can help prevent attacks that expose file contents or execute code for which access should not have been allowed.	100	<input type="radio"/> Add to web ACL
SQL database	Contains rules that allow you to block request patterns associated with exploitation of SQL databases, like SQL injection attacks. This can help prevent remote injection of unauthorized queries.	200	<input type="radio"/> Add to web ACL
Windows operating system	Contains rules that block request patterns associated with exploiting vulnerabilities specific to Windows, (e.g., PowerShell commands). This can help prevent exploits that allow an attacker to run unauthorized commands or execute malicious code.	200	<input type="radio"/> Add to web ACL
WordPress application	The WordPress Applications group contains rules that block request patterns associated with common attacks, such as cross-site scripting (XSS) and denial-of-service (DoS).	100	<input type="radio"/> Add to web ACL

Question 15 of 65

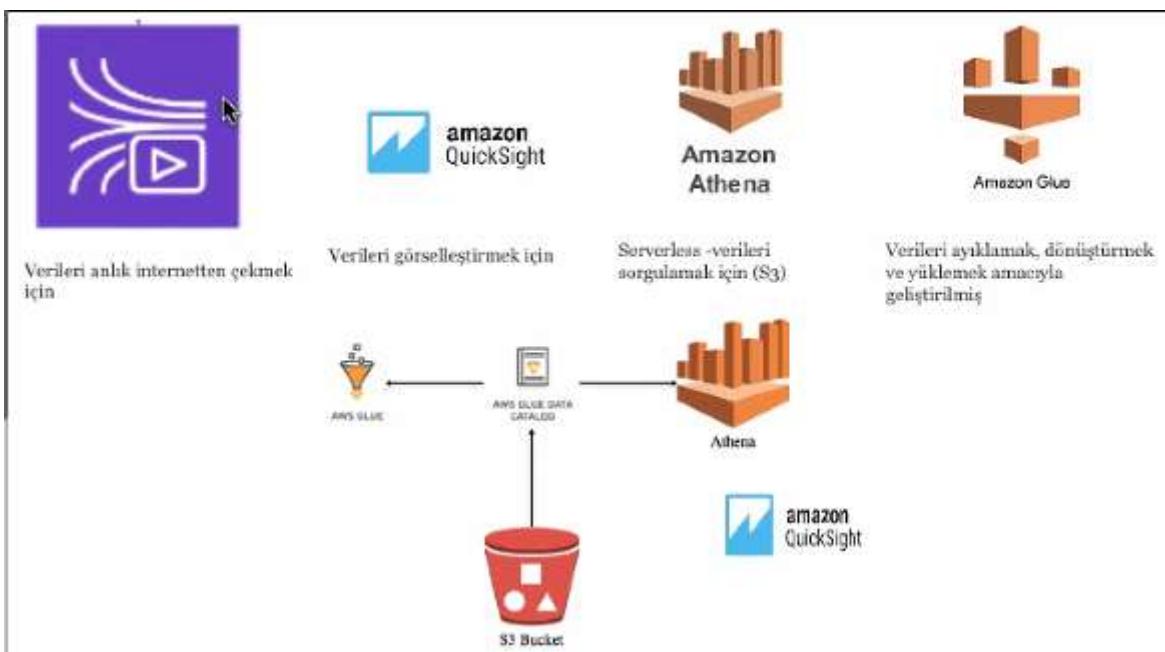
Exit Quiz

Domain: Design Cost-Optimized Architectures

A large amount of structured data is stored in Amazon S3 using the JSON format. You need to use a service to analyze the S3 data directly with standard SQL. In the meantime, the data should be easily visualized through data dashboards. Which of the following services is the most appropriate?

- A. Amazon Athena and Amazon QuickSight.
- B. AWS Glue and Amazon Athena.
- C. AWS Glue and Amazon QuickSight.
- D. Amazon Kinesis Data Stream and Amazon QuickSight.

Cevap A



Question 16 of 65

Exit

Domain: Design Secure Architectures

An AWS Organization has the below hierarchy of **Organizational Units (OUS)**:

Root → Project_OU → Dev_OU

The Root is attached to the default **Service Control Policy (SCP)**.

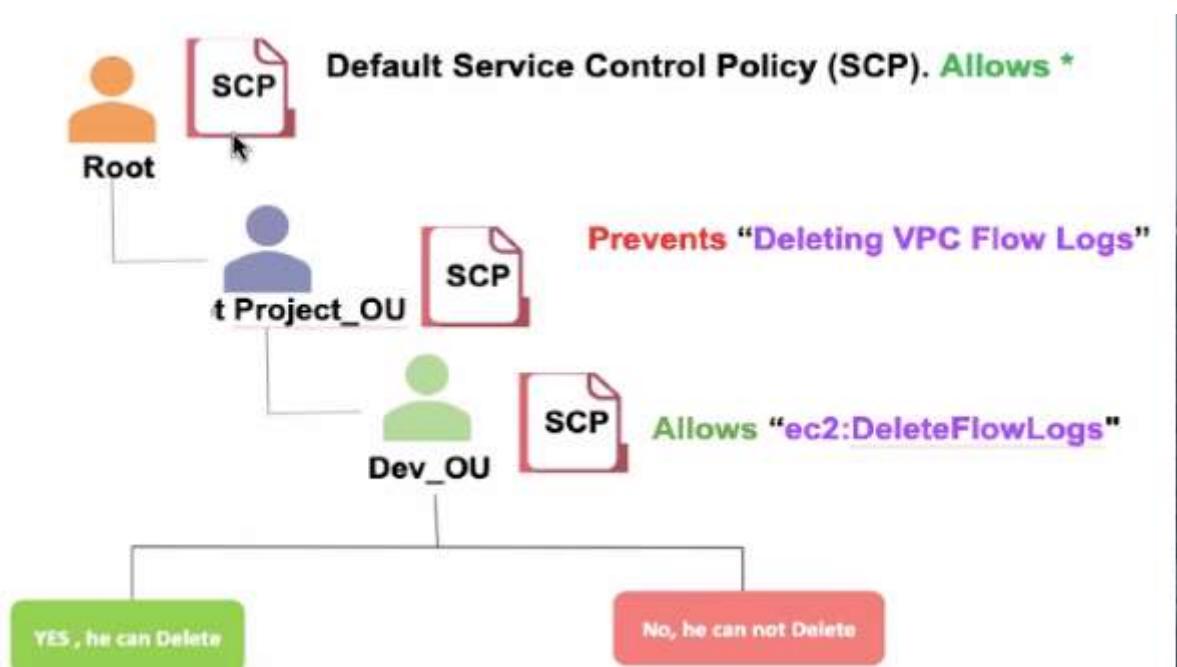
Project_OU is attached to an SCP that prevents users from deleting VPC Flow Logs.

Dev_OU has an SCP that allows the action of "ec2:DeleteFlowLogs".

Are the IAM users/roles in Dev_OU AWS accounts allowed to delete VPC Flow Logs?

- A. It is permitted because the SCP in Dev_OU allows it.
- B. It is allowed because the Root has the default SCP that allows all actions.
- C. It is not allowed as the SCP in Project_OU restricts the action.
- D. It is not allowed as the default SCP in Root denies the action.

Cevap C bu hiyerarşide göre silemez çünkü ortada yasak var.



Question 17 of 65

[Exit Quiz](#)

Domain: Design Secure Architectures

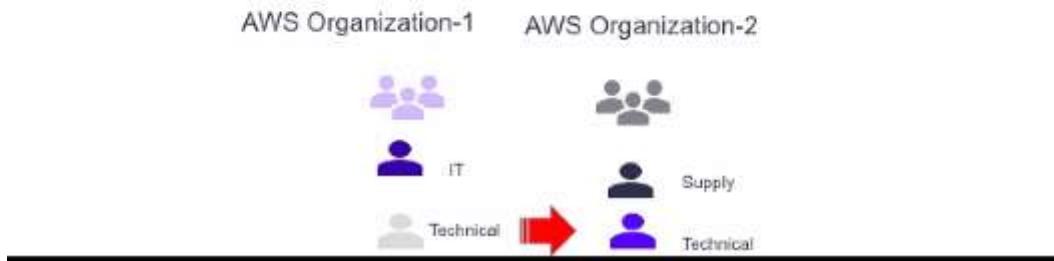
One AWS Organization owns several AWS accounts. Recently, due to a change of company organizations, one member account needs to be moved from this AWS Organization to another one. How can you achieve this?

- A. In the AWS console, drag and drop this account from one Organization to another.
- B. In the AWS console, select the member account and migrate it to the destination AWS Organization.
- C. Delete the old AWS Organization. Send an invite from the new Organization and accept the invite for the member account.
- D. Remove the member account from the old Organization. Send an invite from the new Organization to the member account and accept the invite.

Cevap D

MEMBER AND MASTER ACCOUNT LEAVING PROCESS

1. Remove the **member account** from the old Organization.
2. Send an invite to the **member account** from the new Organization.
3. Accept the invite to the new Organization from the **member account**.
4. Delete the old Organization.
5. **Send an invite to the master account**
6. Accept the invite to the new Organization from the **master account**



Question 18 of 65

Exit Quiz

Domain: Design Secure Architectures

While managing permissions for the API Gateway, what could be used to ensure that the right level of permissions is given to Developers, IT Admins, and end-users? The permissions should be easily managed.

- A. Use the secure token service to manage the permissions for different users.
- B. Use IAM Permissions to create different policies for different types of users.
- C. Use the AWS Config tool to manage the permissions for different users.
- D. Use IAM Access Keys to create sets of keys for different types of users.

Cevap B

Create policy		Policy actions *	
Filter policies		Q API	Showing 12 results
Policy name	Description		
AmazonAPIGatewayAdministrator	Provides full access to create/edit/delete APIs in Amazon API Gateway via the AWS Management Console.		
AmazonAPIGatewayInvokeFullAccess	Provides full access to invoke APIs in Amazon API Gateway.		
AmazonAPIGatewayPushToCloudWatchLogs	Allows API Gateway to push logs to user's account.		
AmazonAugmentedAIIntegratedAPIAccess	Provides access to perform all operations Amazon Augmented AI resources, including RowDefinitions, HumanTaskUs...		
AmazonDynamoDBFullAccessWithDataPipeline	Provides full access to Amazon DynamoDB including Export/Import using AWS Data Pipeline via the AWS Manage...		
AmazonEC2RoleforDataPipelineRole	Default policy for the Amazon EC2 Role for Data Pipeline service role.		
AmazonMQApiFullAccess	Provides full access to AmazonMQ via our API/SDK.		
AmazonMQApiReadOnlyAccess	Provides read only access to AmazonMQ via our API/SDK.		
APIGatewayServiceRolePolicy	Allows API Gateway to manage associated AWS Resources on behalf of the customer.		
AWSDataPipeline_FullAccess	Provides full access to Data Pipeline, list access for S3, DynamoDB, Redshift, RDS, SNS, and IAM roles, and passRole...		
AWSDataPipeline_PowerUser	Provides full access to Data Pipeline, list access for S3, DynamoDB, Redshift, RDS, SNS, and IAM roles, and passRole...		
AWSDataPipelineRole	Default policy for the AWS Data Pipeline service role.		

Question 19 of 65

Exit Quiz

Domain: Design High-Performing Architectures

You have an Amazon Route 53 alias record that routes the traffic to an Application Load Balancer. Later on, the availability zones enabled for the load balancer are changed by a team member. When you check the load balancer using the dig command, you find that the IPs of the ELB have changed. What kind of change do you need to do for the alias record in Route 53?

- A. Change the record type from A to CNAME.
- B. Modify the destination to the DNS name of the Application Load Balancer.
- C. Add the new IP addresses in the destination of the alias record.
- D. Nothing, as Route 53 automatically recognizes changes in the resource for the alias record.

Cevap D

Which type of record?

Create Domain Variations via Sub Domains	Mapping Value	Value Type	Record Type
www.clarusway.us	 Point out	<ul style="list-style-type: none">● IP of Server 1.2.3.4.5● Another Domain www.xxxxx.com● AWS End point S3 Bucket url Load Balancer DNS CloudFront● Etc.	<ul style="list-style-type: none">→ A AAA→ CNAME→ Alias→ MX

Record Value type determines the record type

Question 20 of 65

[Exit Quiz](#)

Domain: Design High-Performing Architectures

There is an urgent requirement to monitor some database metrics for a database hosted on AWS and send notifications. Which AWS services can accomplish this? (Select TWO)

- A. Amazon Simple Email Service
- B. Amazon CloudWatch
- C. Amazon Simple Queue Service
- D. Amazon Route 53
- E. Amazon Simple Notification Service

Cevap B ve E

There is an urgent requirement to monitor some database metrics for a database hosted on AWS and send notifications. Which AWS services can accomplish this? (Select Two)

- A. Amazon Simple Email Service
- B. Amazon CloudWatch
- C. Amazon Simple Queue Service
- D. Amazon Route 53
- E. Amazon Simple Notification Service

Question 21 of 65

Exit Quiz

Domain: Design High-Performing Architectures

You have the following architecture deployed in AWS.

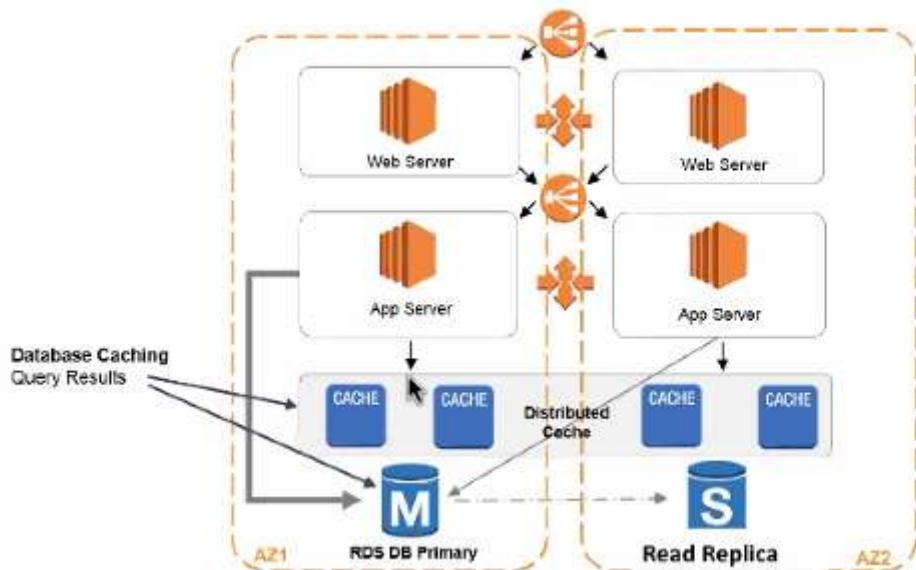
- a) A set of EC2 Instances which sit behind an ELB
- b) A database hosted in Amazon RDS

Of late, the performance on the database has been lacking due to a high number of read requests. Which of the following can be added to the architecture to alleviate the given performance issue? (Select TWO)

- A. Add read replica to the primary database to offload read traffic.
- B. Use ElastiCache in front of the database.
- C. Use AWS CloudFront in front of the database.
- D. Use Amazon DynamoDB to offload all the reads. Populate the common read items in a separate table.

Cevap A ve B

Database Caching Diagram



Question 22 of 65

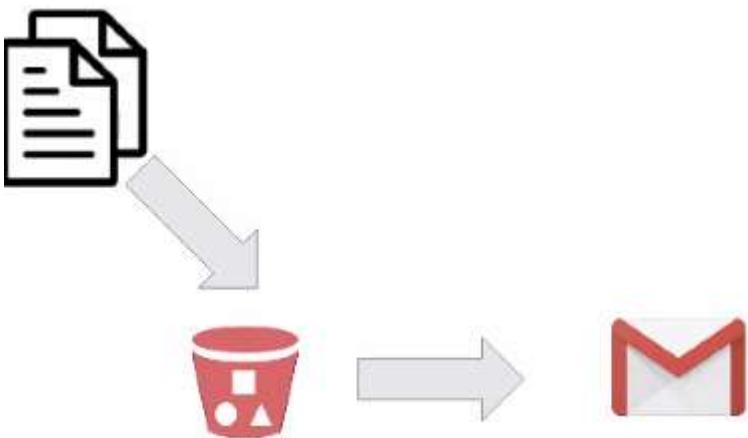
Exit Quiz

Domain: Design Secure Architectures

You have an S3 bucket that is used to store important data for a web application. You want to receive an email notification whenever an object removal event happens in the S3 bucket. How would you configure the S3 bucket to achieve this requirement?

- A. Configure the object-level logging for the S3 bucket and register an SNS topic to provide notifications.
- B. Configure the server access logging for the object removal events. Add an SNS topic to notify the team via emails.
- C. Set up an AWS Config rule to check the object deletion events. Register a Lambda function to send notifications.
- D. Configure an S3 event notification for the object removal events. Send the events to an SNS topic.

Cevap D



Question 23 of 65

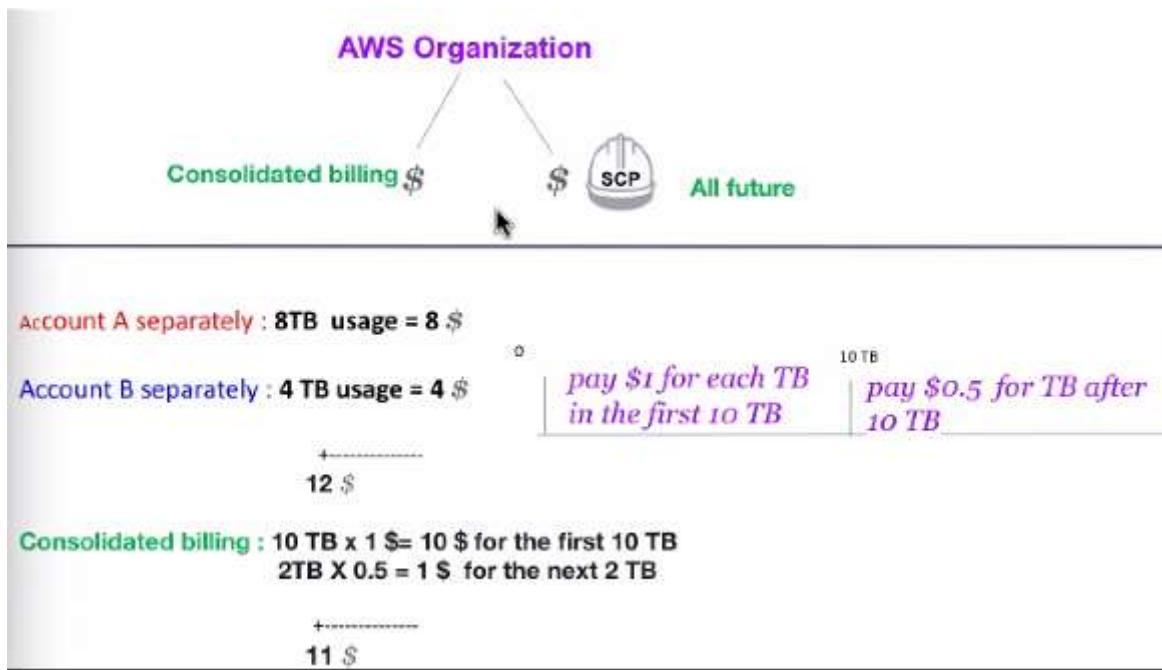
[Exit Quiz](#)

Domain: Design Cost-Optimized Architectures

To manage a large number of AWS accounts in a better way, you create a new AWS Organization and invite multiple accounts. You only enable the "Consolidated billing" out of the two feature sets (**All features** and **Consolidated billing**) available in the AWS Organizations. Which of the following is the primary benefit of using Consolidated billing feature?

- A. Apply SCPs to restrict the services that IAM users can access.
- B. Configure tag policies to maintain consistent tags for resources in the organization's accounts.
- C. Configure a policy to prevent IAM users in the organization from disabling AWS CloudTrail.
- D. Combine the usage across all accounts to share the volume pricing discounts.

Cevap D



Question 24 of 65

Exit Quiz

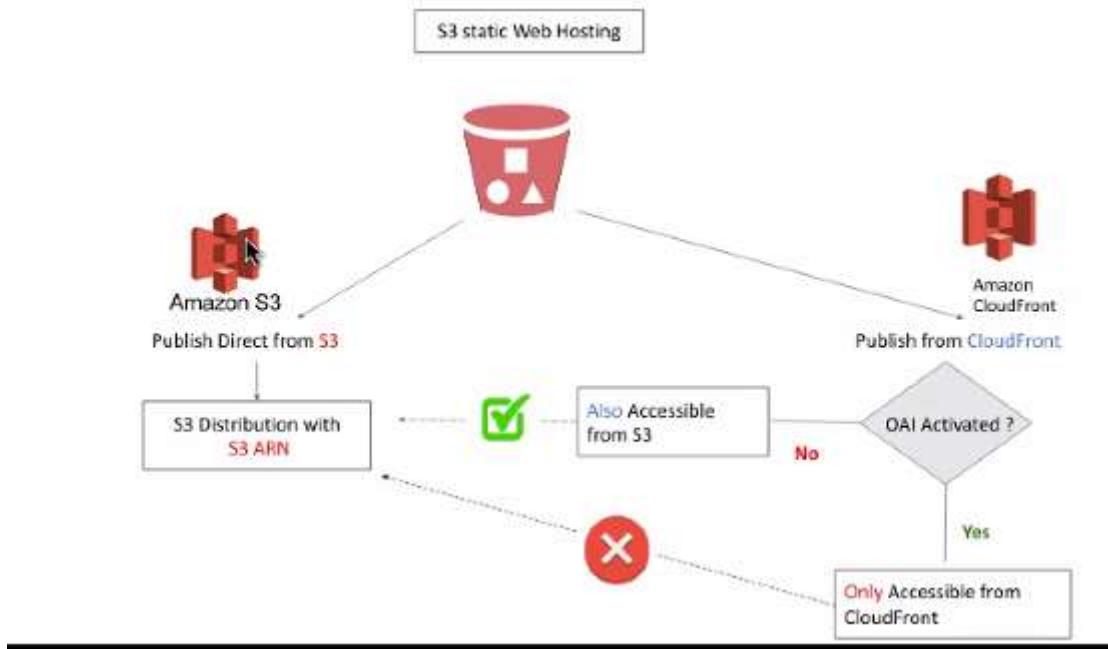
Domain: Design Secure Architectures

A Solutions Architect has been asked to design a solution that will deliver digital content to users through Amazon CloudFront. You have set up an Amazon S3 bucket as the origin and by default, CloudFront never exposes Amazon S3 URLs. The contents should only be accessed through the CloudFront distribution. How can this be achieved?

- A. Store the digital contents in the S3 bucket and create signed URLs to access S3 through CloudFront.
- B. Create OAI in CloudFront. Use the S3 bucket policy to ensure that only the OAI can access the files in the Amazon S3 bucket.
- C. Store the digital contents as private objects in the S3 buckets and use the S3 ACL to ensure that only the CloudFront distribution ARN can access the bucket.
- D. Store the digital contents in the S3 bucket and configure signed cookies for users to access contents through CloudFront.

Cevap B

S3 den gelen veriyi cloudfonttan hızlandırmak istiyorsun. Ama hem s3 url hem cloufront url var nasıl önlersin bunu



Question 25 of 65

[Exit Quiz](#)

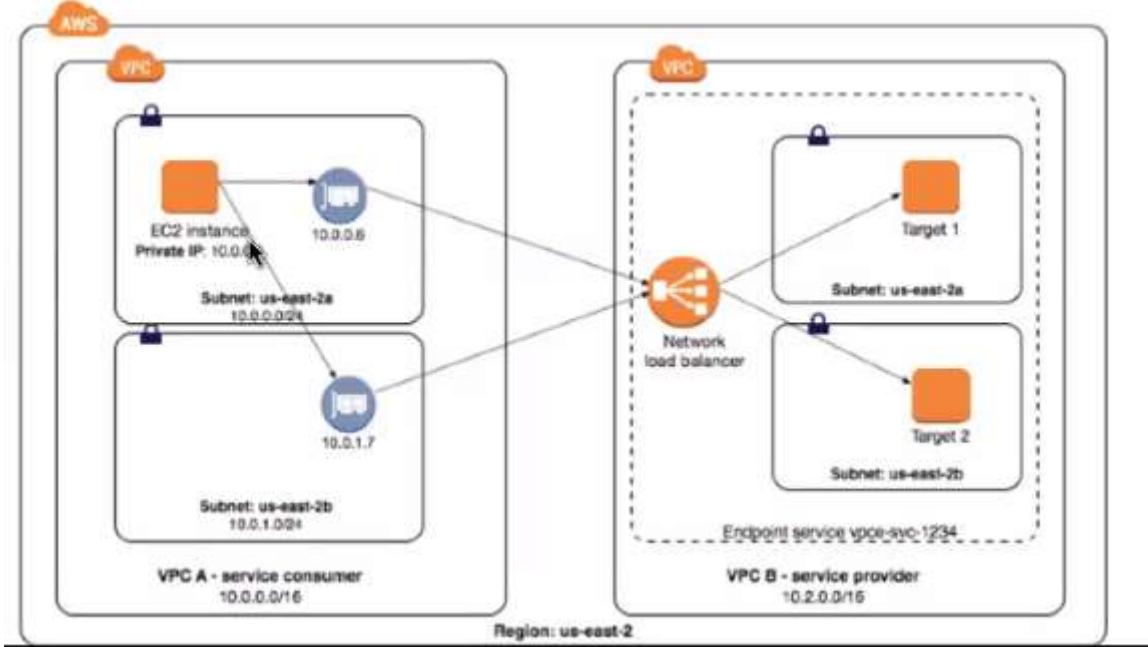
Domain: Design Resilient Architectures

You are working as an AWS Architect for an IT Company. Your Company is using EC2 instances in multiple VPCs spanning Availability Zones in us-east-1 Region. The Development Team has deployed a new Intranet application that needs to be accessed via VPC.

You need to make sure that the connectivity to this particular application uses the internal AWS network between different VPCs, and that the solution is highly scalable and secure. Which of the following solution would you recommend?

- A. Attach an Internet Gateway to all the VPCs in the us-east-1 region and allow all users to access this application over the internet.
- B. Deploy Network Load Balancers along with VPC endpoint service (AWS PrivateLink) to establish connectivity between the VPCs in the us-east-1 region.
- C. Use the VPC Gateway Endpoint service between all the VPCs in the us-east-1 region to provide connectivity between users & servers.
- D. Create a VPN between instances at the various VPCs in the us-east-1 region to establish connectivity.

Cevap B



Question 26 of 65

[Exit Quiz](#)

Domain: Design Secure Architectures

You work in a large organization. Your team creates AWS resources such as Amazon EC2 dedicated hosts and reserved capacities that need to be shared by other AWS accounts. You need an AWS service to centrally manage these resources so that you can easily specify which accounts or Organizations can access the resources. Which AWS service would you choose to meet this requirement?

- A. IAM
- B. Resource Access Manager
- C. Service Catalog
- D. AWS Single Sign-On

Cevap B

Services

- AWS App Mesh
 - Amazon Aurora
 - AWS Certificate Manager Private Certificate Authority
 - AWS CodeBuild
 - **Amazon EC2**
 - EC2 Image Builder
 - AWS Glue
 - AWS License Manager
 - AWS Network Firewall
 - AWS Outposts
 - AWS Resource Groups
 - Amazon Route 53
 - Amazon VPC
-

Question 27 of 65

Exit Quiz

Domain: Design High-Performing Architectures

Your company wants to use an S3 bucket for web hosting but has several different domains to perform operations on the S3 content. In the CORS configuration, you have added CORSRule AllowedOrigin for the following Domains: <http://www.domainnamea.com>, <https://www.secure.domainnamea.com>, and <http://www.domainnameb.com>. Following Domains, <https://domainnameb.com> and <http://www.domainnameb.com:80>, are not allowed to access the S3 bucket.

What could be the most likely cause behind the unexpected access behaviour of the domains?

- A. Both request <https://domainnameb.com> and <http://www.domainnameb.com:80> don't match the allowed origin in your configuration.
- B. HTTPS must contain a specific port in the request, e.g. <https://domainnameb.com:443>
- C. There's a limit of two origin sites per S3 bucket allowed
- D. Adding CORS automatically removes the S3 ACL and bucket policies

Cevap A identical olmalı

CORS Domains:

http://www.domainnamea.com,

https://www.secure.domainnamea.com,

http://www.domainnameb.com.

Attempts

https://www.domainnameb.com

http://www.domainnameb.com:80

Question 28 of 65

[Exit Quiz](#)

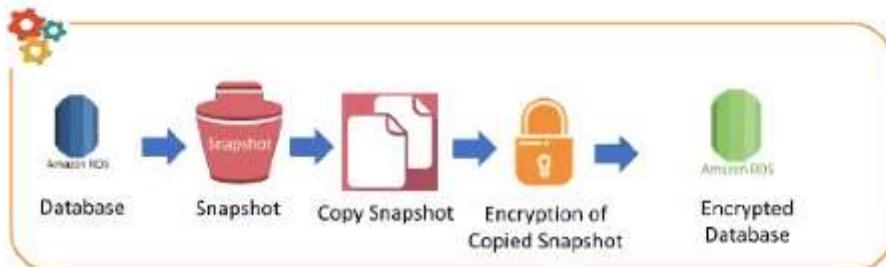
Domain: Design Secure Architectures

A company has a PostgreSQL DB instance in Amazon RDS which is not encrypted. As per security policy, data in the RDS instances should be encrypted at rest with AWS KMS.

Which option is correct for RDS DB encryption?

- A. Amazon RDS for PostgreSQL DB instance can only be encrypted at creation time and not after its creation. There is no way to achieve this requirement.
- B. Take a snapshot of the unencrypted DB instance. Copy the snapshot and encrypt the new snapshot with AWS KMS. Restore the DB instance with the new encrypted snapshot.
- C. Take a snapshot of the unencrypted DB instance. Encryption can be enabled by restoring a DB instance from the unencrypted snapshot.
- D. Stop the existing RDS instance and encrypt the DB with a KMS CMK.

Cevap B snapot alınır ve kopyalarken şifreleme çıkar



Question 29 of 65

[Exit Quiz](#)

Domain: Design Resilient Architectures

You currently manage a set of web servers hosted on EC2 instances with public IP addresses. These IPv4 addresses are mapped to domain names. There was an urgent maintenance activity that need to be carried out on the servers. The servers had to be stopped and restarted. After the maintenance, the web application hosted on these EC2 instances is not accessible via the domain names configured earlier.

Which of the following could be a reason for this?

- A. The Route 53 hosted zone needs to be restarted.
- B. The Elastic IP address needs to be initialized again.
- C. The public IP addresses need to be associated with the ENI (Elastic network interfaces) again.
- D. The public IP addresses have changed after the instance was stopped and started again.

Cevap D

You currently manage a set of web servers hosted on EC2 Servers with public IP addresses. These IP addresses are mapped to domain names. There was an urgent maintenance activity that had to be carried out on the servers. **The servers had to be stopped and restarted.** Now the web application hosted on these EC2 Instances is not accessible via the domain names configured earlier. Which of the following could be a reason for this?

- A. The Route 53 hosted zone needs to be restarted.
- B. The network interfaces need to initialized again.
- C. The public IP addresses need to be associated with the ENI again.
- D. The public IP addresses have changed after the instance was stopped and started again.

Question 30 of 65

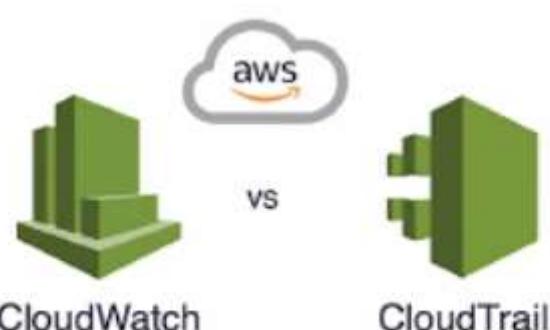
[Exit Quiz](#)

Domain: Design Secure Architectures

You are responsible for deploying a critical application to AWS. It is required to monitor web application logs to identify any malicious activity. Also, there is a need to store log data in highly durable storage. Which of the following services could be used to fulfill this requirement?

- A. Amazon CloudWatch Logs
- B. AWS Personal Health Dashboard
- C. Amazon Trusted Advisor
- D. Amazon CloudTrail

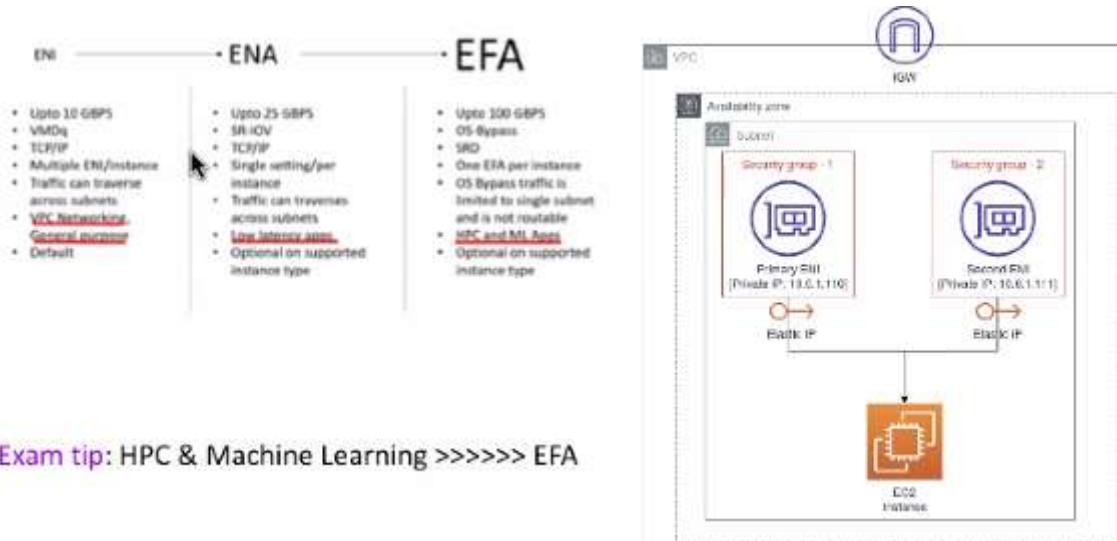
Cevap A cloud trail aktiviteleri kontrol eder. Kim yapmış gibi



Domain: Design High-Performing Architectures

You need to deploy a high performance computing (HPC) and machine learning application in AWS Linux EC2 instances. The performance of inter-instance communication is very critical for the application. You want to attach a network device to the instance so that the computing performance can be greatly improved. Which of the following options can achieve the best performance?

- A. Enable enhanced networking feature in the EC2 instance.
- B. Configure Elastic Fabric Adapter (EFA) in the instance.
- C. Attach high speed Elastic Network Interface (ENI) in the instance.
- D. Create Elastic File System (EFS) and mount the file system in the instance.

Cevap B

Domain: Design High-Performing Architectures

A company is planning on testing a large set of IoT-enabled devices. These devices will generate a large amount of data every second. You need a scalable and durable real-time data streaming service to capture the data generated from these devices. Which AWS service would be the most appropriate for this purpose?

- A. AWS EMR.
- B. AWS Kinesis Data Streams.
- C. AWS SQS.
- D. AWS SNS.

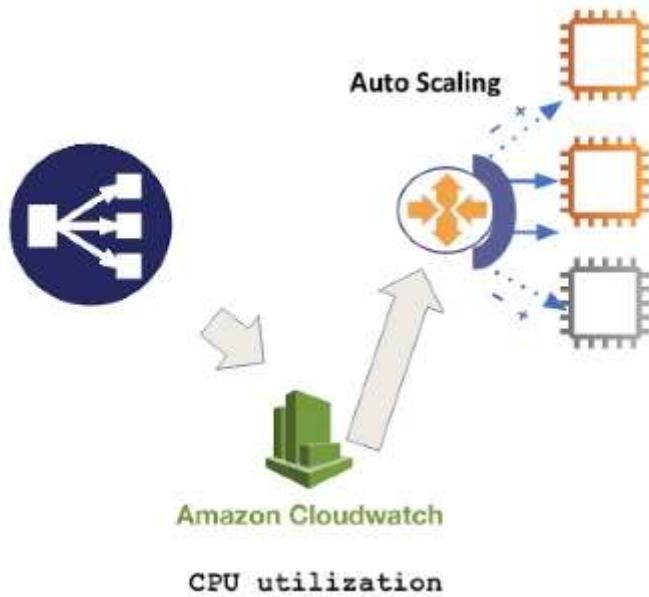
Cevap B

Domain: Design Resilient Architectures

Your company currently has a set of non-production EC2 instances hosted in AWS. To save costs, you want to stop the EC2 instance when the average CPU utilization percentage has been lower than 10 percent for 24 hours, signaling that it is idle and no longer in use. Which step could be helpful to fulfill this requirement?

- A. Use CloudWatch Logs to store the state change of the instances.
- B. Create Amazon CloudWatch alarms that monitor the CPU utilization metric and stop the instances when the alarms are triggered.
- C. Use SQS to monitor the metric and add the record to a DynamoDB table.
- D. Use AWS Lambda to monitor the metric and store the state in a DynamoDB table.

Cevap B



Question 34 of 65

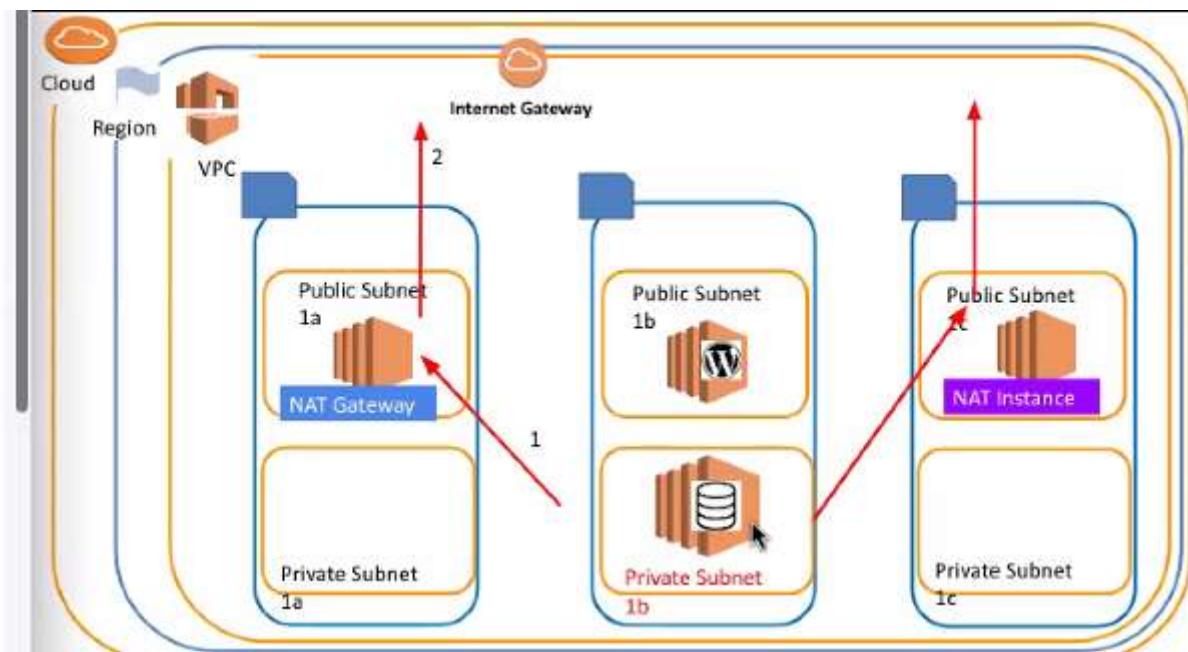
[Exit Quiz](#)

Domain: Design Secure Architectures

You have instances hosted in a private subnet in a VPC. There is a need for instances to download updates from the Internet. As an architect, what change would you suggest to the IT Operations team that would also be the most efficient and secure?

- A. Create a new public subnet and move the instance to that subnet.
- B. Create a new EC2 Instance to download the updates separately and then push them to the required instance.
- C. Use a NAT Gateway to allow the instances in the private subnet to download the updates.
- D. Create a VPC link to the Internet to allow the instances in the private subnet to download the updates.

Cevap C



Question 35 of 65

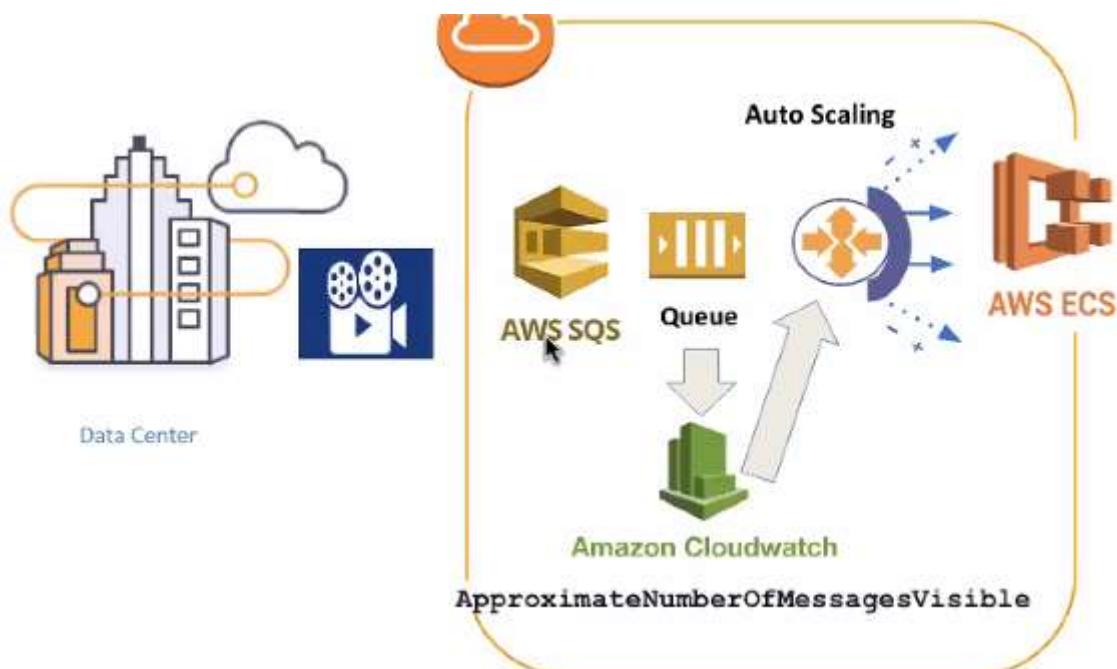
[Exit Quiz](#)

Domain: Design Resilient Architectures

You have an S3 bucket that receives photos uploaded by customers. When an object is uploaded, an event notification is sent to an SQS queue with the object details. You also have an ECS cluster that gets messages from the queue to do the batch processing. Each of the batch processing job takes the same amount of time to get executed. The queue size may change greatly depending on the number of incoming messages and backend processing speed. Which metric would you use to scale up/down the ECS cluster capacity?

- A. The number of messages in the SQS queue.
- B. Memory usage of the ECS cluster.
- C. Number of objects in the S3 bucket.
- D. Number of containers in the ECS cluster.

Cevap A



Domain: Design High-Performing Architectures

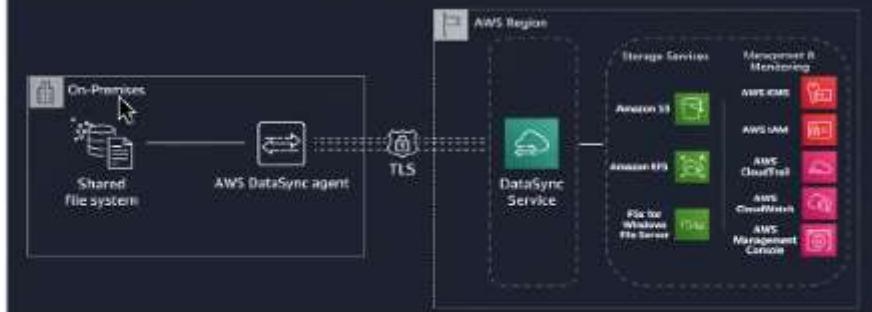
You are a solutions architect working for a regional bank that is moving its data center to the AWS cloud. You need to migrate your data center storage to a new S3 and EFS data store in AWS. Since your data includes Personally Identifiable Information (PII), you have been asked to transfer data from your data center to AWS without traveling over the public internet. Which option gives you the most efficient solution that meets your requirements?

- A. Migrate your on-prem data to AWS using the DataSync agent using NAT Gateway.
- B. Create a private VPC endpoint, and configure the DataSync agent to communicate to the private DataSync service endpoints via the VPC endpoint using Direct Connect
- C. Migrate your on-prem data to AWS using the DataSync agent using Internet Gateway.
- D. Create a public VPC endpoint, and configure the DataSync agent to communicate to the DataSync private service endpoints via your VPC endpoint via your VPN.

Cevap B

How does AWS DataSync work?

Simplifies, automates, and accelerates data transfer to or from AWS



- Public endpoints:
- Federal Information Processing Standard (FIPS) endpoints: USA and Canada
- Virtual private cloud (VPC) endpoints: If you use a VPC endpoint, all communication from DataSync to AWS occurs through the endpoint in your AWS VPC. This establishes a private connection between your self-managed storage system, your VPC, and AWS services, providing extra security as your data is copied over the network.

Question 37 of 65

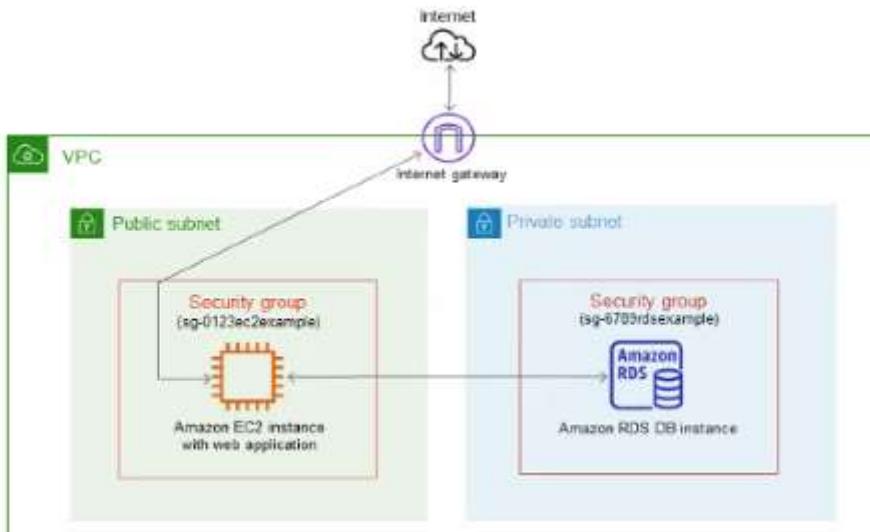
[Exit Quiz](#)

Domain: Design Secure Architectures

You have planned to host a web application on AWS. You create an EC2 instance in a public subnet that needs to connect to an EC2 instance that will host an Oracle database. Which steps would ensure a secure setup? (SELECT TWO)

- A. Place the EC2 instance with the Oracle database in the same public subnet as the Webserver for faster communication.
- B. Place the ec2 instance that will host the Oracle database in a private subnet.
- C. Create a database Security group which allows incoming traffic only from the Web server's security group.
- D. Ensure that the database security group allows incoming traffic from 0.0.0.0/0.

Cevap B ve C db private olur.

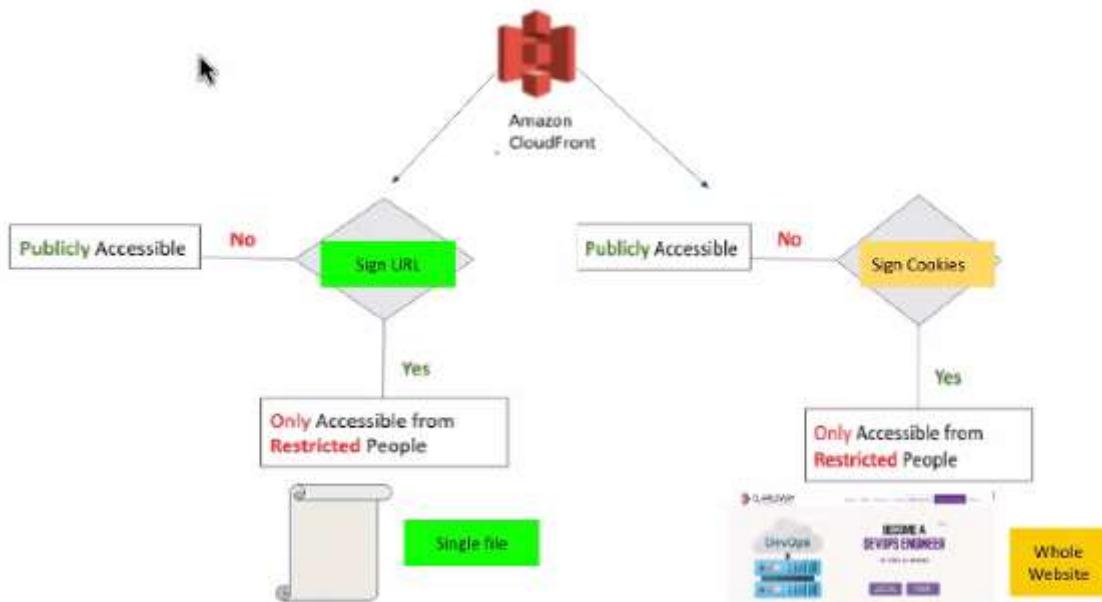


Domain: Design Secure Architectures

You are designing a website for a company that streams anime videos. You serve this content through CloudFront. The company has implemented a section for premium subscribers. This section contains more videos than the free section. You want to ensure that only premium subscribers can access this premium section. How can you achieve this easily?

- A. Using bucket policies.
- B. Requiring HTTPS for communication between users and CloudFront.
- C. Using CloudFront origin with signed URLs.
- D. Using CloudFront origin with signed cookies.

Cevap D birden fazla file olduğu için



Question 39 of 65

[Exit Quiz](#)

Domain: Design High-Performing Architectures

A company has set up an application in AWS that interacts with DynamoDB. It is required that when an item is modified in a DynamoDB table, an immediate entry has to be made to the associating application. How can this be accomplished? (SELECT TWO)

- A. Set up CloudWatch to monitor the DynamoDB table for changes. Then trigger a Lambda function to send the changes to the application.
- B. Set up CloudWatch logs to monitor the DynamoDB table for changes. Then trigger AWS SQS to send the changes to the application.
- C. Use DynamoDB streams to monitor the changes to the DynamoDB table.
- D. Trigger a lambda function to make an associated entry in the application as soon as the DynamoDB streams are modified.

Cevap C



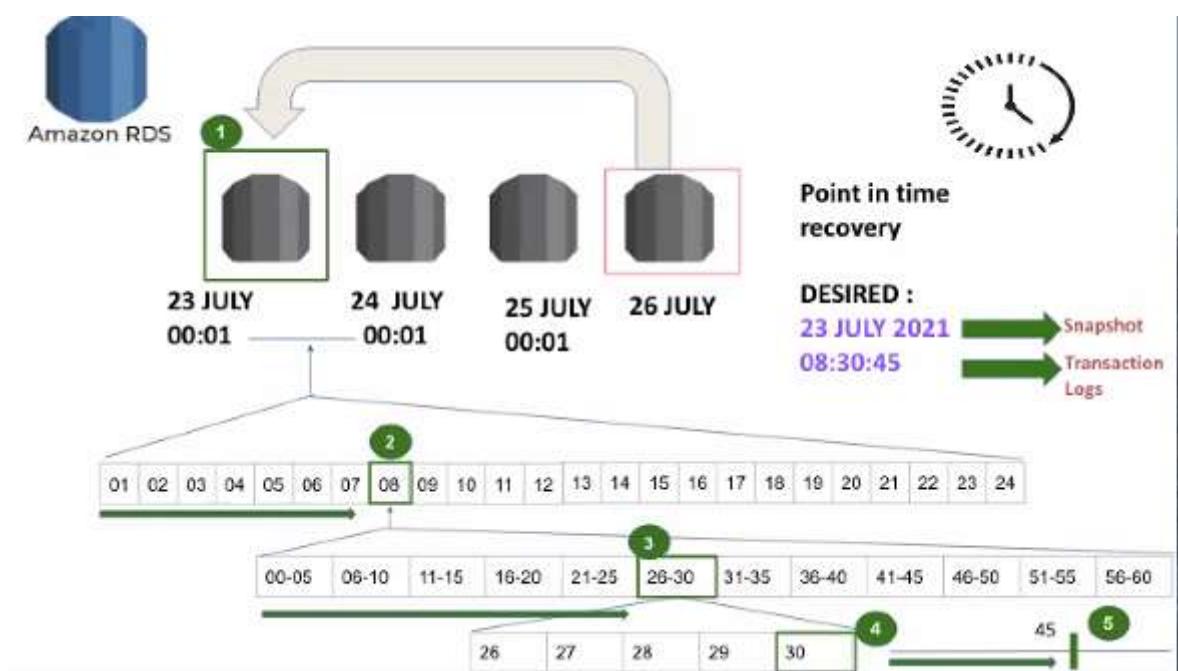
When enabled, DynamoDB Streams captures a time-ordered sequence of item-level modifications in a DynamoDB table and durably stores the information for **up to 24 hours**. Applications can access a series of **stream records**, which contain an item change, from a **DynamoDB stream** in near real time.

Domain: Design Resilient Architectures

You are working as an AWS consultant in an E-Commerce organization. Your organization is planning to migrate its database from on-premises data centers to Amazon RDS. The automated backup helps to restore the Database to any specific time during the backup retention period in Amazon RDS. Which of the following actions are performed as a part of the Amazon RDS automated backup process?

- A. AWS creates a storage volume snapshot of the database instance during the backup window once a day. AWS RDS also captures transactions logs and uploads them to S3 buckets every 5 minutes.
- B. AWS creates a full snapshot of the database every 12 hours during the backup window, captures transactions logs throughout the day, and stores them in S3 buckets.
- C. AWS creates a full daily snapshot during the backup window. With the snapshot, the RDS instance can be restored at any time.
- D. AWS creates a storage volume snapshot of the database instance every 12 hours during the backup window, captures transactions logs throughout the day, and stores them in S3 buckets.

Cevap A



Question 41 of 65

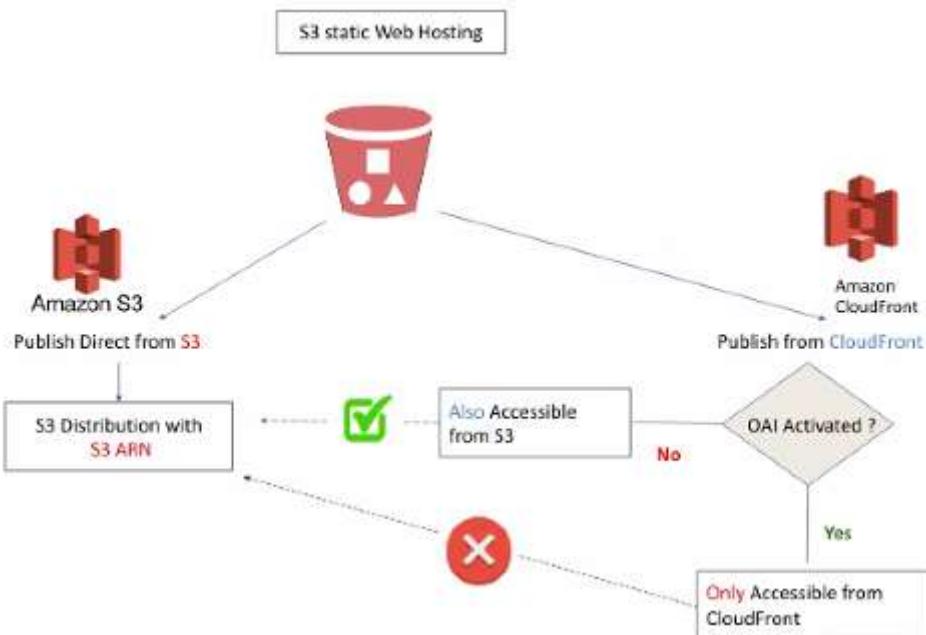
[Exit Quiz](#)

Domain: Design Secure Architectures

You configure an Amazon S3 bucket as the origin for a new CloudFront distribution. The traffic should not hit the S3 URLs directly instead, they should be directed to the CloudFront distribution and the files should be fetched through the CloudFront URL. Which method is the most appropriate?

- A. Configure Signed URLs to serve private content by using CloudFront.
- B. Configure Signed Cookies to restrict access to S3 files.
- C. Create the origin access identity (OAI) and associate it with the distribution.
- D. Configure the CloudFront web distribution to ask viewers to use HTTPS to request S3 objects.

Cevap C



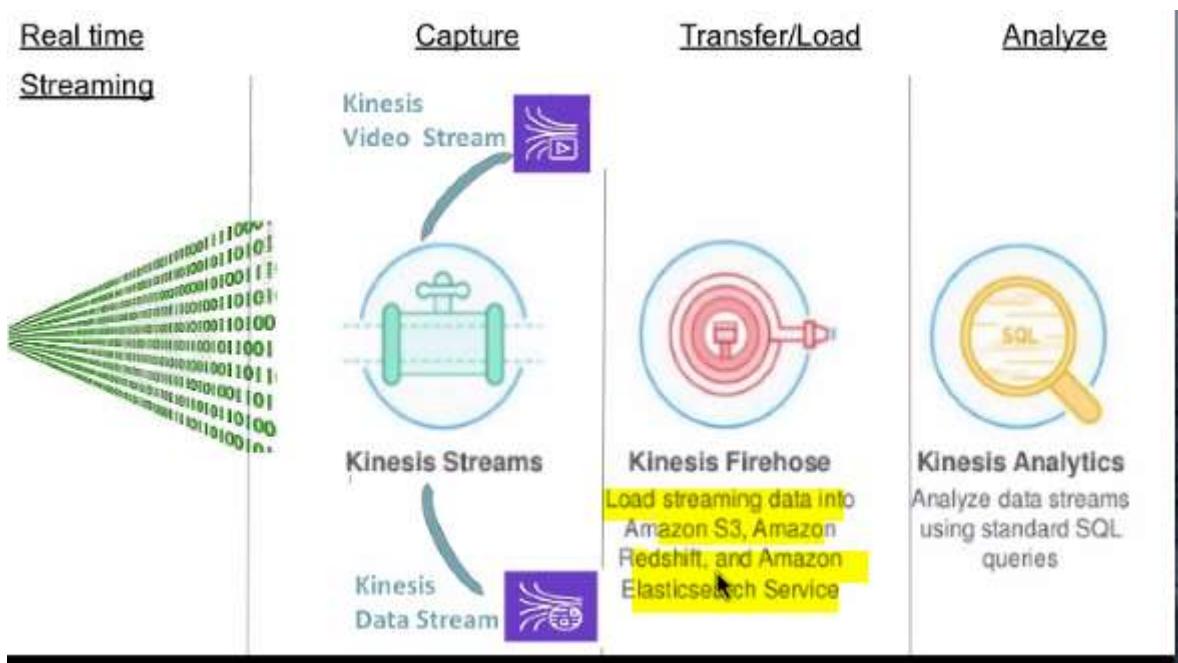
As a Solutions Architect for a multinational organization with more than 150000 employees, management has decided to implement a real-time analysis for their employees' time spent in offices worldwide. You are tasked to design an architecture that will receive the inputs from 10000+ sensors with swipe machine sending in and out data across the globe, each sending 20KB data every 5 Seconds in JSON format. The application will process and analyze the data and upload the results to dashboards in real-time.

Other application requirements will include the ability to apply real-time analytics on the captured data. Processing of captured data will be parallel and durable. The application must be scalable as per the requirement as the load varies and new sensors are added or removed at various facilities. The analytic processing results are stored in a persistent data storage for data mining.

What combination of AWS services would be used for the above scenario?

- A. Use EMR to copy the data coming from Swipe machines into DynamoDB and make it available for analytics.
- B. Use Amazon Kinesis Data Streams to ingest the Swipe data coming from sensors, Use custom Kinesis Data Streams Applications to analyze the data and then move analytics outcomes to RedShift using AWS EMR.
- C. Use SQS to receive the data coming from sensors, Kinesis Firehose to analyze the data from SQS, then save the results to a Multi-AZ RDS instance.
- D. Use Amazon Kinesis Data Streams to ingest the sensors' data, Use custom Kinesis Streams applications to analyze the data, and move analytics outcomes to RDS using AWS EMR.

Cevap B



Question 43 of 65

[Exit Quiz](#)

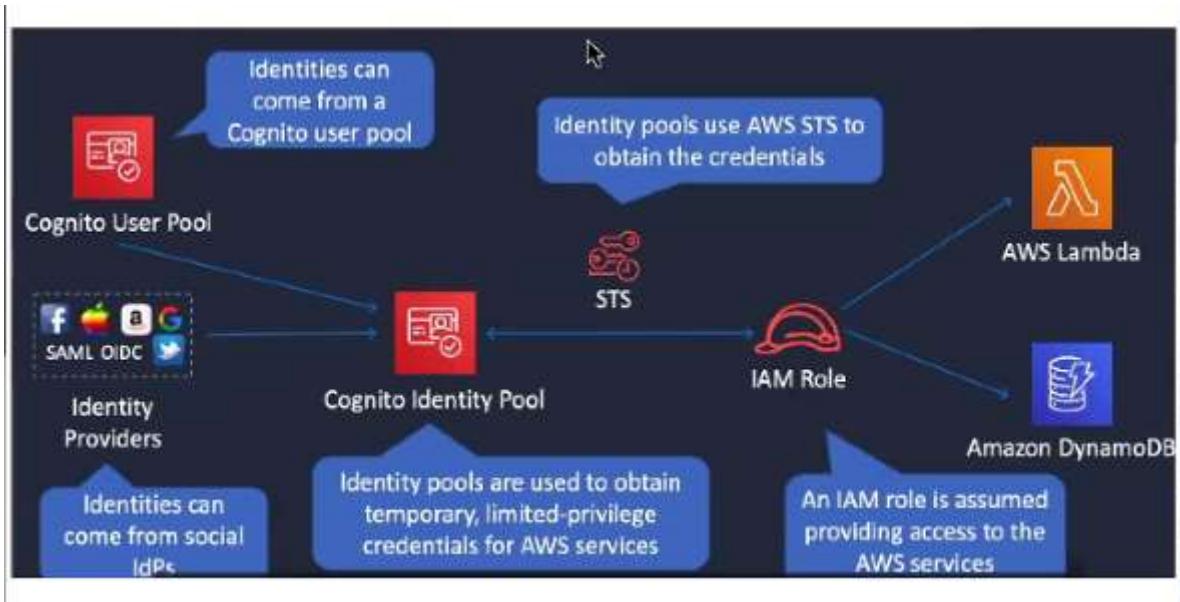
Domain: Design Secure Architectures

You have designed an application that uses AWS resources, such as S3, to operate and store users' documents. You currently use Cognito identity pools and user pools. To increase usage and ease of signing up, you decide that adding social identity federation is the best path forward.

How would you differentiate the Cognito identity pool and the federated identity providers (e.g. Google)?

- A. They are the same and just called different things.
- B. First, you sign-in via Cognito then through a federated site, like Google.
- C. Federated identity providers and identity pools are used to authenticate services.
- D. You can choose a federated identity provider to authenticate users and associate a Cognito identity pool to authorize the users.

Cevap D Cognito kimliklendirme işlenmelerini yapar Active directory gibi



Question 44 of 65

[Exit Quiz](#)

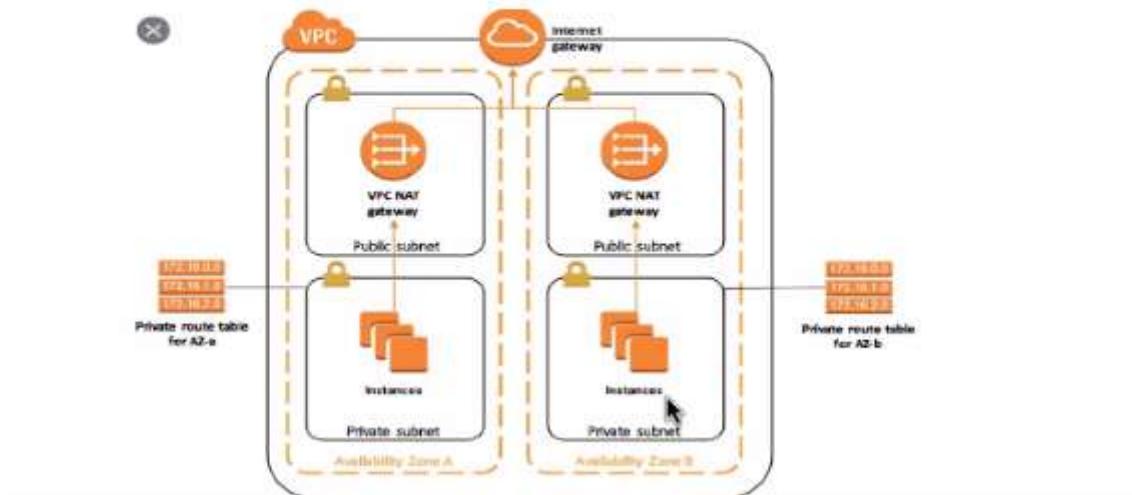
Domain: Design Resilient Architectures

You currently have your EC2 instances running in multiple availability zones in an AWS region. You need to create NAT gateways for your private instances to access internet. How would you set up the NAT gateways so that they are highly available?

- A. Create two NAT Gateways and place them behind an ELB.
- B. Create a NAT Gateway in each Availability Zone.
- C. Create a NAT Gateway in another region.
- D. Use Auto Scaling groups to scale the NAT Gateways.

Cevap B

If you have resources in multiple Availability Zones and they share one NAT Gateway, in the event that the NAT Gateway's Availability Zone is down, resources in the other Availability Zones lose internet access. To create an Availability Zone-independent architecture, create a NAT Gateway in each Availability Zone and configure your routing to ensure that resources use the NAT Gateway in the same Availability Zone.



Question 45 of 65

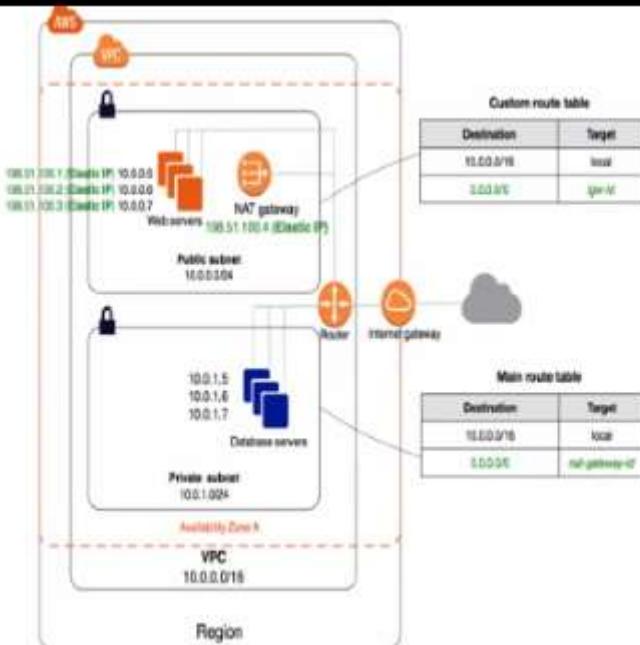
[Exit Quiz](#)

Domain: Design High-Performing Architectures

A Solutions Architect is designing an online shopping application running in a VPC on EC2 Instances behind an Elastic Load Balancer. The instances run in an Auto Scaling group across multiple Availability Zones. The application tier must read and write data to a customer-managed database cluster. There should be no access to the database from the Internet. But the cluster must be able to obtain software patches from the Internet. Which of the following VPC design meets the requirements?

- A. Create public subnets for the application tier and the database cluster.
- B. Create public subnets for the application tier and private subnets for the database cluster.
- C. Create public subnets with NAT Gateway for the application tier and private subnets for the database cluster.
- D. Create private subnets for the application tier, and private subnets with NAT Gateway for the database cluster.

Cevap C



Question 46 of 65

[Exit Quiz](#)

Domain: Design High-Performing Architectures

It is expected that only certain specified customers can upload images to the S3 bucket for a certain period of time. What would you suggest as an architect to fulfill this requirement?

- A. Create a secondary S3 bucket. Then, use an AWS Lambda to sync the contents to the primary bucket.
- B. Use pre-signed URLs for uploading the images.
- C. Use ECS Containers to upload the images.
- D. Upload the images to SQS and then push them to the S3 bucket.

Cevap B

The screenshot shows the AWS S3 console with a single object listed:

Name	Last modified	Type
Yeni Metin Belgesi.txt	May 23, 2021, 15:49:52 (UTC+01:00)	txt

Below the table, a command-line interface (CLI) command is shown:

```
$ aws s3 presign s3://mycloudtrails3bucket-mk/"Yeni Metin Belgesi.txt" --expires-in 30
```

The URL generated by the CLI command is highlighted with a red arrow:

https://mycloudtrails3bucket-mk.s3.us-east-1.amazonaws.com/Yeni%20Metin%20Belgesi.txt?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=AKIAY3YMFVJJYLFDLWV5%2F20210525%2Fs-east-1%2Fs3%2Faws4_request&X-Amz-Date=20210525T125743Z&X-Amz-Expires=30&X-Amz-SignedHeaders=host&X-Amz-Signature=b02d2c978829ff3bed34001380b8b9e0b51262727d2b37391c4b5c12e0d6f31c

Question 47 of 65

Exit Quiz

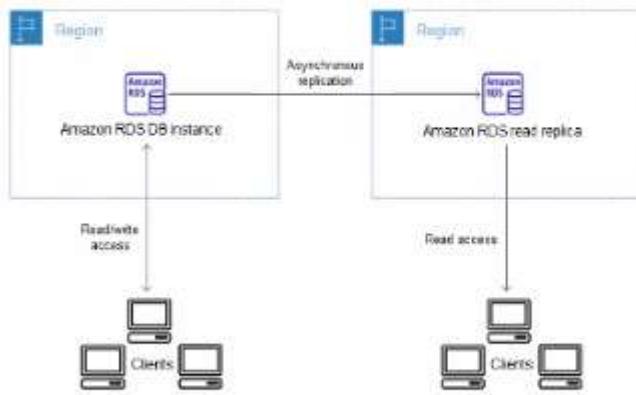
Domain: Design High-Performing Architectures

You own a MySQL RDS instance in AWS Region us-east-1. The instance has a Multi-AZ instance in another availability zone for high availability. As business grows, more and more clients come from Europe (eu-west-2), and most of the database workload is read-only. What is the proper way to reduce the load on the source RDS instance?

- A. Create a snapshot of the instance and launch a new instance in eu-west-2.
- B. Promote the Multi-AZ instance to be a Read Replica and move the instance to eu-west-2 region.
- C. Configure a read-only Multi-AZ instance in eu-west-2 as Read Replicas cannot span across regions.
- D. Create a Read Replica in the AWS Region eu-west-2.

Cevap D

With Amazon RDS, you can create a MariaDB, MySQL, Oracle, or PostgreSQL read replica in a different AWS Region from the source DB instance. Creating a cross-Region read replica isn't supported for SQL Server on Amazon RDS.



Availability - Failover-Senkron

Performance- Asenkron

Question 48 of 65

[Exit Quiz](#)

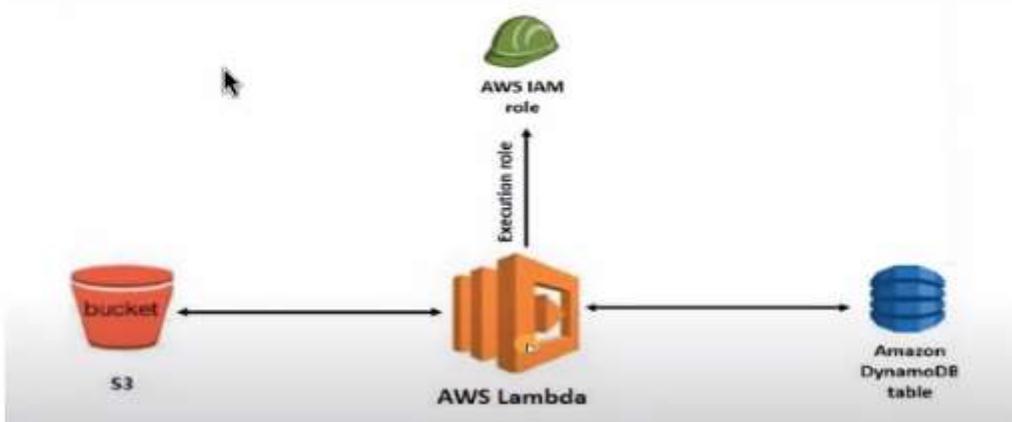
Domain: Design Resilient Architectures

Your company manages an application that currently allows users to upload images to an S3 bucket. These images are picked up by EC2 Instances for processing and then placed in another S3 bucket. You need an area where the metadata for those images can be stored. What would be an ideal data store for this?

- A. AWS Redshift
- B. AWS Glacier
- C. AWS DynamoDB
- D. AWS SQS

Cevap C

Exam tip: DynamoDb and Metadata



Question 49 of 65

[Exit Quiz](#)

Domain: Design High-Performing Architectures

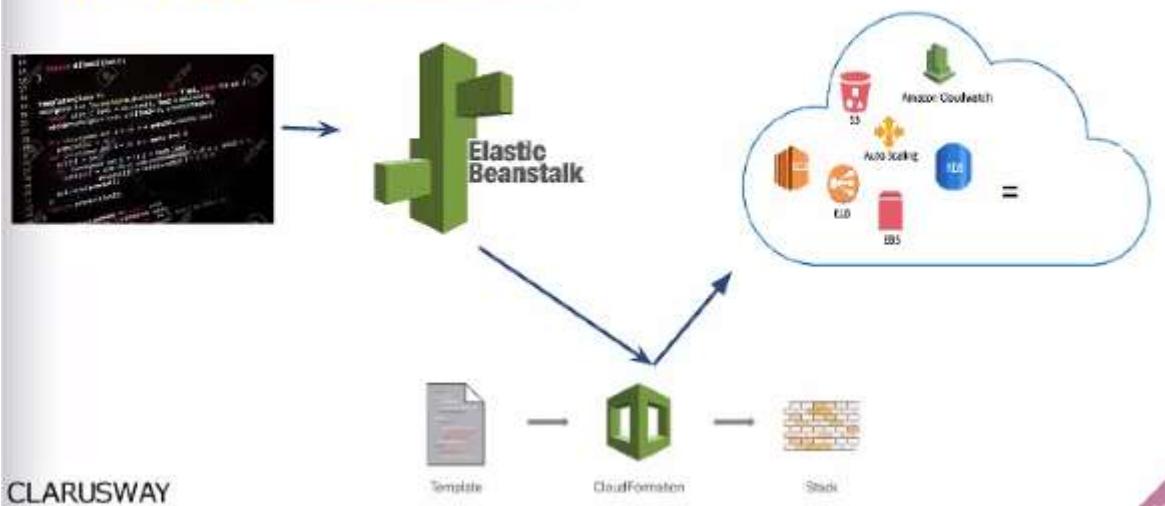
An application team needs to quickly provision a development environment consisting of a web and database layer. What would be the quickest and most ideal way to get this set up in place?

- A. Create Spot Instances and install the web and database components.
- B. Create Reserved Instances and install the web and database components.
- C. Use AWS Lambda to create the web components and AWS RDS for the database layer.
- D. Use Elastic Beanstalk to quickly provision the environment.

Cevap D

▶ Introduction to Elastic Beanstalk

Why AWS Elastic Beanstalk?



CLARUSWAY

Question 50 of 65

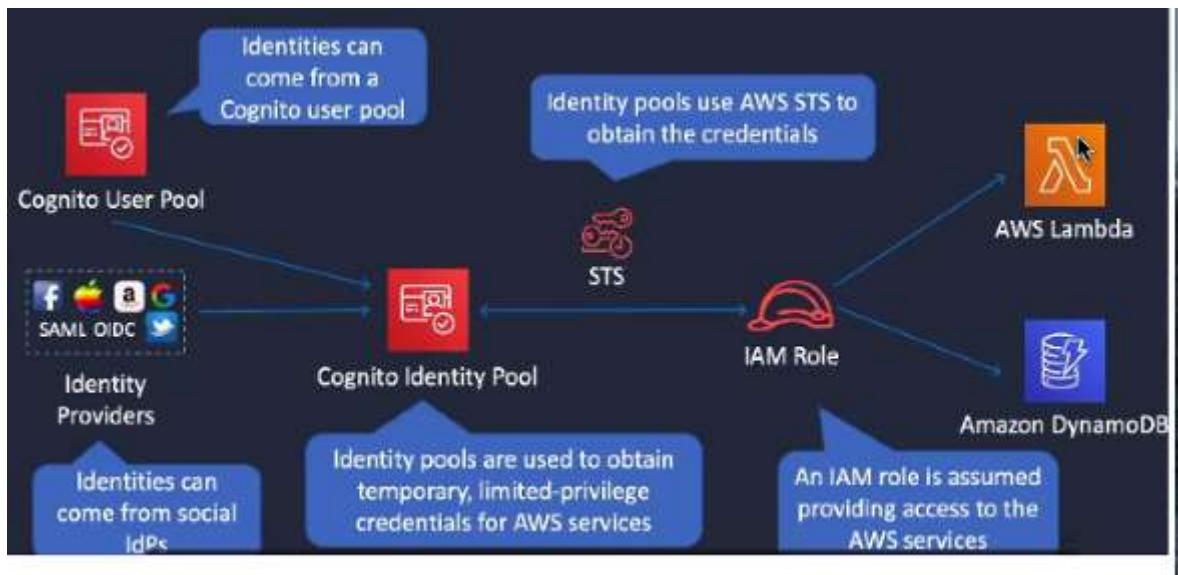
Exit Quiz

Domain: Design Secure Architectures

Third-party sign-in (Federation) has been implemented in your web application to allow users who need access to AWS resources. Users have been successfully logging in using Google, Facebook, and other third-party credentials. Suddenly, their access to some AWS resources has been restricted. What is the most likely cause of the restricted use of AWS resources?

- A. IAM policies for resources were changed, thereby restricting access to AWS resources.
- B. Federation protocols are used to authorize services and need to be updated.
- C. IAM groups for accessing the AWS resources were changed, thereby restricting their access via federated login.
- D. The identity providers no longer allow access to AWS services.

Cevap A

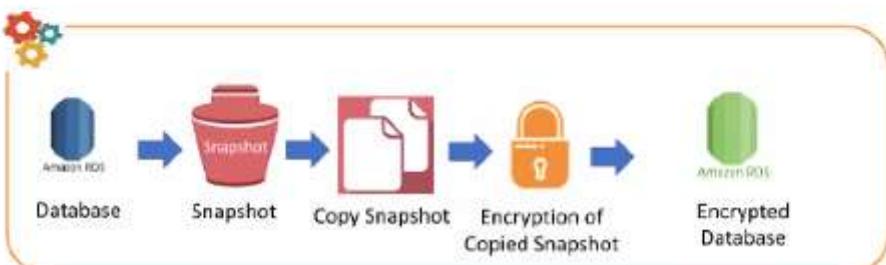
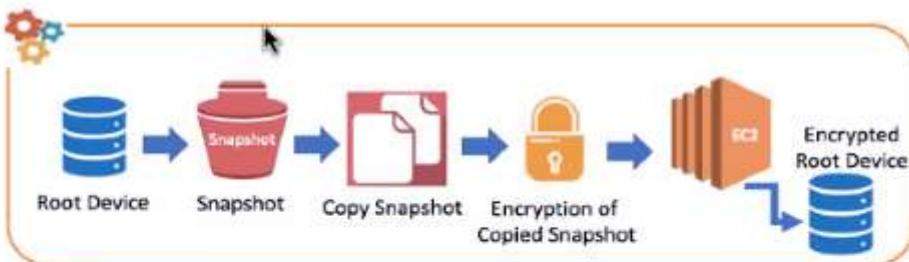


Domain: Design Secure Architectures

A security audit discovers that one of your RDS MySQL instances is not encrypted. The instance has a Read Replica in the same AWS region which is also not encrypted. You need to fix this issue as soon as possible. What is the proper way to add encryption to the instance and its replica?

- A. Create a DB snapshot from the instance. Copy the DB snapshot with encryption enabled. Restore a new DB instance from the new encrypted snapshot and configure a Read Replica in the new DB instance.
- B. Encrypt the DB instance. Launch a new Read Replica and the replica is encrypted automatically.
- C. Create a DB snapshot from the RDS instance and encrypt the newly-created snapshot. Launch a new instance and its Read Replica from the snapshot.
- D. Promote the Read Replica to be a standalone instance and encrypt it. Add a new Read Replica to the standalone instance.

Cevap A



A Media firm Firm_A uses AWS infrastructure and has a global presence for its sports programming & broadcasting network. It uses AWS Organization to manage multiple AWS accounts. Recently it was acquired by Firm_B which also uses AWS Infrastructure. Firm_B also has its own sets of AWS accounts.

After the merger, AWS Accounts of both organizations need to merge to create & manage policies more effectively.

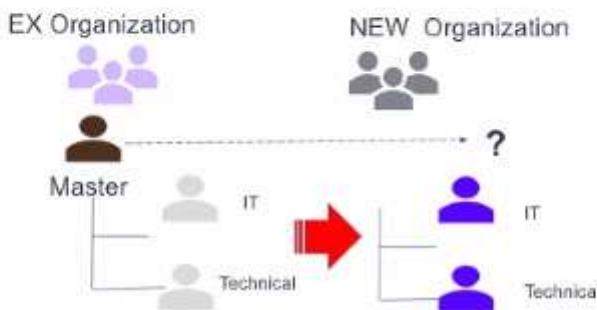
As an AWS Consultant, which of the following steps would you suggest to the client to move the management account of the Firm_A to the organization used by the merged entity? (Select THREE)

- A. Remove all member accounts from the organization in Firm_A.
- B. Configure another member account as the management account in the Firm_A organization.
- C. Delete the organization in Firm_A.
- D. Invite the Firm_A management account to join the new organization (Firm_B) as a member account.
- E. Invite the Firm_A management account to join the new organization (Firm_B) as a management account.

Cevap A,C ve D e olmaz çünkü 1 gemide 2 kaptan olmaz

MEMBER AND MASTER ACCOUNT LEAVING PROCESS

1. Remove the **member account** from the old Organization.
2. Send an invite to the **member account** from the new Organization.
3. Accept the invite to the new Organization from the **member account**.
4. Delete the old Organization.
5. Send an invite to the **master account**
6. Accept the invite to the new Organization from the **master account**



Question 53 of 65

Exit Quiz

Domain: Design Secure Architectures

Your company has designed an app and requires it to store data in DynamoDB. The company has registered the app with identity providers for users to sign-in using third-parties like Google and Facebook. What must be in place such that the app can obtain temporary credentials to access DynamoDB?

- A. Multi-factor authentication must be used to access DynamoDB.
- B. AWS CloudTrail needs to be enabled to audit usage.
- C. An IAM role allowing the app to have access to DynamoDB.
- D. The user must additionally log into the AWS console to gain database access.

Cevap C

Role Credentials

Once an entity assumes a role, it receives **temporary credentials** in the form of an **access key**, **secret key** and **session token**

Question 54 of 65

Exit Quiz

Domain: Design Resilient Architectures

A company has an entire infrastructure hosted on AWS. It requires to create code templates used to provide the same set of resources in another region in case of a disaster in the primary region. Which AWS service can be helpful in this regard?

- A. AWS Beanstalk
 - B. AWS CloudFormation
 - C. AWS CodeBuild
 - D. AWS CodeDeploy

Cevap B

Component	Description
Templates	The JSON or YAML text file that contains the instructions for building out the AWS environment.
Stacks	The entire environment described by the template and created, updated, and deleted as a single unit.
StackSets	AWS CloudFormation StackSets extends the functionality of stacks by enabling you to create, update, or delete stacks across multiple accounts and regions with a single operation.
Change Sets	A summary of proposed changes to your stack that will allow you to see how those changes might impact your existing resources before implementing them.

awsWorks Stacks **CloudFormation** **Elastic Beanstalk**

Question 55 of 65

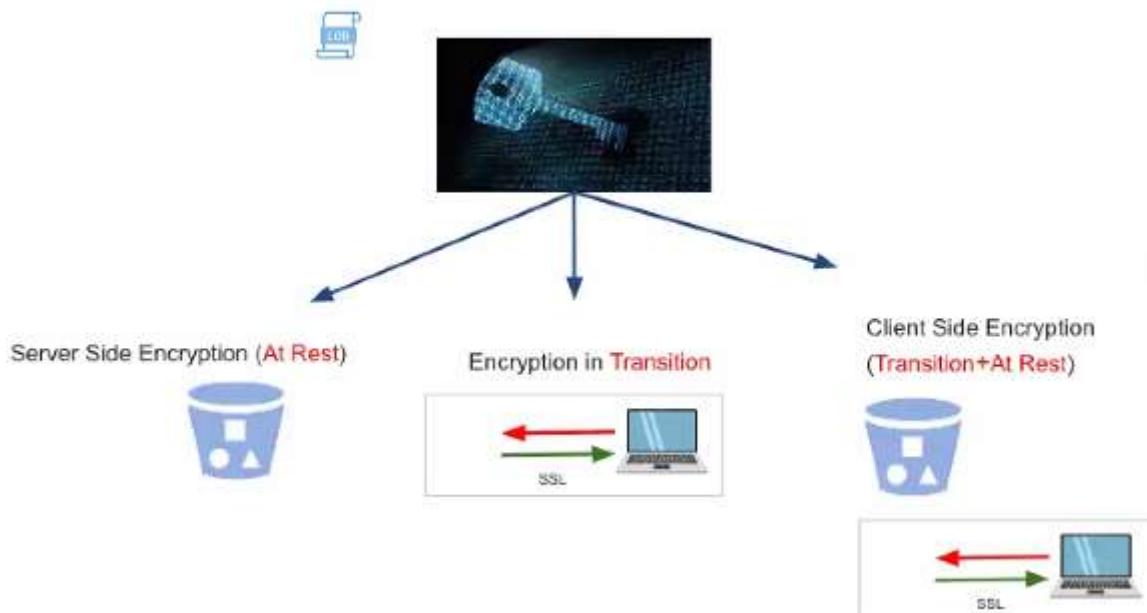
[Exit Quiz](#)

Domain: Design Secure Architectures

Your recent security reviews revealed a large spike in logins attempted to your AWS account. With respect to sensitive data stored in encryption enabled S3, the data has not been encrypted and is susceptible to fraud if it was to be stolen. You've recommended AWS Key Management Service as a solution. Which of the following is true regarding the operations of KMS?

- A. Only KMS generated keys can be used to encrypt or decrypt data.
- B. Data is encrypted at rest with KMS.
- C. KMS allows all users and roles to use the keys by default.
- D. Data is encrypted in transit with the KMS key.

Cevap B



Question 56 of 65

Exit Quiz

Domain: Design Resilient Architectures

Your company has a set of EC2 Instances hosted in AWS. It is required to prepare for regional disasters and come up with the necessary disaster recovery procedures. Which of the following steps would help to mitigate the effects of disaster in the future on EC2 Instances?

- A. Place an ELB in front of the EC2 Instances.
- B. Use Auto Scaling to ensure that the minimum number of instances are always running.
- C. Use CloudFront in front of the EC2 Instances.
- D. Create AMIs from the EC2 Instances. Use them to recreate the EC2 Instances in another region.

Cevap D



Question 57 of 65

[Exit Quiz](#)

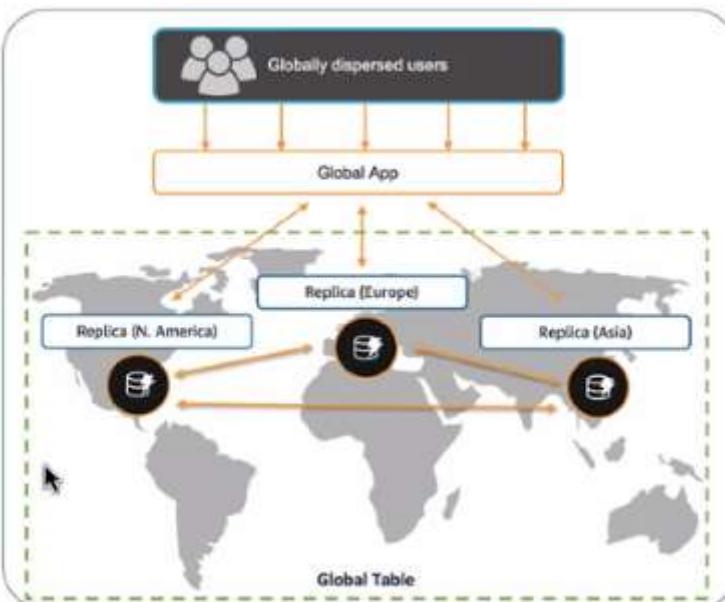
Domain: Design High-Performing Architectures

Your company currently has setup their data store on AWS DynamoDB. One of your main revenue generating applications uses the tables in this service. Your application is now expanding to 2 different other locations and you want to ensure that the latency for data retrieval is the least from the new regions.

Which of the following can help accomplish this?

- A. Place a CloudFront distribution in front of the database
- B. Enable Multi-AZ for DynamoDB
- C. Place an ElastiCache in front of DynamoDB
- D. Enable global tables for DynamoDB

Cevap D gecikmeleri önlemek için global table kullanılabilir



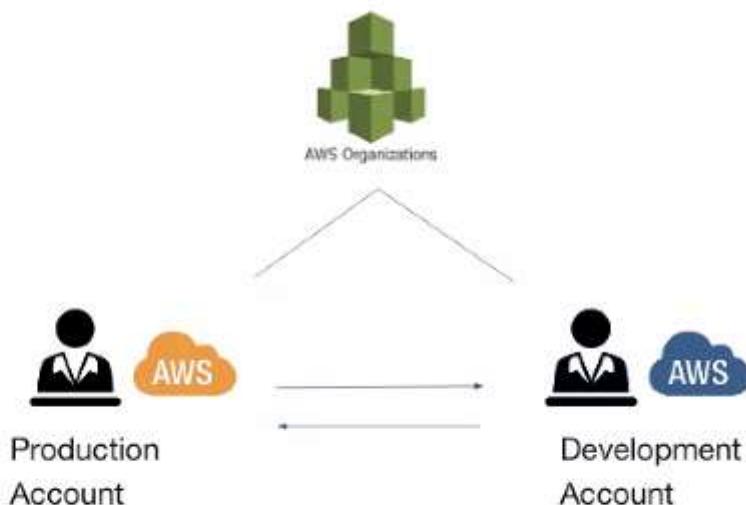
Domain: Design Resilient Architectures

A start-up firm is using an AWS Organization for managing policies across its Development and Production accounts. The development account needs an EC2 dedicated host. The Production account has subscribed to an EC2 dedicated host for its application but is not currently using it. Sharing has NOT been enabled with the AWS Organization in AWS RAM.

Which of the following can be done to share the Amazon EC2 dedicated host from the Production account to the Development account?

- A. Remove both Development & Production Accounts from Organization & then share resources between them.
- B. Resources in the same organization are automatically shared without the need to accept the invitation of sharing resources.
- C. Create a resource share in the production account and accept the invitation in the development account.
- D. Remove the destination Development account from an Organization & then share resources with it.

Cevap C



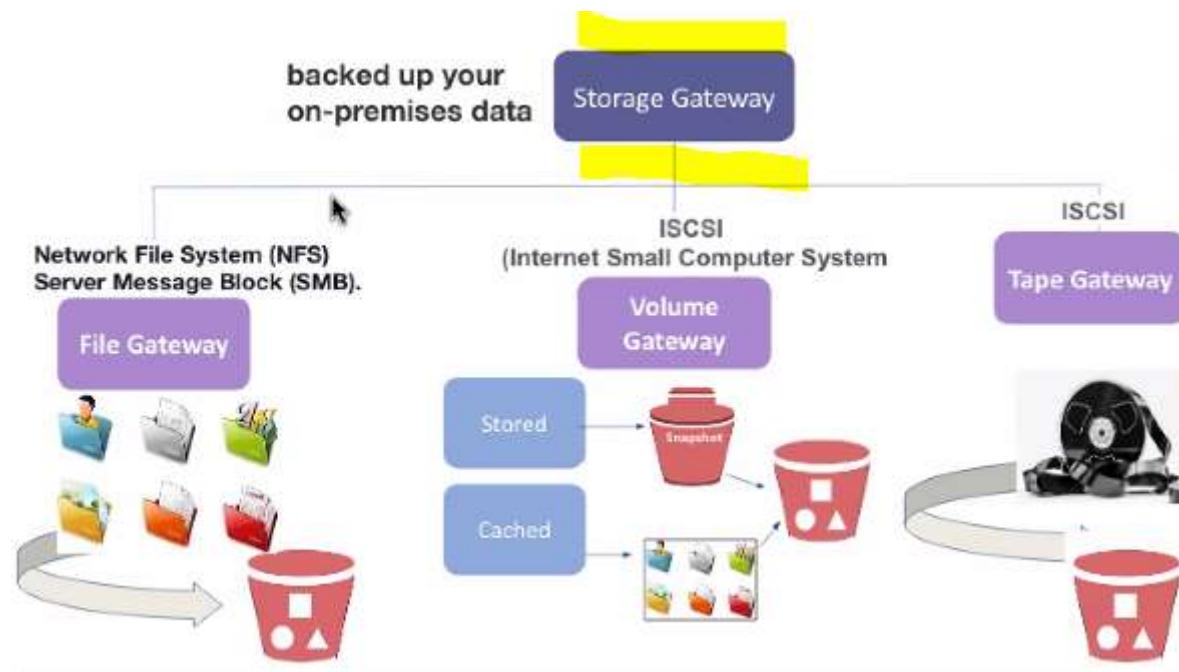
Enable resource sharing in AWS Organizations with RAM

Domain: Design High-Performing Architectures

A company has a lot of data hosted on their On-premises infrastructure. Running out of storage space, the company wants a quick win solution using AWS. There should be low latency for the frequently accessed data. Which of the following would allow the easy extension of their data infrastructure to AWS?

- A. The company could start using Gateway Cached Volumes.
- B. The company could start using Gateway Stored Volumes.
- C. The company could start using the Amazon S3 Glacier Deep Archive storage class.
- D. The company could start using Amazon S3 Glacier.

Cevap A on primedeki verileri aws de turma stogare gatewaydir.. cahced çok kullanılan



Domain: Design Resilient Architectures

A Large Medical Institute is using a legacy database for saving all its patient details. Due to compatibility issues with the latest software, they plan to migrate this database to AWS cloud infrastructure. This large size database will be using a NoSQL database Amazon DynamoDB in AWS. As an AWS consultant, you need to ensure that all the current legacy database tables are migrated without a glitch to Amazon DynamoDB. Which of the following is the most cost-effective way of transferring legacy databases to Amazon DynamoDB?

- A. Use AWS DMS with AWS Schema Conversion Tool to save data to Amazon S3 bucket & then upload all data to Amazon DynamoDB.
- B. Use AWS DMS with engine conversion tool to save data to Amazon S3 bucket & then upload all data to Amazon DynamoDB.
- C. Use AWS DMS with engine conversion tool to save data to Amazon EC2 & then upload all data to Amazon DynamoDB.
- D. Use AWS DMS with AWS Schema Conversion Tool to save data to Amazon EC2 instance & then upload all data to Amazon DynamoDB.

Cevap A

İki tip data migration vardır. Homogeneous (oracle dan oracela) and heterogenous(sqlden Aurora farklı db) farklı tip ise SCT kullanılır

Info: The term **legacy database** commonly refers to a database that has been in use for many years and is therefore unsuitable for modern apps and environments.

- DMS create a replication instance for migrating. Your databases are fully operational when migrating.
- **Homogeneous** database migrations (**Oracle to Oracle**) (**need Engine conversion**)
Heterogenous database migrations (**Microsoft SQL to Aurora**) (**Use SCT**)
- Use **SCT (Schema conversion tool)** from one schema to another schema for **Heterogenous migrations**

From **Legacy** to **DynamoDb** >>>>> **Heterogenous** database migrations



SCT (Schema conversion tool)

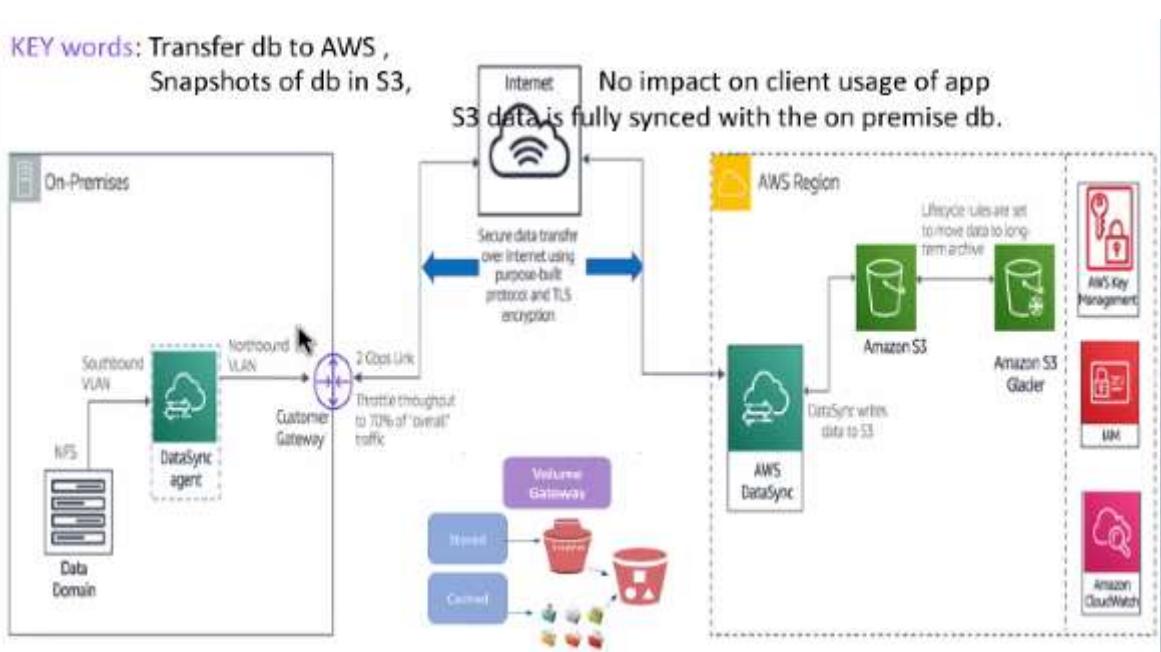
A Financial firm is planning to build a highly resilient application with primary database servers located at on-premises data centers while maintaining its DB snapshots in an S3 bucket. The IT Team is looking for a cost-effective and secure way of transferring the large customer financial databases from on-premises servers to the Amazon S3 bucket with no impact on the client usage of these applications. Also, post this data transfer, the on-premises application will be fetching data from the Amazon S3 bucket in case of a primary database failure.

So, your solution should ensure that the Amazon S3 data is fully synced with the on-premises database. Which of the following can be used to meet this requirement?

- A. Use Amazon S3 Transfer Acceleration for transferring data between the on-premises & Amazon S3 bucket while using AWS Data Sync for accessing these S3 bucket data from the on-premises application.
- B. Use AWS Data Sync for transferring data between the on-premises & Amazon S3 bucket while using AWS Storage Gateway for accessing these S3 bucket data from the on-premises application.
- C. Use AWS Snowball Edge for transferring data between the on-premises & Amazon S3 bucket while using AWS Storage Gateway for accessing these S3 bucket data from the on-premises application.
- D. Use AWS Transfer for transferring data between the on-premises & Amazon S3 bucket while using AWS Data Sync for accessing these S3 bucket data from the on-premises application.

CevaP B

KEY words: Transfer db to AWS ,
S snapshots of db in S3,



Question 62 of 65

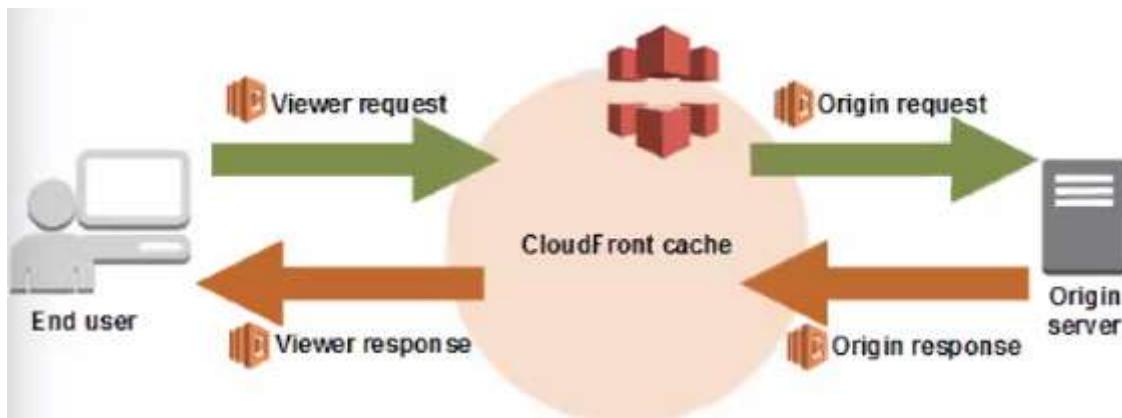
Exit Quiz

Domain: Design Cost-Optimized Architectures

A company has an application that delivers objects from S3 to global users. Of late, some users have been complaining of slow response times. Which additional step would help to build a cost-effective solution and ensure that the users get an optimal response to objects from S3?

- A. Use S3 Replication to replicate the objects to regions closest to the users.
- B. Ensure S3 Transfer Acceleration is enabled to ensure that all users get the desired response times.
- C. Place an ELB in front of S3 to distribute the load across S3.
- D. Place the S3 bucket behind a CloudFront distribution.

Cevap D



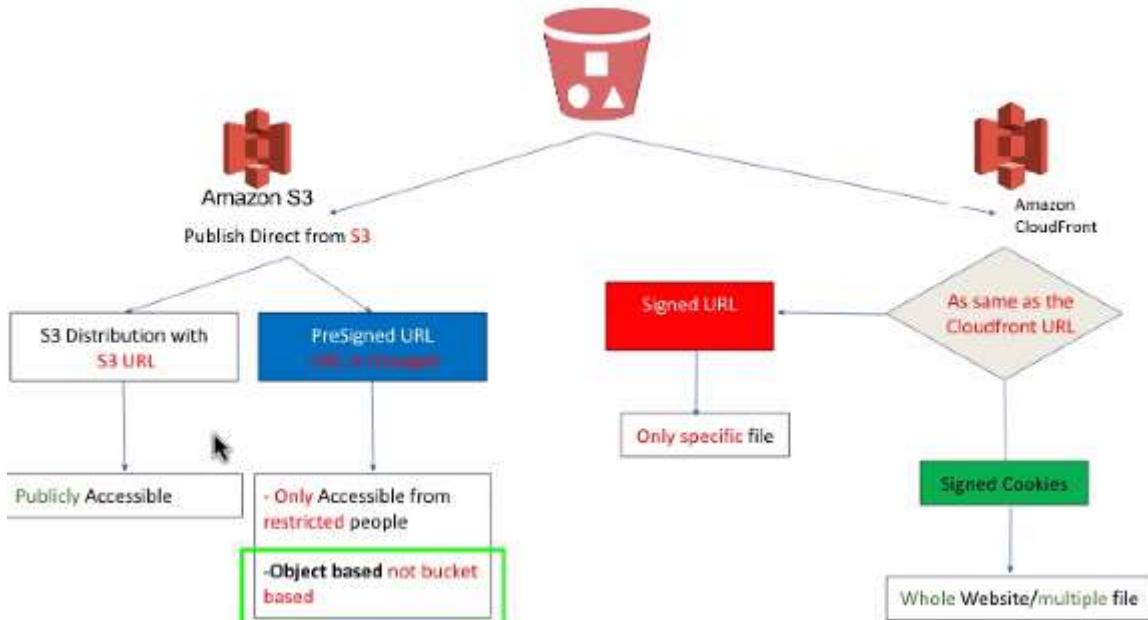
Domain: Design Secure Architectures

A large IT company is using Amazon CloudFront for its web application. Static Content for this application is saved in the Amazon S3 bucket. Amazon CloudFront is configured for this application to provide faster access to these files for global users.

IT Team is concerned about some critical files that need to be accessed only by users from certain white-list countries that you have defined in Amazon CloudFront geo-restriction. There is a requirement that no users should access these files directly using the Amazon S3 URL. Which of the following is the best way to achieve the given requirement?

- A. Create an OAI user to associate with distribution & modify permission on Amazon S3 bucket using bucket policy.
- B. Create Amazon CloudFront Signed URLs to limit access to these files & modify permission on Amazon S3 bucket using bucket policy.
- C. Create an OAI user to associate with distribution & modify permission on Amazon S3 bucket using object ACL's.
- D. Create Amazon CloudFront Signed URLs to limit access to these files & modify permission on Amazon S3 bucket using object ACL's.

Cevap A



Question 64 of 65

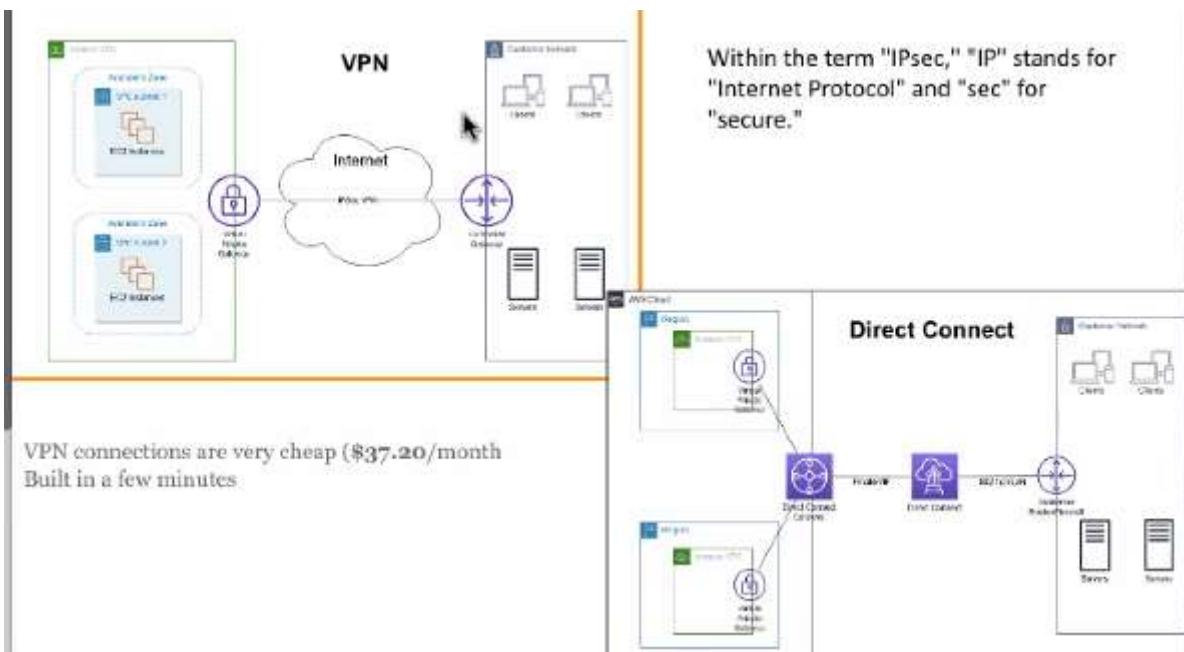
[Exit Quiz](#)

Domain: Design Secure Architectures

A start-up firm has a corporate office in New York & a regional office in Washington & Chicago. These offices are interconnected over Internet links. Recently they have migrated a few application servers to EC2 instance launched in the AWS US-east-1 region. The Developer Team located at the corporate office requires secure access to these servers for initial testing & performance checks before go-live of the new application. Since the go-live date is approaching soon, the IT team is looking for quick connectivity to be established. As an AWS consultant, which link option will you suggest as a cost-effective & quick way to establish secure connectivity from on-premise to servers launched in AWS?

- A. Use AWS Direct Connect to establish IPSEC connectivity from On-premise to VGW.
- B. Install a third party software VPN appliance from AWS Marketplace in the EC2 instance to create a VPN connection to the on-premises network.
- C. Use Hardware VPN over AWS Direct Connect to establish IPSEC connectivity from On-premise to VGW.
- D. Use AWS Site-to-Site VPN to establish IPSEC VPN connectivity between VPC and the on-premises network.

Cevap D



Question 65 of 65

[Exit Quiz](#)

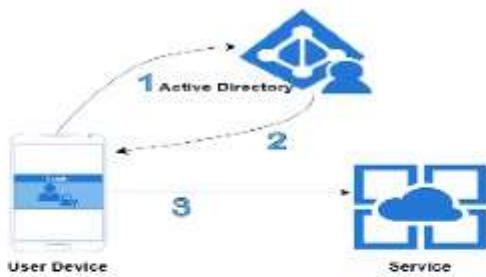
Domain: Design Secure Architectures

You are a Solutions Architect in a startup company that is releasing the first iteration of its app. Your company doesn't have a directory service for its intended users but wants the users to sign in and use the app. Which of the following solutions is the most cost-efficient?

- A. Create an IAM role for each end user and the user will assume the IAM role when he signs in the APP.
- B. Create an AWS user account for each customer.
- C. Invest heavily in Microsoft Active Directory as it's the industry standard.
- D. Use Cognito Identity along with a User Pool to securely save users' profile attributes.

Cevap D

Basic Active Directory Authentication



1
2
3

AWS Directory Service for Microsoft Active Directory

Simple AD

AD Connector

Amazon Cognito





CSAA 03 Practice Test 2

Question 1 of 65

Exit Quiz

Domain: Design High-Performing Architectures

A global beverage company is using AWS cloud infrastructure for hosting its web application. For a new beverage, the company plans to use a unique voice using Amazon Polly to help market this product to a wide range of customers.

Which feature can be used with Amazon Polly for this purpose?

- A. Create custom Lexicons with Amazon Polly
- B. Use a custom SSML tag with Amazon Polly
- C. Build a Brand Voice using Amazon Polly
- D. Use a Newscaster Speaking Style with Amazon Polly

Cevap C

Amazon Polly

You can use Custom Lexicons to customize pronunciation of specific words

You can use Speech Synthesis Markup Language (SSML) and SSML tag
For example <break>, <emphasis>



You can select Newscaster Speaking Style

Brand Voice



Amazon Polly is a service that turns **text into lifelike speech**.



Text metni verip çeşitli dillerde seslendirmesi yapılıyor

Domain: Design High-Performing Architectures

An airline company needs to analyze customer feedback about services provided. A large number of files are created which have texts such as Ticket returns and Flight complaints based upon customer feedback. This text should be categorized, and a label needs to be assigned for further analysis. The company is planning to use Amazon Comprehend for this purpose. The project team is looking for your suggestions for storing files that Amazon Comprehend can use for processing and saving the results.

How can a solution be implemented for this purpose?

- A. Store the files in Amazon EFS volumes. Use Amazon Comprehend to read the data from the Amazon EFS volumes and write results to Amazon Redshift
- B. Store the files in Amazon EBS volumes. Use Amazon Comprehend to read the data from the Amazon EBS volumes and write results to Amazon S3 buckets
- C. Store the files in Amazon Redshift. Use Amazon Comprehend to read the data from the Amazon Redshift and write results to Amazon S3 buckets
- D. Store the files in Amazon S3 bucket. Use Amazon Comprehend to read the data from the Amazon S3 bucket and write results to Amazon Redshift

Cevap D Burada Comprehend verileri S3 te tutar Anketleri yapay zeka ile analiz ediyor.

Explanation:

Correct Answer: D

Amazon Comprehend is a fully managed NLP (natural language processing) service that uses machine learning to analyze, understand and interpret meanings from a text document. Amazon Comprehend can be commonly used for the following cases,

1. Voice of Customer Analytics
2. Semantic Search
3. Knowledge management and discovery

The Amazon Comprehend can read data from the Amazon S3 bucket and write results to any AWS storage, database, or data warehousing resources.

For the above case, the company can store all the files in an Amazon S3 bucket. The Amazon Comprehend will read the data from these files and process the text files to extract the keywords such as Ticket returns and Flight complaints. It will store the output in Amazon Redshift which can be



Amazon Comprehend

Example: In this example, a restaurant is reviewing a customer review to understand where they can improve their business.

Sample Text: I loved the burger, but the service was slow.

Text	Entity Type	Entity Confidence Score	Sentiment	Score
I	Person	0.99	Neutral	0.99
Burger	Other	0.99	Positive	0.99
Service	Attribute	0.99	Negative	0.99

Müşteri verilerini alıp analiz ediyor.

Question 3 of 65

[Exit Quiz](#)

Domain: Design High-Performing Architectures

A stock broking company has deployed a stock trading web application on the Amazon EC2 instance. The company is looking for virtual agents to be integrated with this application to provide conversational channels to its premium customers. Real-time personalized stock recommendations should be provided for premium customers during market hours.

Which service is best suited to integrate with this application?

- A. Amazon Lex
- B. Amazon Translate
- C. Amazon Transcribe
- D. Amazon Personalize

Cevap A



Amazon Lex

Chatbots with conversational AI

Size nasıl yardımcı olabilirim şeklinde çıkan robot Yapay zeka mantığıyla çalışır.

A Company has provisioned a website in the US West (N. California) region using Amazon EC2 instance and Amazon CloudFront. The company is using IAM to control access to these resources. The client plans to use a third-party SSL/TLS certificate to support HTTPS. As an AWS consultant, you have been engaged to advise importing the certificate and using it along with Amazon CloudFront. Certificates should be easily imported, and you should monitor the expiration dates of imported certificates.

How can certificates be provisioned to meet this requirement?

- A. Import the third-party certificate in the US East (N. Virginia) region of AWS Certificate Manager and use it for Amazon CloudFront in the US West (N. California) region
- B. Import the third-party certificate in the US West (N. California) region of AWS Certificate Manager and use it for Amazon CloudFront in the US West (N. California) region
- C. Import the third-party certificate in the US East (N. Virginia) region of IAM Certificate Manager and use it for Amazon CloudFront in the US West (N. California) region
- D. Import the third-party certificate in the US West (N. California) region of IAM Certificate Manager and use it for Amazon CloudFront in the US West (N. California) region

Cevap A dışarıdan bir ssl sertifikası yüklenmesi N. Virginian'dan olur.

Import Custom SSL Certificate To AWS Certificate Manager



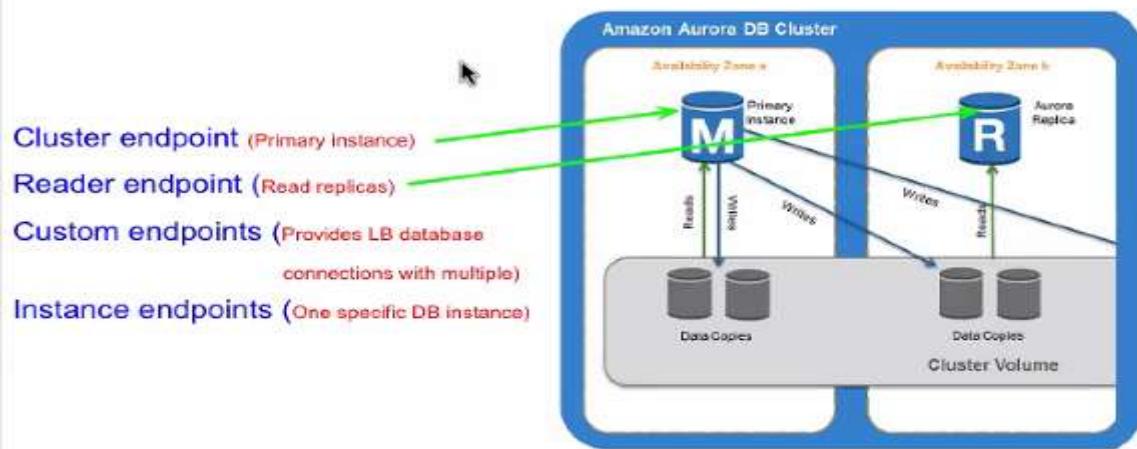
Domain: Design Resilient Architectures

The Developer Team has deployed a new application using Amazon Aurora DB cluster. This cluster has a primary instance and five Aurora Replicas. While testing this application, it was observed that the primary instance is getting overutilized with the read requests. Before going to production, the Project Manager wants you to analyze the issue and suggest changes.

Which design changes can a solution architect propose to the application?

- A. Point application to the custom endpoint of the Amazon Aurora
- B. Point application to cluster endpoint of the Amazon Aurora
- C. Point application to reader endpoint of the Amazon Aurora
- D. Point application to instance endpoint of the Amazon Aurora

Cevap C 5 tane read replica yapılabiliyor. Datanın tutulduğu volümeler sanki birbiri ile sync bağlanmış gibidir.



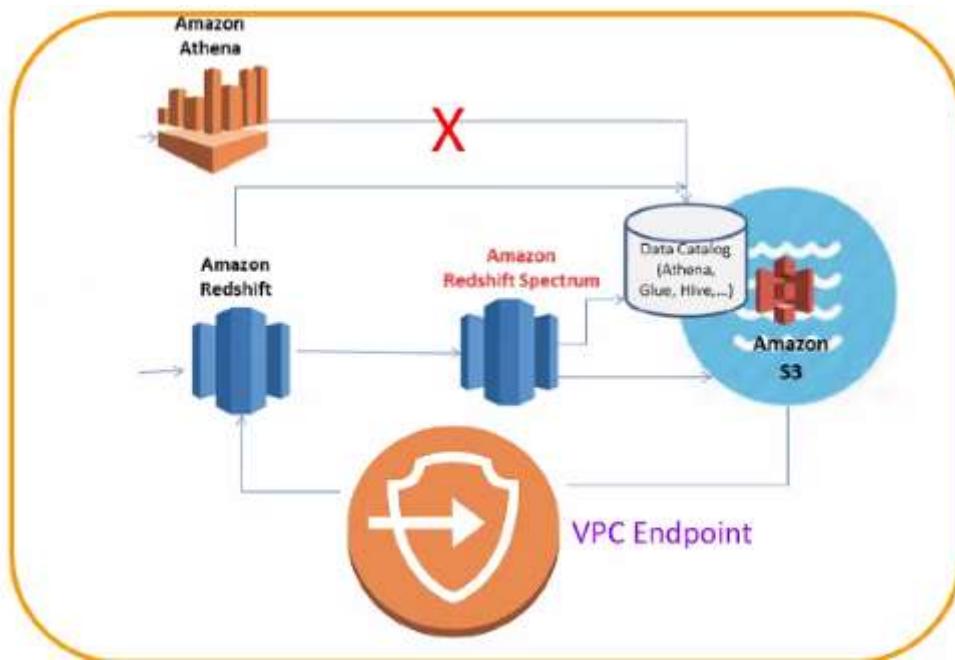
Domain: Design Secure Architectures

A gaming company stores large size (terabytes to petabytes) of clickstream events data into their central S3 bucket. The company wants to analyze this clickstream data to generate business insight. Amazon Redshift, hosted securely in a private subnet of a VPC, is used for all data warehouse-related and analytical solutions. Using Amazon Redshift, the company wants to explore some solutions to securely run complex analytical queries on the clickstream data stored in S3 without transforming/copying or loading the data in the Redshift.

As a Solutions Architect, which of the following AWS services would you recommend for this requirement, knowing that security and cost are two major priorities for the company?

- A. Create a VPC endpoint to establish a secure connection between Amazon Redshift and the S3 central bucket and use Amazon Athena to run the query
- B. Use NAT Gateway to connect Amazon Redshift to the internet and access the S3 static website. Use Amazon Redshift Spectrum to run the query
- C. Create a VPC endpoint to establish a secure connection between Amazon Redshift and the S3 central bucket and use Amazon Redshift Spectrum to run the query
- D. Create Site-to-Site VPN to set up a secure connection between Amazon Redshift and the S3 central bucket and use Amazon Redshift Spectrum to run the query

Cevap C



Domain: Design Secure Architectures

The drug research team in a Pharmaceutical company produces highly sensitive data and stores them in Amazon S3. The team wants to ensure top-notch security for their data while it is stored in Amazon S3. To have better control of the security, the team wants to use their own encryption key but doesn't want to maintain any code to perform data encryption and decryption. Also, the team wants to be responsible for storing the Secret key.

As a Solutions Architect, which of the following encryption types will suit the above requirement?

- A. Server-side encryption with customer-provided encryption keys (SSE-C)
- B. Server-Side Encryption with Amazon S3-Managed Keys (SSE-S3)
- C. Server-Side Encryption with KMS keys Stored in AWS Key Management Service (SSE-KMS)
- D. Protect the data using Client-Side Encryption

Cevap A



Server-Side Encryption with Amazon S3-Managed Keys (SSE-S3)



Server-Side Encryption with KMS keys
Stored in AWS Key Management Service (SSE-KMS)



Same as SSE-S3 but additionally you can track the key when and who use etc. Also it may charge.

Server-Side Encryption with Customer-Provided Keys (SSE-C)



You provide the key AWS provide the encryption

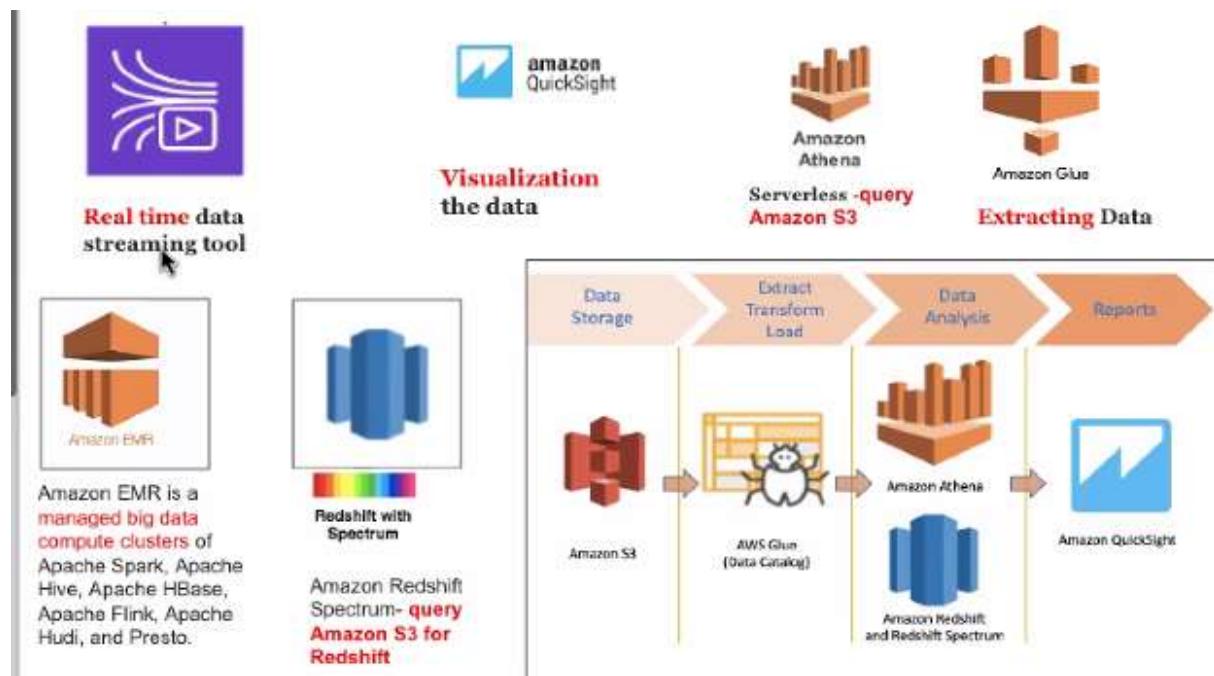
Domain: Design Cost-Optimized Architectures

An online retail company stores a large number of customer data (terabytes to petabytes) into Amazon S3. The company wants to drive some business insight out of this data. They plan to securely run SQL-based complex analytical queries on the S3 data directly and process it to generate business insights and build a data visualization dashboard for the business and management review and decision-making.

You are hired as a Solutions Architect to provide a cost-effective and quick solution to this. Which of the following AWS services would you recommend?

- A. Use Amazon Redshift Spectrum to run SQL-based queries on the data stored in Amazon S3 and then process it to Amazon Kinesis Data Analytics for creating a dashboard
- B. Use Amazon Redshift to run SQL-based queries on the data stored in Amazon S3 and then process it on a custom web-based dashboard for data visualization
- C. Use Amazon EMR to run SQL-based queries on the data stored in Amazon S3 and then process it to Amazon Quicksight for data visualization
- D. Use Amazon Athena to run SQL-based queries on the data stored in Amazon S3 and then process it to Amazon Quicksight for dashboard view

Cevap D



Question 9 of 65

Exit Quiz

Domain: Design Cost-Optimized Architectures

An organization has archived all their data to Amazon S3 Glacier for a long term. However, the organization needs to retrieve some portion of the archived data regularly. This retrieval process is quite random and incurs a good amount of cost for the organization. As expense is the top priority, the organization wants to set a data retrieval policy to avoid any data retrieval charges.

Which one of the following retrieval policies suits this in the best way?

- A. No Retrieval Limit
- B. Free Tier Only
- C. Max Retrieval Rate
- D. Standard Retrieval

Cevap B

Glacier Retrieve Policies**Free Tier Only Policy****Max Retrieval Rate Policy****No Retrieval Limit Policy****Glacier Retrieve Types****Expedited**

1–5 minutes

Standard

3–5 hours

Bulk

5–12 hours.

Domain: Design High-Performing Architectures

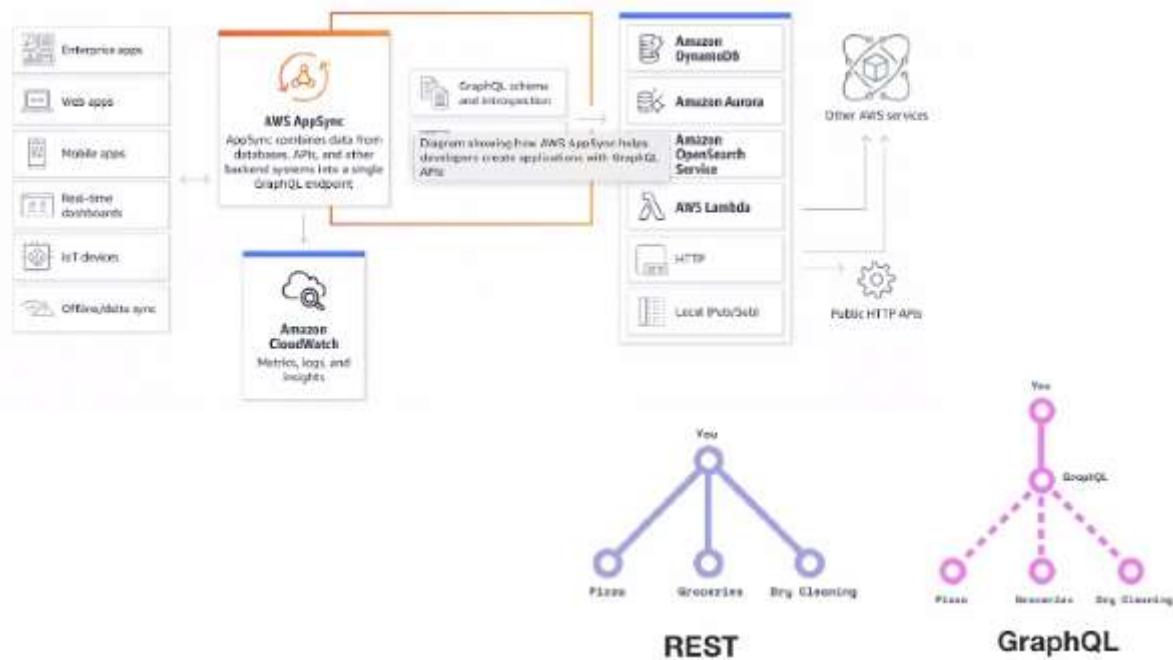
A gaming company planned to launch their new gaming application that will be in both web and mobile platforms. The company considers using GraphQL API to securely query or update data through a single endpoint from multiple databases, microservices, and several other API endpoints. They also want some portions of the data to be updated and accessed in real-time.

The customer prefers to build this new application mostly on serverless components of AWS.

As a Solutions Architect, which of the following AWS services would you recommend the customer to develop their GraphQL API?

- A. Kinesis Data Firehose
- B. Amazon Neptune
- C. Amazon API Gateway
- D. AWS AppSync

Cevap D



Question 11 of 65

[Exit Quiz](#)

Domain: Design High-Performing Architectures

A weather forecasting company comes up with the requirement of building a high-performance, highly parallel POSIX-compliant file system that stores data across multiple network file systems to serve thousands of simultaneous clients, driving millions of IOPS (Input/Output Operations per Second) with sub-millisecond latency. The company needs a cost-optimized file system storage for short-term, processing-heavy workloads that can provide burst throughput to meet this requirement.

What type of file systems storage will suit the company in the best way?

- A. FSx for Lustre with Deployment Type as Scratch File System
- B. FSx for Lustre with Deployment Type as Persistent file systems
- C. Amazon Elastic File System (Amazon EFS)
- D. Amazon FSx for Windows File Server

Cevap A



Amazon FSx
for Windows



Amazon FSx for
Lustre



- For windows Instance
- Can't write/read S3
- Used with Windows Active directory

- For Linux Instance
- Can write/read S3
- No Windows Active directory solution
- HPC!!!!!!

- For Linux Instance
- Can't write/read S3
- Increases and decreases the storage capacity automatically
- NFS

Deployment options for FSx for Lustre

Scratch File Systems are designed for temporary storage and shorter-term processing of data.

Persistent File Systems are designed for longer-term storage and workloads.

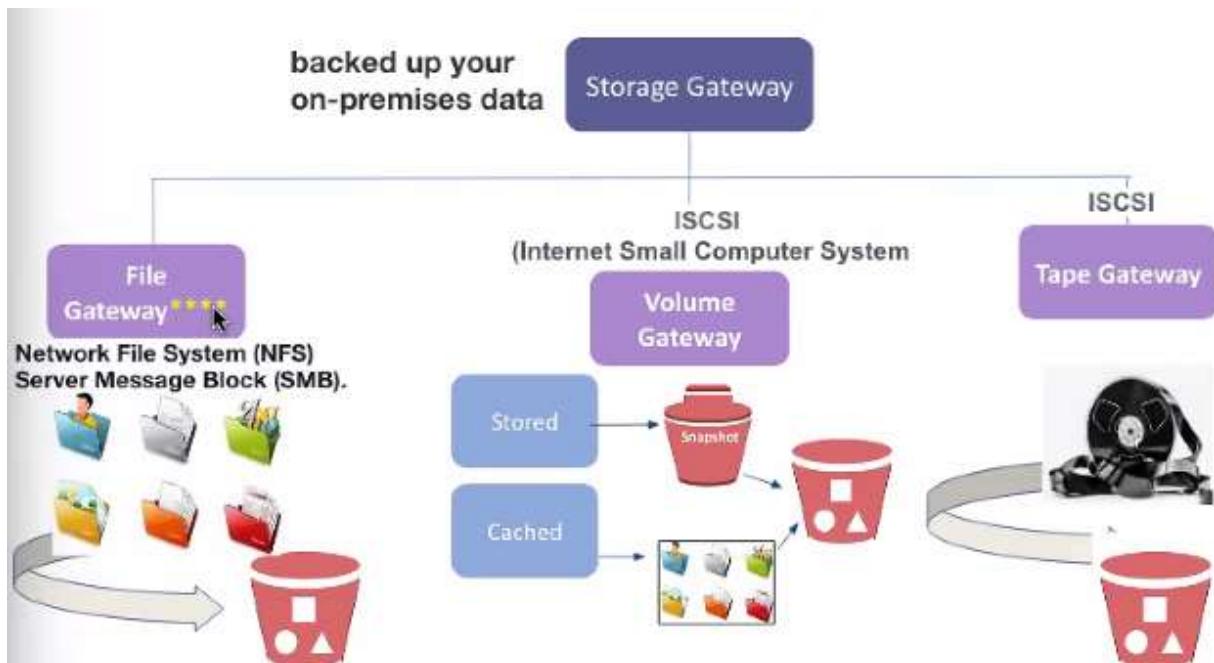
Domain: Design Resilient Architectures

A premier educational institute is stepping up to AWS Cloud as part of their 'Go Digital' initiative. The goal is to design a hybrid cloud storage service that gives them on-premises access to virtually unlimited cloud storage. Hence the institute plans to use AWS native service to back up their on-premises files using NFS and SMB file protocol. The institute wants to minimize latency in accessing the files in AWS by the students and users located across the globe. Hence the institute prefers to have multi-region replicas of the storage class so that the users can access it from AWS Regions that are geographically closer to them.

You are hired as a Solutions Architect to design this hybrid cloud facility. What will be your choice for the low-cost, secure and durable AWS Storage that supports multi-region data availability?

- A. Tape Gateway
- B. Amazon S3 File Gateway
- C. Amazon FSx File Gateway
- D. Volume Gateway

Cevap B



Question 13 of 65

[Exit Quiz](#)

Domain: Design Cost-Optimized Architectures

To comply with industry regulations, a Healthcare organization wants to keep their large volume of lab records in some durable, secure, lowest-cost AWS storage for a long period of time (say about five years). The data will be rarely accessed once per quarter but requires immediate retrieval (in milliseconds). You are a Solutions Architect in the organization and the organization wants your suggestion to select a suitable storage class here. Which of the following would you recommend for the given requirement?

- A. S3 Glacier Flexible Retrieval
- B. S3 Glacier Deep Archive
- C. S3 Glacier Instant Retrieval
- D. S3 Standard-Infrequent Access

Cevap C

Storage Classes

Archive Retrieval Options



Domain: Design Secure Architectures

An online retail company recently tied up with an external audit firm.

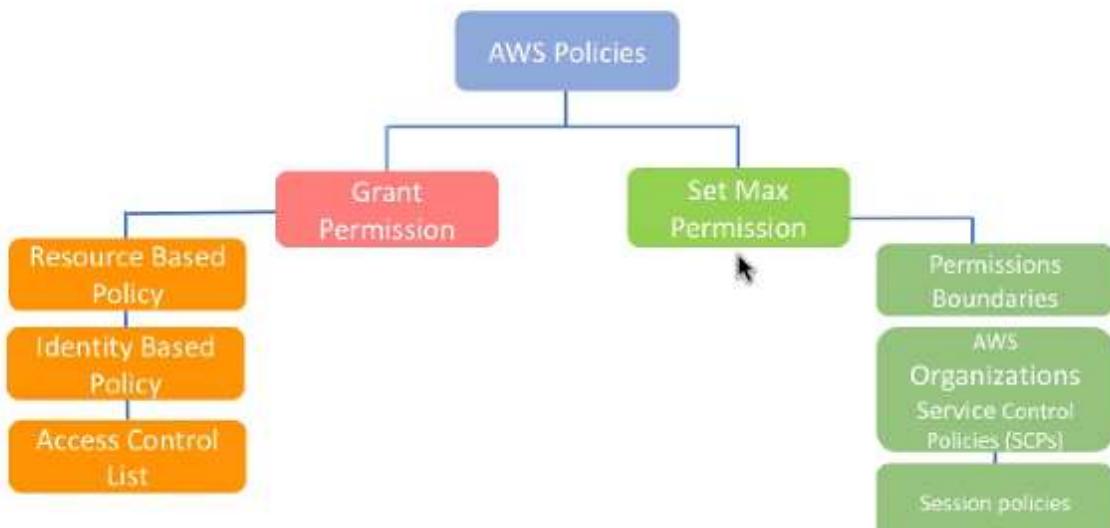
The retail company maintains all logs and saves them to an Amazon S3 bucket and the bucket access is restricted by Service Control Policy (SCP) for all other accounts of the company except the Admin account. The audit firm maintains all their audit related application and services in their own AWS account, and needs immediate access to retail company's S3 bucket to kick off their audit.

You are a solution architect and part of the admin team of the retail company, how should you enable the bucket access for the external audit team in the most secure manner, with no operational and management overload, and easy to deploy?

- A. Create a common IAM user in the retail company's AWS account and attach it to the IAM Role that has AmazonS3FullAccess. Share the IAM user details with the audit team
- B. Allow access to the audit firm in IAM Policy, Service Control Policy (SCP), and Amazon S3 Bucket Policy in the retail company's AWS account
- C. Allow S3 access to the audit firm in the retail company's SCP, just like the Admin account
- D. Add a bucket policy on the S3 bucket granting access to the aws account of the audit firm. Add HTTPS-only calls using "aws:SecureTransport" in the bucket policy

Cevap D

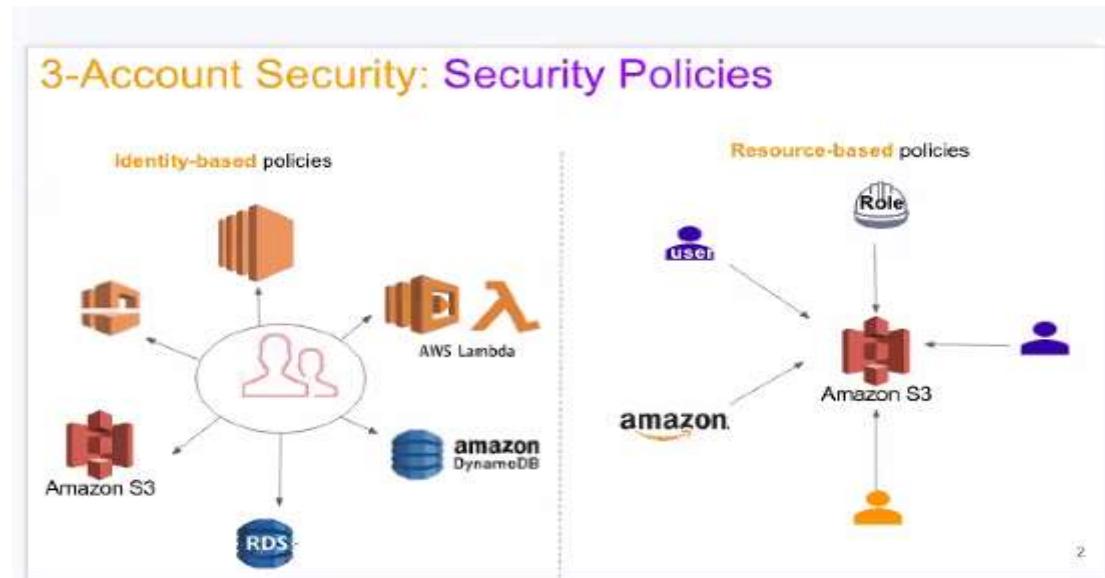
► Policy Types



3-Account Security: Security Policies



3-Account Security: Security Policies



Domain: Design Cost-Optimized Architectures

A customer is looking for file storage in AWS that supports Network File System version 4 (NFSv4.1 and NFSv4.0) protocol. They want a simple, serverless, set-and-forget service that can grow and shrink automatically as they add and remove files. Accessibility of these files is going to be random and infrequent. The customer is also looking for an option to save cost by transitioning the files that have not been accessed for quite some time automatically to a low-cost storage tier. As a Solution Architect, which storage service will you choose that fits the customer's requirement perfectly?

- A. Use Amazon S3 for storage as it is for better cost-saving option
- B. Use Amazon S3 Glacier for the storage
- C. Use Amazon Elastic File System (EFS) with lifecycle policy
- D. Amazon FSx for Lustre

Cevap C

General

Name - optional
Name your file system.
Optional. Apply a name to your file system.

Name can include letters, numbers, and -=_:/ symbols, up to 256 characters.

Storage class [Learn more](#)

Standard
Stores data redundantly across multiple AZs

One Zone
Stores data redundantly within a single AZ

Automatic backups

Automatically backup your file system data with AWS Backup using recommended settings. Additional pricing applies. [Learn more](#)

Enable automatic backups

Lifecycle management

EFS Intelligent-Tiering uses Lifecycle Management to automatically achieve the right price and performance blend for your application by moving your files between the Standard and Standard-Infrequent Access storage classes. [Learn more](#)

Transition into IA
Transition files from Standard to Standard-Infrequent Access.

30 day(s) since last access ▾

Transition out of IA
Transition files from Standard-Infrequent Access to Standard.

None ▾

Domain: Design Resilient Architectures

A drug research team in a Medical Company has decided to use Amazon Elastic File System (EFS) as shared file system storage for their Linux workloads. All these files are related to new drug discoveries in the field of cancer treatment and are critically important for the next six months. The customer is looking for a solution to protect the data by backing up the EFS file system and simplifying the creation, migration, restoration, and deletion of backups while providing improved reporting and auditing.

As a Solution Architect, what would be your suggestions for a centralized and easy-to-develop backup strategy for the above requirement?

- A. Use Amazon S3 File Gateway to back up the Amazon EFS file system
- B. Use AWS Backup to back up the Amazon EFS file systems
- C. Amazon FSx File Gateway to back up the Amazon EFS file systems
- D. Use Amazon S3 Transfer Acceleration to copy the files from EFS into a centralized S3 bucket and then configure Cross-Region Replication of the bucket

Cevap B

AWS Backup



Domain: Design Secure Architectures

The CIO of a Start-up company is very much concerned with the performance of their DevOps team as they take a long time to detect and investigate issues using AWS resources. Sometimes this leads to a revenue loss for the company. In addition to this, the same team had several compliance issues in a recent security audit as the audit team found secure information like application configurations, custom environment variables, product keys, credentials, and sensitive AMI IDs being mentioned directly in the code.

The company has hired you as a Solution Architect, and the CIO instructed you to resolve all the challenges faced by the DevOps team on priority and get them back to their efficiency.

He has also emphasized implementing a cost-effective solution for storing secure information that does not need automated secret rotation.

What would you do to achieve the requirement? (Select TWO)

- A. For quick failure analysis and investigation, use AWS Systems Manager Application Manager
- B. For quick failure analysis and investigation, use AWS AppSync
- C. Resolve the compliance issue by storing all the secrets in AWS Secrets Manager and changing the code to access the secrets from there
- D. Resolve the compliance issue by storing all the secrets in AWS Systems Manager Parameter Store and changing the code to access the secrets from there
- E. Resolve the compliance issue by storing all the secrets in a private Amazon S3 bucket. Then create a Gateway VPN Endpoint to access the secret from the bucket securely

Cevap A ve D

The screenshot shows the AWS Systems Manager console with the 'Application Management' section highlighted by a yellow box. Inside this box, the 'Application Manager' link is also highlighted with a green underline. Other options like 'AppConfig' and 'Parameter Store' are visible below it.

Application Manager, a capability of AWS Systems Manager, helps DevOps engineers investigate and remediate issues with their AWS resources in the context of their applications and clusters.

Parameter Store can store data such as passwords, database strings, Amazon Machine Image (AMI) IDs, and license codes as parameter values. You can store values as plain text or encrypted data.

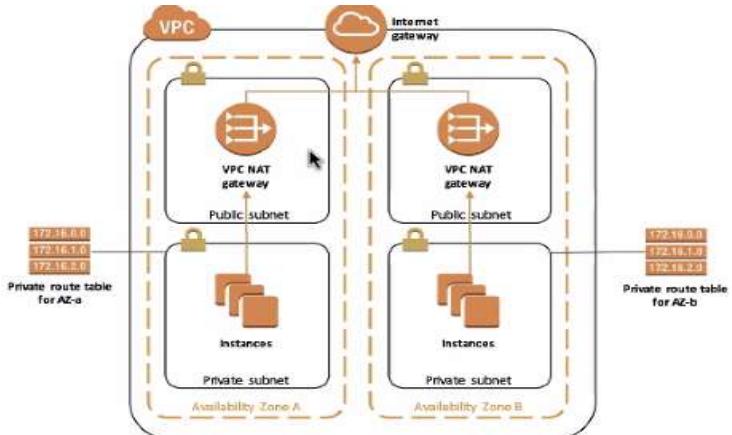
Secrets Manager enables you to replace hardcoded credentials in your code, including DB passwords, with an API call to Secrets Manager to retrieve the secret programmatically.

An organization in the banking sector has got their AWS resources distributed in multiple Availability Zones (AZ) in a region. They share one NAT gateway to connect to all their RDS instances placed in Private Subnet in different AZs to the internet. While trying to perform some Database related operations, the employee of the organization reported intermittent connectivity issues. However, on checking the logs in AWS, they found that the RDS instances were all up and running during the time when the issue was reported.

You are hired as a Solution Architect to identify the root cause of this connectivity issue and remediate it without compromising security and resiliency. How would you approach this?

- A. The customer experiences a connectivity issue when the NAT gateway's Availability Zone goes down. To remediate this, create a NAT gateway in each Availability Zone and configure your routing to ensure that resources use the NAT gateway in the same Availability Zone
- B. The customer experiences a connectivity issue due to insufficient IAM Policy and Roles. Resolve this by adding the required permissions in the IAM Policy and attach them to the role used for the AWS resources
- C. The customer experiences a connectivity issue due to low bandwidth in their network. Ask the customer to sign up for a higher bandwidth plan with the network provider/operator
- D. The customer experiences the connectivity issue when the NAT gateway's Availability Zone goes down. To remediate this, remove the NAT gateway and use VPC Peering to connect to the Multi-AZ RDS instances

Cevap A



Domain: Design Resilient Architectures

An online advertising company wants to build a solution in AWS that will understand the interest of the customer and provide relevant and personalized recommendations of new products to them. The company is looking for a self-managed, highly durable, and available Database engine that stores relationships between information such as customer's interests, community networks, friend circle, likes, purchase history, etc., and quickly maps them to identify the pattern.

As a Solution Architect in the company, what would be your recommendations for the database here?

- A. Amazon DocumentDB
- B. Amazon DynamoDB
- C. Amazon Aurora
- D. Amazon Neptune

Cevap D

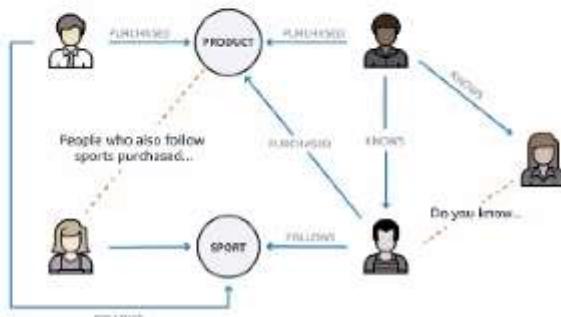


SQL

NoSQL

MongoDB-compatible Databases

Fully managed Graph Database



Question 20 of 65

[Exit Quiz](#)**Domain:** Design Resilient Architectures

In response to the high demand and increase in load, a customer plans to migrate his on-premises and native MongoDB to AWS Cloud. The customer is looking for a compatible Database solution in AWS for easy and fast migration with minimum operation and management overhead. The new database should also be compatible with existing MongoDB so that the applications don't require code changes.

As a Solution Architect in the company, what would be your suggestion for this scenario?

- A. Amazon DocumentDB
- B. Amazon DynamoDB
- C. Amazon Keyspaces
- D. Amazon Neptune

Cevap A



Domain: Design Secure Architectures

You are an engineer in charge of the FinOps department of your organization. The multi-account strategy has been created using AWS Control Tower as part of the best practices and recommendations. You want to ensure that all accounts under the 'OU=development' of the AWS Organizations should not create resources outside of Ireland (eu-west-1). How can you make sure that the required condition is applied using the below Policy Statement?

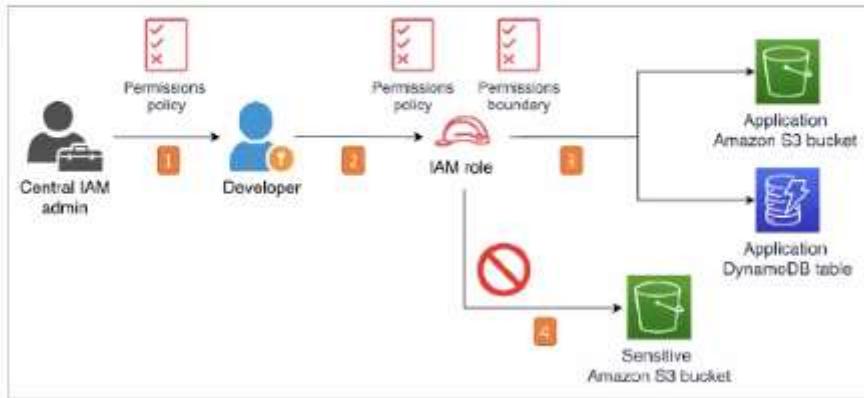
```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "IrelandDeny",
      "Effect": "Deny",
      "Action": [
        "*"
      ],
      "Resource": "*",
      "Condition": {
        "StringNotEquals": {
          "aws:RequestedRegion": [
            "eu-west-1"
          ]
        }
      }
    }
  ]
}
```



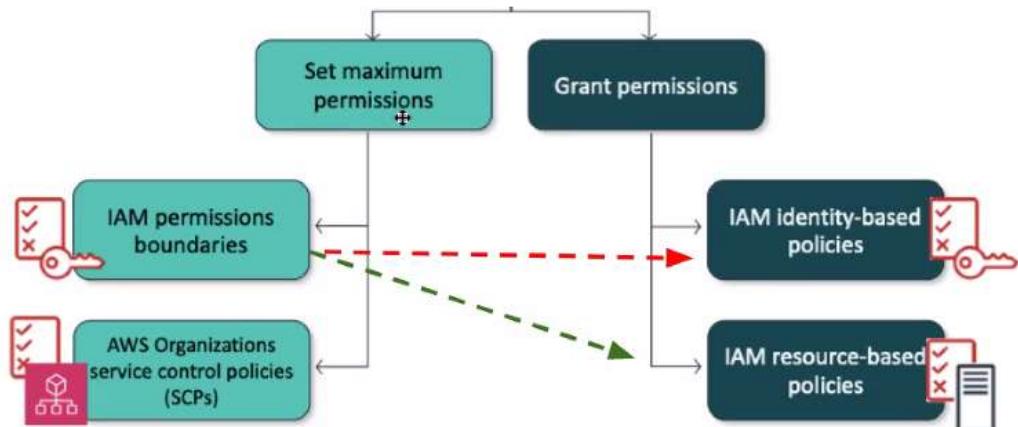
- A. Add the Policy Statement to a Permissions Boundary to the developer IAM role
- B. Add the Policy Statement to the SCP (Service control policy) and attach it to the Organizational Unit OU=development
- C. Add the Policy Statement to the SCP (Service control policy) and attach it to the Organizational Unit OU=root
- D. Add the Policy Statement as a managed policy for the role

Cevap A

Permissions Boundaries.



3-Account Security: Security Policies



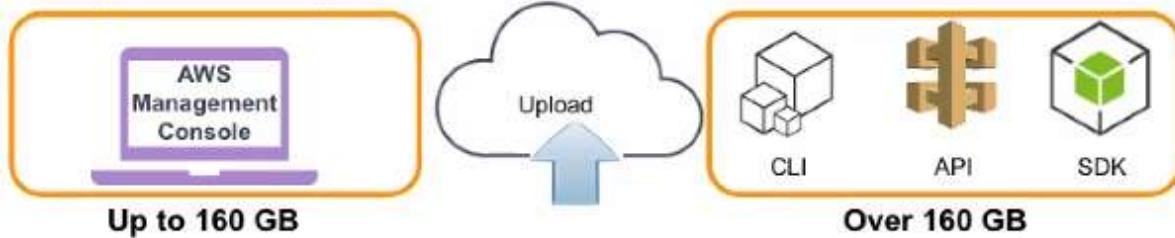
Domain: Design Resilient Architectures

You are working in a multimedia company and want to transfer a massive amount of data to Amazon S3. You have heard that SnowBall Edge could be the right tool for this purpose. When you are trying to transfer data using the AWS CLI, one of our biggest files of 12TB, you get an error with the client validating this transfer. What could be the main cause of the error?

- A. The maximum file size that could be transferred using the AWS CLI is 150GB
- B. The maximum file size that could be transferred using the AWS CLI is 5TB
- C. The role that you are using to transfer files can't transfer more than 5TB
- D. Amazon S3 doesn't support files larger than 10TB

Cevap B

Object in S3-Upload

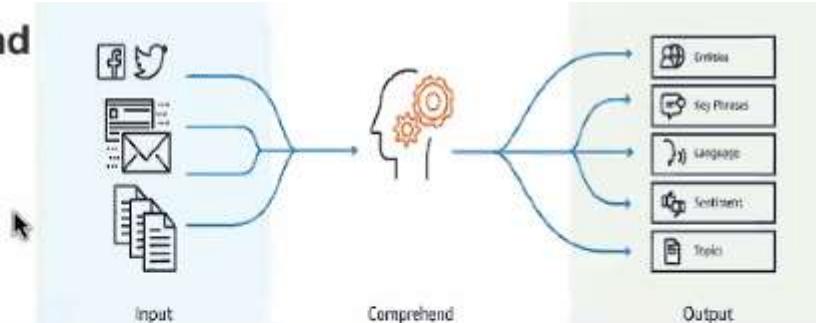


- The max. size of an object you can upload via **AWS Management Console** is 160 GB.
- For uploading a file **greater than 160 GB**, the AWS CLI, AWS SDK, or API is needed to be used
- The max size of a **single file** can be put in a bucket is 5 TB.

Domain: Design Secure Architectures

You are a Research Scientist working on NLP (Natural Language Processing) models. You are planning to use Amazon Comprehend to do some research about multiple texts that are currently stored in your email, but you have some ownership concerns about this service. Who would retain the ownership of the content of the emails that you will analyze using Amazon Comprehend?

- A. Customer
- B. AWS
- C. A third-party company that is in charge of managing the Amazon Comprehend service
- D. Both, AWS and the Customer

Cevap A**Amazon Comprehend****Q: Who has access to my content that is processed and stored by Amazon Comprehend?**

Only authorized employees will have access to your content that is processed by Amazon Comprehend. Your trust, privacy, and the security of your content are our highest priority and we implement appropriate and sophisticated technical and physical controls, including encryption at rest and in transit, designed to prevent unauthorized access to, or disclosure of, your content and ensure that our use complies with our commitments to you. Please see the AWS [data privacy FAQs](#) for more information.

[Show less](#)**Q: Do I still own my content that is processed and stored by Amazon Comprehend?**

You always retain ownership of your content and we will only use your content with your consent.

[Show less](#)

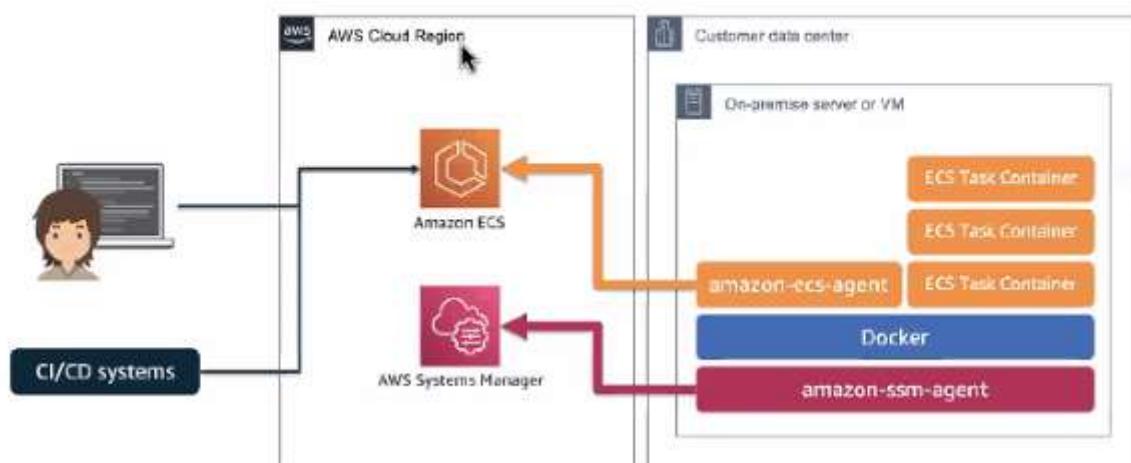
Domain: Design Resilient Architectures

You are a DevOps Engineer specializing in Containers. You want to run your container workloads running in a Linux environment in your current Data Center for regulations with easy management of your Docker tasks using the AWS Console. You have heard about Amazon ECS Anywhere, and you would like to explore it deeply. What are the main components you should have in your current Linux Servers in an ECS Cluster?

- A. AWS Systems Manager Agent, Amazon ECS container agent, and Docker must be installed on these Linux Servers
- B. AWS Systems Manager Agent Docker must be installed as part of the process
- C. Amazon ECS container agent and Docker must be installed on these external instances
- D. The Amazon CloudWatch Agent, Amazon ECS container agent, and Docker must be installed on these external instances

Cevap A

ECS Anywhere Overview



Domain: Design Resilient Architectures

You are a DevOps Engineer for a company specializing in Container technology. You are currently running your container workloads **nodes in your bare metal nodes** due to regulatory compliance rules. You would like to continue running your workloads on **hardware (bare metal)** without any **virtualization**, but want to use an AWS-managed service for managing multiple clusters of nodes. What is the best solution to integrate your container workloads from your Data Center using a **Kubernetes ecosystem?**

- A. Use Amazon ECS Anywhere
- B. Move all your workloads from your current Data Center to Amazon EKS
- C. Deploy an Amazon EKS Anywhere technology on bare metal nodes
- D. Create a VMware vSphere cluster and integrate it with Amazon EKS Anywhere

Cevap C

EKS Anywhere on Bare Metal

NO NEED FOR VIRTUALIZATION



Question 26 of 65

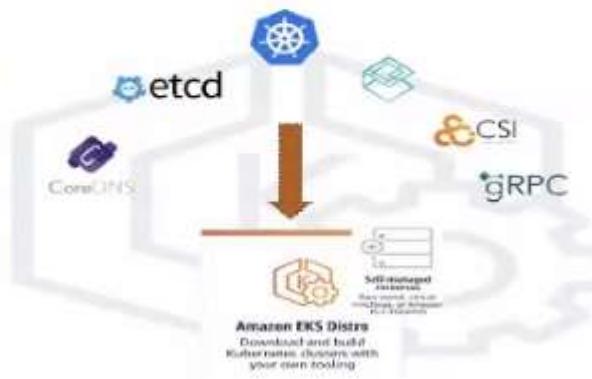
Exit Quiz

Domain: Design High-Performing Architectures

Which of the following components are included with Amazon EKS Distro?

- A. Kubelet, CNI plugins, CoreDNS, etcd, CSI Sidecars, Kubernetes Metrics Server
- B. CNI plugins, CoreDNS, etcd, CSI Sidecars, Kubernetes Metrics Server, Amazon Systems Manager Agent
- C. FluxCD operator, Kubernetes, CNI plugins, CoreDNS ,etcd,CSI Sidecars
- D. Amazon CloudWatch agent, Kubernetes, CNI plugins, CoreDNS, etcd, CSI Sidecars, Kubernetes Metrics Server

Cevap A ????



Question 27 of 65

[Exit Quiz](#)

Domain: Design High-Performing Architectures

You are the owner of a Microservices application that has a poor latency when it runs into the ECS cluster. Which AWS services could help you analyze the root cause by tracing different calls into the application?

- A. Amazon CloudWatch
- B. AWS X-Ray
- C. Amazon Event Bridge
- D. Amazon CloudTrail

[Answer](#)

Cevap B



MySQL and PostgreSQL databases (self-hosted, Amazon RDS, Amazon Aurora),

Amazon DynamoDB.

Amazon Simple Queue Service and Amazon Simple Notification Service.

Question 28 of 65

[Exit Quiz](#)

Domain: Design Secure Architectures

You are a Security Consultant who wants to prevent the application from DDoS attacks. You have heard about AWS Shield and DDoS attack vectors. Which of the following classes of attacks can be detected using AWS Shield?

- A. Layer 7 (Data Link Layer attacks)
- B. Layer 5 (Session Layer attacks)
- C. Layer 4 (Network protocol attacks)
- D. Layer 3 (Network volumetric attacks)
- E. Layer 1 (Physical Layer attacks)

Cevap C ve D ??

AWS WAF

- Web Application Firewall (WAF) operates specifically at protocol **layer 7** and monitors **HTTP traffic**
- Typically protects web applications against specific attacks:
 - cross-site forgery**
 - cross-site-scripting (XSS)**
 - SQL injection**
 - Distributed-denial-of-service (DDOS)**

AWS Shield Advanced

AWS Shield Advanced provide protections against **Distributed Denial of Service (DDoS)** attacks for AWS resources at the

- Network layer (**layer 3**)
- Transport layers (**layer 4**)
- The application layer (**layer 7**).

Question 29 of 65

[Exit Quiz](#)

Domain: Design Secure Architectures

You are a Cloud Database Administrator and want to enable IAM authentication on your Aurora Databases for the user 'iam_db_user'. Which of the following is the correct method to enable the IAM authentication in your Aurora MySQL Database using AWSAuthenticationPlugin?

- A. CREATE USER iam_db_user IDENTIFIED WITH AWSAuthenticationPlugin;
- B. CREATE USER iam_db_user_plugin;
- C. Rename IAM iam_db_user as a Database user
- D. CREATE USER iam_db_user IDENTIFIED WITH AWSAuthenticationPlugin AS 'RDS';

Cevap D

Database authentication

Database authentication options: [Info](#)

Password authentication

Authenticates using database passwords.

Password and IAM database authentication (not available for Multi-AZ DB cluster)

Authenticates using the database password and user credentials through AWS IAM users and roles.

Password and Kerberos authentication (not available for Multi-AZ DB cluster)

Choose a directory in which you want to allow authorized users to authenticate with the DB instance using Kerberos Authentication.

Using IAM authentication with Aurora MySQL

With Aurora MySQL, authentication is handled by `AWSAuthenticationPlugin`—an AWS-provided plugin that works seamlessly with IAM to authenticate your IAM users. Connect to the DB cluster and issue the `CREATE USER` statement, as shown in the following example.

```
CREATE USER jane_doe IDENTIFIED WITH AWSAuthenticationPlugin AS 'RDS';
```

11 IC1

Question 30 of 65

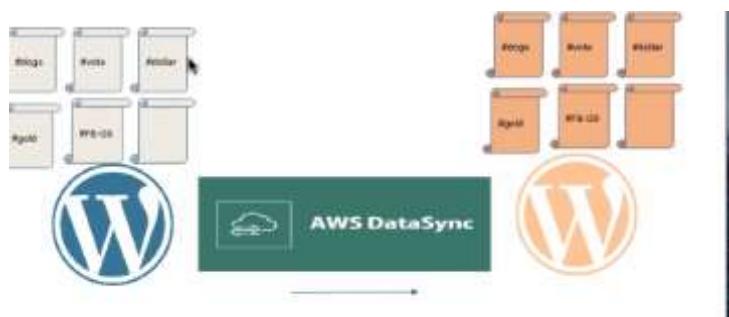
[Exit Quiz](#)

Domain: Design Resilient Architectures

You are using AWS DataSync to migrate more than 8TB from on-prem to Amazon S3. After the first DataSync task runs, you notice that some files were not copied. After reviewing the CloudWatch logs, you noticed that the files were skipped. What could be the main cause?

- A. The source file was locked and couldn't be opened by AWS DataSync
- B. The source file was opened and modified while it was transferred
- C. The source file's owner has been changed after it was transferred during the VERIFYING phase
- D. The source file's permissions are changed after it was transferred and couldn't be read during the VERIFYING phase

Cevap A ve B



Domain: Design Secure Architectures

Your company plans to use AppSync with Amazon DynamoDB as a data source. The operations you need to manage between AWS AppSync DynamoDB resolver and DynamoDB are BatchGetItem, BatchPutItem, and BatchDeleteItem. Which policy should you assign to AppSync to use DynamoDB as a source of data and apply the above operations?

A. {
 "Version": "2012-10-17",
 "Statement": [
 {
 "Action": [
 "dynamodb:BatchGetItem",
 "dynamodb:BatchWriteItem"
],
 "Effect": "Allow",
 "Resource": [
 "arn:aws:dynamodb:region:account:table/TABLENAME",
 "arn:aws:dynamodb:region:account:table/TABLENAME/*"
]
 }
]
}

- B. {
 "Version": "2012-10-17",
 "Statement": [
 {
 "Action": [
 "dynamodb:BatchGetItem",
 "dynamodb:BatchPutItem",
 "dynamodb:BatchDeleteItem"
],
 "Effect": "Allow",
 "Resource": [
 "arn:aws:dynamodb:region:account:table/TABLENAME",
 "arn:aws:dynamodb:region:account:table/TABLENAME/*"
]
 }
]
}
- C. {
 "Version": "2012-10-17",
 "Statement": [
 {
 "Action": [
 "dynamodb:BatchGetItem",
 "dynamodb:BatchPutItem",
 "dynamodb:BatchDeleteItem"
],
 "Effect": "Allow",
 "Resource": [
 "arn:aws:dynamodb:region:account:table/TABLENAME/*"
]
 }
]
}
- D. {
 "Version": "2012-10-17",
 "Statement": [
 {
 "Action": [
 "dynamodb:BatchGetItem",
 "dynamodb:BatchWriteItem"
],
 "Effect": "Allow",
 "Resource": [
 "arn:aws:dynamodb:region:account:table/TABLENAME"
]
 }
]
}

Cevap B

```
    "dynamodb:BatchDeleteItem"
  ],
  "Effect": "Allow",
  "Resource": [
    "arn:aws:dynamodb:region:account:table/TABLENAME",
    "arn:aws:dynamodb:region:account:table/TABLENAME/*"
  ]
}
}
```

* demek içindeki

bilgiye erişimdir. Bucketin kendisi için de ayrıca policy tanımlanır tabloya erişim.

Domain: Design Resilient Architectures

You are a solutions architect working for an online retailer. Your online website uses REST API calls via API Gateway and Lambda from your Angular SPA front-end to interact with your DynamoDB data store. Your DynamoDB tables are used for customer preferences, account, and product information. When your web traffic spikes, some requests return a 429 error response. What might be the reason your requests are returning a 429 response?

- A. Your Lambda function has exceeded the concurrency limit
- B. DynamoDB concurrency limit has been exceeded
- C. Your Angular service failed to connect to your API Gateway REST endpoint
- D. Your Angular service cannot handle the volume spike
- E. Your API Gateway has exceeded the steady-state request rate and burst limits

Cevap A ve E



When request submissions exceed the steady-state request rate and burst limits, API Gateway begins to throttle requests. Clients may receive 429 Too Many Requests error responses at this point.

GET



API



DynamoDB

Concurrency is the number of requests that your function is serving at any given time.

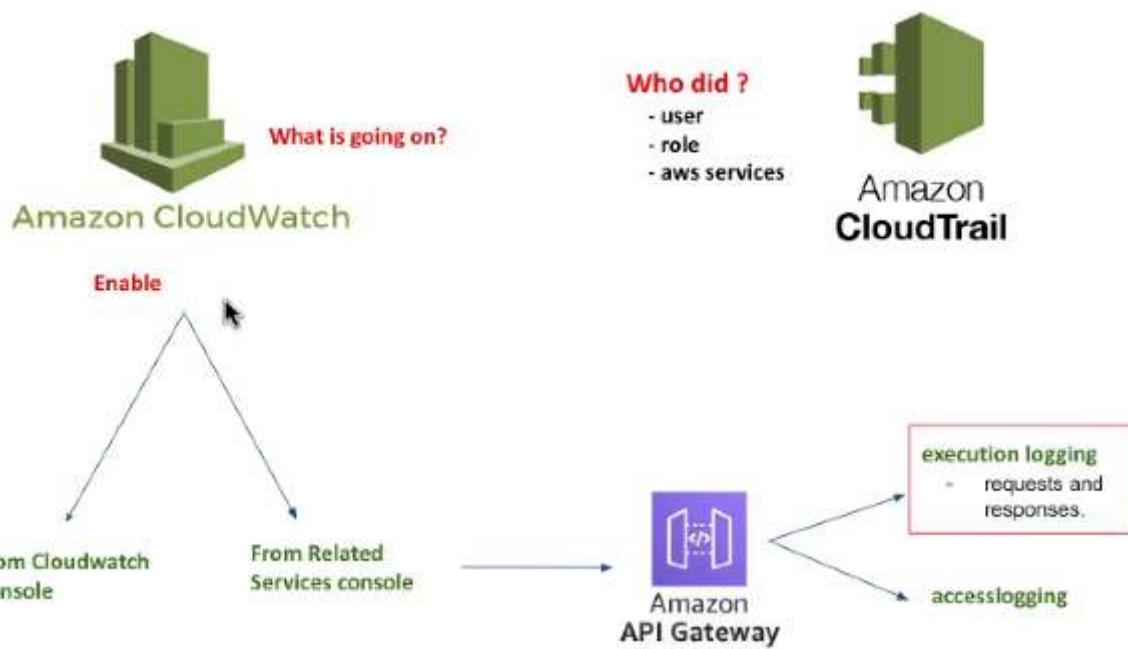
When requests come in faster than your function can scale, or when your function is at maximum concurrency, additional requests fail with a throttling error (429 status code).

Domain: Design Resilient Architectures

You are a solutions architect working for a data analytics company that delivers analytics data to politicians that need the data to manage their campaigns. Political campaigns use your company's analytics data to decide on where to spend their campaign money to get the best results for the efforts. Your political campaign users access your analytics data through an Angular SPA via API Gateway REST endpoints. You need to manage the access and use of your analytics platform to ensure that the individual campaign data is separate. Specifically, you need to produce logs of all user requests and responses to those requests, including request payloads, response payloads, and error traces. Which type of AWS logging service should you use to achieve your goals?

- A. Use CloudWatch access logging
- B. Use CloudWatch execution logging
- C. Use CloudTrail logging
- D. Use CloudTrail execution logging

Cevap B



Question 34 of 65

[Exit Quiz](#)

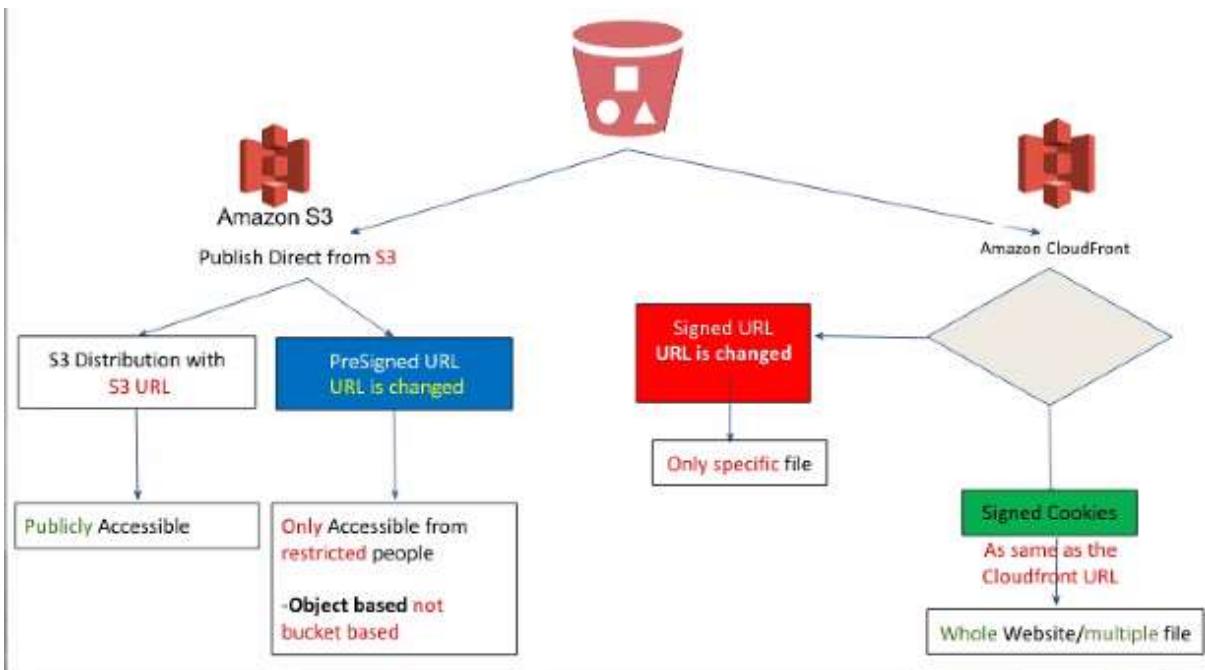
Domain: Design Secure Applications and Architectures

You are a solutions architect working for a media company that produces stock images and videos for sale via a mobile app and website. Your app and website allow users to gain access only to stock content they have purchased. Your content is stored in S3 buckets. You need to restrict access to multiple files that your users have purchased. Also, due to the nature of the stock content (purchasable by multiple users), you don't want to change the URLs of each stock item.

Which access control option best fits your scenario?

- A. Use CloudFront signed URLs
- B. Use S3 Presigned URLs
- C. Use CloudFront Signed Cookies
- D. Use S3 Signed Cookies

Cevap C mutiple files olduğu için cevap Cookies

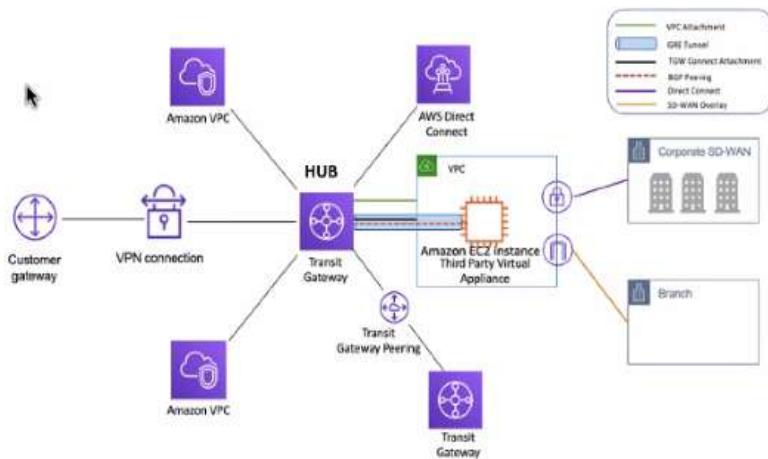


You are a solutions architect working for a financial services firm that operates applications in the hybrid cloud model. You have applications running on EC2 instances in your VPC that communicate with resources in your on-prem data center. You have a workload on an EC2 network interface in one subnet and a transit gateway association in a different subnet. Also, these two subnets are associated with different NACLs. You have set up Network Access Control List (NACL) rules to control the traffic to and from your EC2 instances and transit gateway.

Which of the following is true about the NACL rules for traffic from your EC2 instances to the transit gateway?

- A. Outbound rules use the source IP address to evaluate traffic from the instances to the transit gateway
- B. Outbound rules use the destination IP address to evaluate traffic from the instances to the transit gateway
- C. Outbound rules are not evaluated for the transit gateway subnet
- D. Inbound rules use the destination IP address to evaluate traffic from the transit gateway to the instances

Cevap B



Domain: Design High-Performing Architectures

A start-up firm has created account A using the Amazon RDS DB instance as a database for a web application. The operations team regularly creates manual snapshots for this DB instance in unencrypted format. The Projects Team plans to create a DB instance in other accounts using the snapshots. They are looking for your suggestion for sharing this snapshot and restoring it to DB instances in other accounts. While sharing this snapshot, it must allow only specific accounts specified by the project teams to restore DB instances from the snapshot.

What actions can be initiated for this purpose?

- A. From Account A, share the manual snapshot by setting the 'DB snapshot' visibility option as private. In other Accounts, directly restore to DB instances from the snapshot.
- B. From Account A, share the manual snapshot by setting the 'DB snapshot' visibility option as public. In other Accounts, directly restore to DB instances from the snapshot.
- C. From Account A, share the manual snapshot by setting the 'DB snapshot' visibility option as private. In other Accounts, create a copy from the snapshot and then restore it to the DB instance from that copy.
- D. From Account A, share the manual snapshot by setting the 'DB snapshot' visibility option as public. In other Accounts, create a copy from the snapshot and then restore it to the DB instance from that copy.

Cevap A

Snapshot permissions

Preferences
You are sharing an unencrypted DB snapshot. When you share an unencrypted DB snapshot, you give the other account permission to make a copy of the DB snapshot and to restore a database from your DB snapshot.

DB snapshot: database-01-snapshot
 Private Public

AWS account ID: 123456789123 **Add**

AWS account ID	Delete
123456789123	Delete

Cancel **Save**

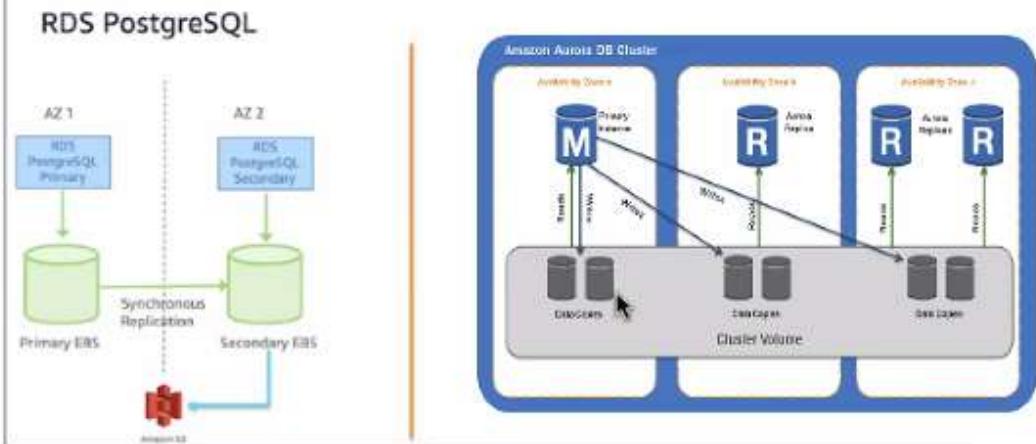
Domain: Design Resilient Architectures

An electronic manufacturing company plans to deploy a web application using the Amazon Aurora database. The Management is concerned about the disk failures with DB instances and needs your advice for increasing reliability using Amazon Aurora automatic features. In the event of disk failures, data loss should be avoided, reducing additional work to perform from the point-in-time restoration.

What design suggestions can be provided to increase reliability?

- A. Add Aurora Replicas to primary DB instances by placing them in different regions. Aurora's crash recovery feature will avoid data loss post disk failure
- B. Add Aurora Replicas to primary DB instances by placing them in different availability zones. Aurora storage auto-repair feature will avoid data loss post disk failure
- C. Add Aurora Replicas to the primary DB instance by placing them in different regions. Aurora Survivable page cache feature will avoid data loss post disk failure
- D. Add Aurora Replicas to the primary DB instance by placing them in different availability zones. Aurora's crash recovery feature will avoid data loss post disk failure

Cevap B



- Storage auto-repair >>> Prevent disk failure and data loss >>> Back up with cluster volume
- Survivable page cache >>> Cache and database mechanism is separate.
So even the DB fails, cache still works
- Crash recovery >>> Designed to recover from a crash almost instantaneously

Question 38 of 65

[Exit Quiz](#)

Domain: Design Cost-Optimized Architectures

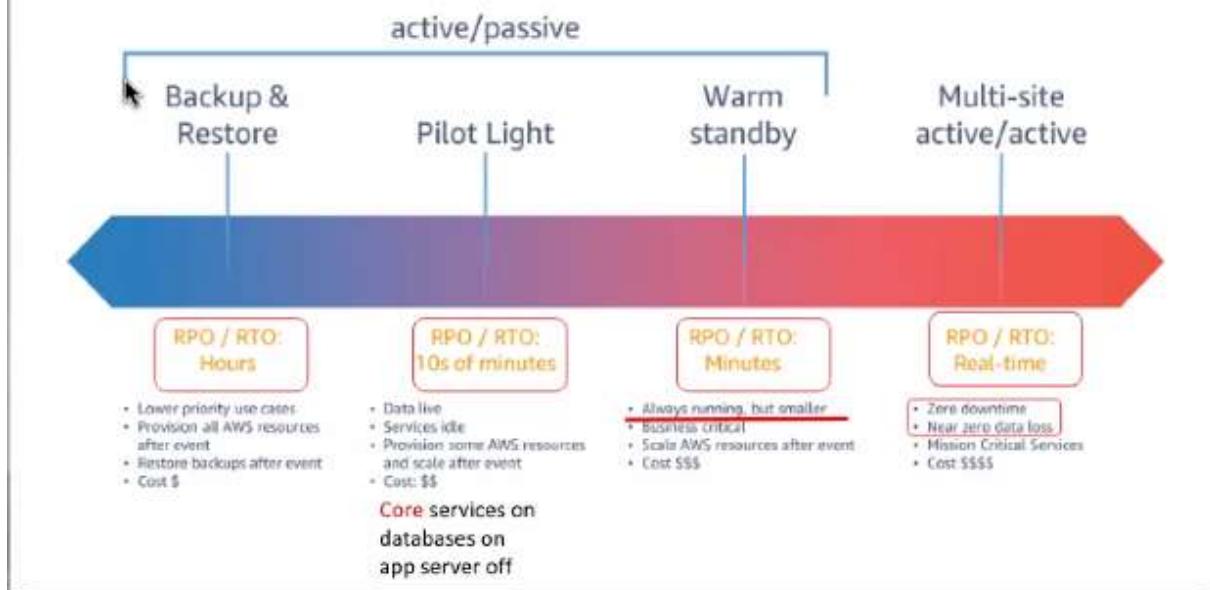
A financial institute has deployed a critical web application in the AWS cloud. The management team is looking for a resilient solution with RTO/RPO in 20 minutes during a disaster. They have budget concerns, and the cost of provisioning the backup infrastructure should not be very high. As a solution architect, you have been assigned to work on setting a resilient solution meeting the RTO/RPO requirements within the cost constraints.

Which strategy is suited perfectly?

- A. Multi-Site Active/Active
- B. Warm Standby
- C. Backup & Restore
- D. Pilot Light

Cevap D 20 dk içinde

Recovery Time Objective (RTO) >>>> How much time it takes ?
Recovery Point Objective (RPO) >>>> Avoiding losing data



Domain: Design Resilient Architectures

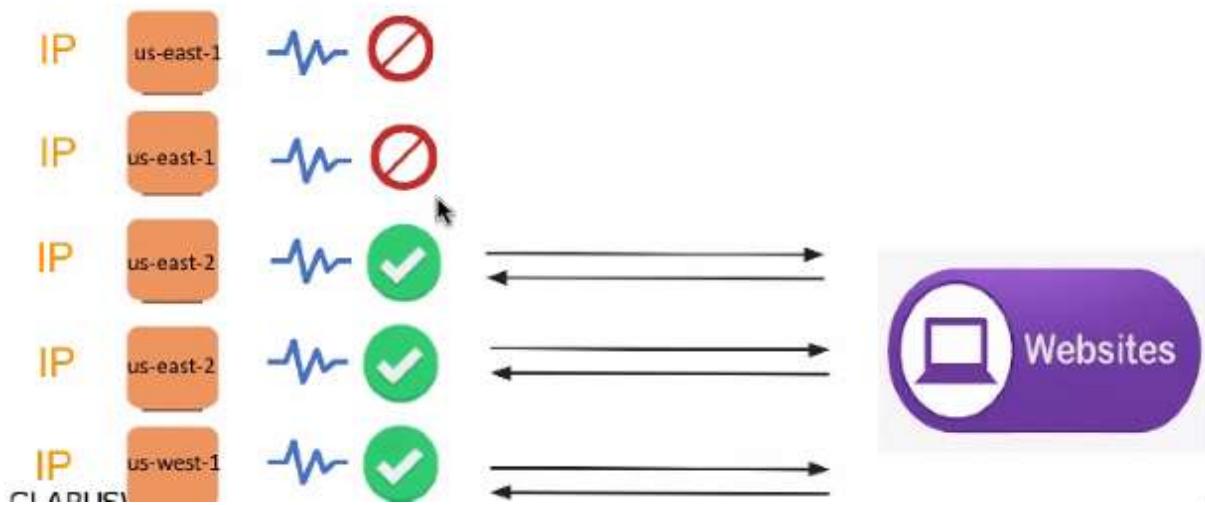
A critical application deployed in AWS Cloud requires maximum availability to avoid any outages. The project team has already deployed all resources in multiple regions with redundancy at all levels. They are concerned about the configuration of Amazon Route 53 for this application which should complement higher availability and reliability. Route 53 should be configured to use failover resources during a disaster.

What solution can be implemented with Amazon Route 53 for maximum availability and increased reliability?

- A. Associate multiple IP endpoints in different regions to Route 53 hostname.
Use a weighted route policy to change the weights of the primary and failover resources. So, all traffic is diverted to failover resources during a disaster
- B. Create two sets of public-hosted zones for resources in multiple regions.
During a disaster, update Route 53 public-hosted zone records to point to a healthy endpoint
- C. Create two sets of private hosted zones for resources in multiple regions.
During a disaster, update Route 53 private hosted zone records to point to a healthy endpoint
- D. Associate multiple IP endpoints in different regions to Route 53 hostname.
Using health checks, configure Route 53 to automatically failover to healthy endpoints during a disaster

Cevap D

► Multivalue Answer Routing Policy (8 IPs)



Domain: Design Cost-Optimized Architectures

An IT company is using EBS volumes for storing projects related work. Some of these projects are already closed. The data for these projects should be stored long-term as per regulatory guidelines and will be rarely accessed. The operations team is looking for options to store the snapshots created from EBS volumes. The solution should be cost-effective and incur the least admin work.

What solution can be designed for storing data from EBS volumes?

- A. Create EBS Snapshots from the volumes and store them in the EBS Snapshots Archive
- B. Use Lambda functions to store incremental EBS snapshots to AWS S3 Glacier
- C. Create EBS Snapshots from the volumes and store them in a third-party low-cost, long-term storage
- D. Create EBS Snapshots from the volumes and store them in the EBS standard tier

Cevap A

Snapshots (1/20)

Name	Snapshot ID	Size	Description
AWS Application Migration Service Base Snapshot	snap-0536abcf5b9947b64	1 GiB	AWS Appli
terraform_example	snap-079a82f599cc388b8	8 GiB	Created by

Snapshot ID: snap-0536abcf5b9947b64 (AWS Application Migration Service Base Snapshot)

Details	Permissions	Storage tier	Tags
Snapshot ID snap-0536abcf5b9947b64 (AWS Application Migration Service Base Snapshot)	Size 1 GiB	Progress Available (100%)	
Owner 046402772087	Volume ID vol-02c5c7f9191916b77	Started Fri Dec 10 2021 23:04:32 GMT+0300 (GMT+03:00)	Product codes
Encryption Not encrypted	KMS key ID	KMS key alias	KMS key ARN

When you need to access an archived snapshot, you can restore it from the archive tier to the standard tier, and then use it

Domain: Design Cost-Optimized Architectures

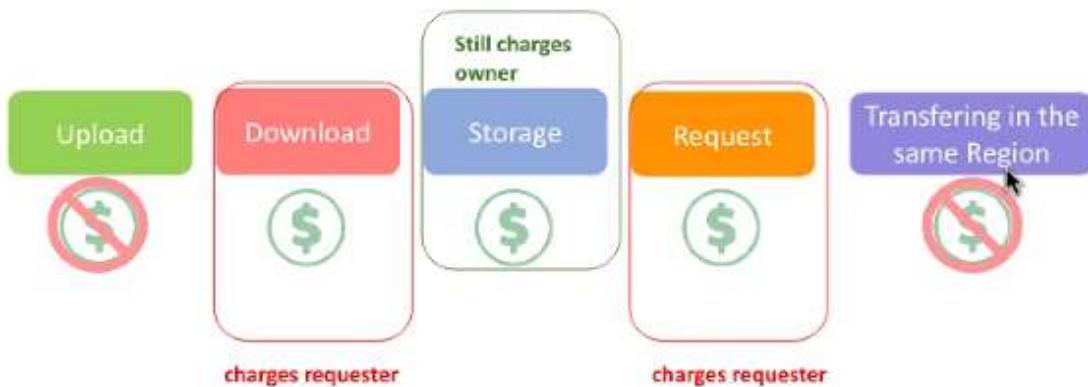
A company is sharing geospatial data with users in different AWS accounts for commercial purposes. Users from these accounts access a large amount of data stored in the Amazon S3 bucket in the us-east-1 region from different AWS regions. The company has incurred high charges for this data sharing belonging to transfer charges. The Finance Team is looking for an option to minimize these charges.

What solution can be designed to minimize this cost?

- A. Configure Amazon CloudFront in front of the Amazon S3 bucket to share data from the nearest edge locations
- B. Configure the Requester Pays option on the Amazon S3 bucket
- C. Share pre-signed URLs with the users to access data from the Amazon S3 bucket
- D. Replicate data to Amazon S3 buckets in all the regions to enable users to download data from local Amazon S3 buckets

Cevap B

► Requester Pays buckets for storage transfers and usage



After you configure a bucket to be a Requester Pays bucket, requesters must include `x-amz-request-payer` in their API request header, for **DELETE, GET, HEAD, POST, and PUT requests**, or as a parameter in a REST request to show that they understand that they will be charged for the **request and the data download**.

Domain: Design Secure Architectures

An IT firm is planning to store all its critical project-related documents in an Amazon S3 bucket. All these files should be encrypted at rest. As per security guidelines, firms need to manage the encryption process internally, but keys used for the encryption should not be stored locally.

How can encryption solutions be designed to meet the data encryption guidelines?

- A. Use a key stored within the application for client-side encryption while uploading/downloading data from the Amazon S3 bucket
- B. Use an AWS KMS key for client-side encryption while uploading/downloading data from the Amazon S3 bucket
- C. Use a customer-provided key for client-side encryption while uploading/downloading data from the Amazon S3 bucket
- D. Use an Amazon S3 bucket key for client-side encryption while uploading/downloading data from the Amazon S3 bucket

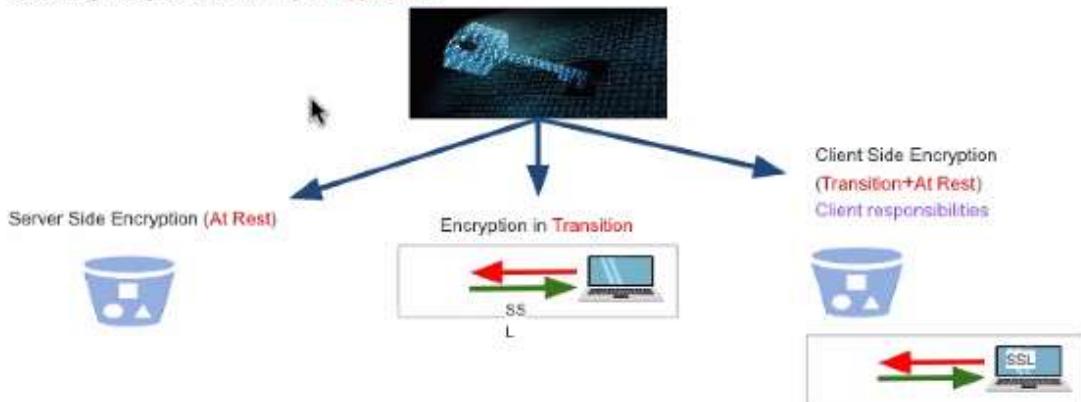
Cevap B

kendi internal yapacağım . client side anahtar sende dursun

Client-side encryption is the act of encrypting your data locally to ensure its security as it passes to the Amazon S3 service. The Amazon S3 service receives your encrypted data; it does not play a role in encrypting or decrypting it.

To enable client-side encryption, you have the following options:

1. Use a key stored in AWS Key Management Service (AWS KMS). You encrypt data with KMS key . S3 it does not play a role in encrypting or decrypting it.
2. Use a key that you store within **your application**.



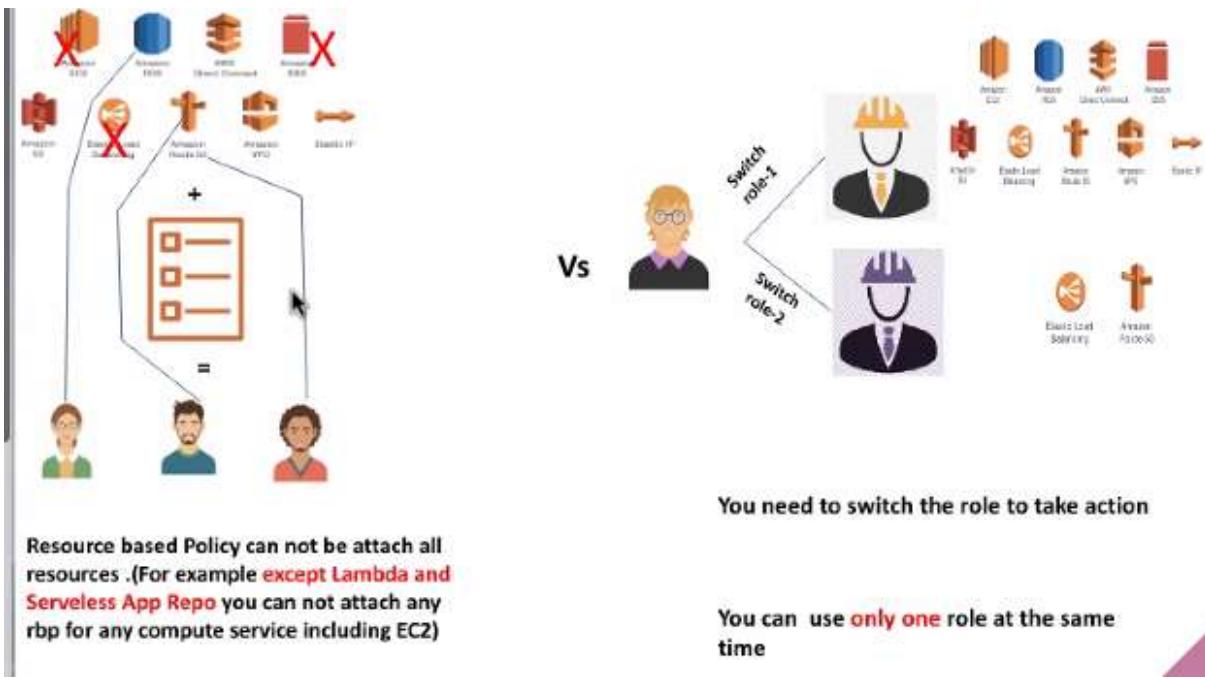
Domain: Design Secure Architectures

The operations team requires users to grant permission to access AWS resources. Resources are categorized as critical and non-critical resources. For granting these permissions, the principle of least privilege should be followed. Users will be accessing critical resources only on a required basis and non-critical resources on a daily basis. The proposed solution should be an efficient solution considering a large number of resources and users.

How can access permissions be designed for this purpose?

- A. Assign permissions to roles for access to critical resources. Users can switch to these roles to access critical resources. Users can use regular credentials to access non-critical resources
- B. Create separate user ids for accessing non-critical and critical resources. Users can switch between two user ids while accessing non-critical and critical resources
- C. Grant full access to all the users to both critical and non-critical resources
- D. Create a resource-based policy on critical resources. Grant permission to specific users who require access to these resources. Users can use regular credentials to access non-critical resources

Cevap A



Question 44 of 65

[Exit Quiz](#)

Domain: Design Secure Architectures

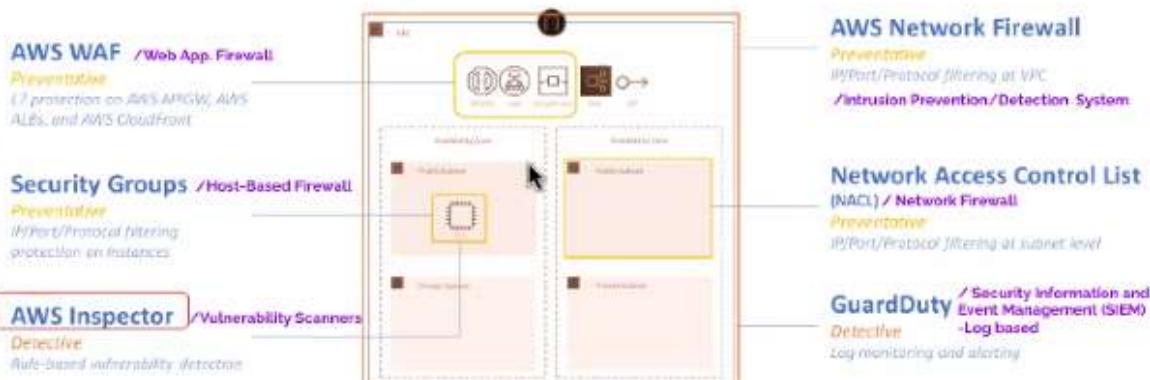
A company has recently installed multiple software on an Amazon EC2 instance for its new web application. Recently there was an incident that caused a major outage to this web application. The root cause was identified as network ports were open which were exploited. The Security Team is looking for a detailed report on the Amazon EC2 instance which should gather information on all network ports open unintentionally and not used for any service. The solution should provide a list that can help remediate these findings based on criticality.

Which of the following approaches can be initiated to get the required reports?

- A. Implement Amazon Inspector on the Amazon EC2 instance
- B. Implement Amazon GuardDuty on the Amazon EC2 instance
- C. Implement Amazon Detective on Amazon EC2 instance
- D. Implement AWS Artifact on Amazon EC2 instance

Cevap A

▶ Summary of AWS Security Services



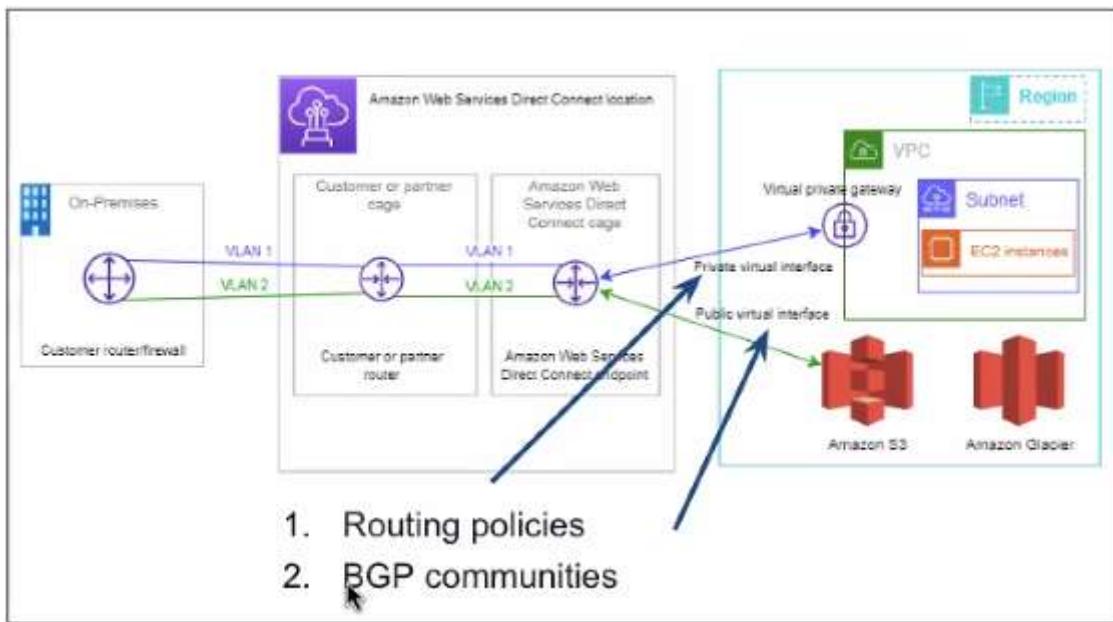
Domain: Design Resilient Architectures

A Direct Connect connection is established from an on-premises location to the AWS us-west-1 region using BGP (Border Gateway Protocol). This connectivity will only be used for accessing Amazon S3 from on-premises. IT head wants to ensure on-premises prefixes are advertised only to the us-west-1 region. These prefixes should not be exported to any other AWS region.

How can the BGP routing policy be designed to control these advertisements?

- A. Set NO_EXPORT BGP community tags while advertising prefixes to AWS in the us-west-1 region
- B. Do not apply any BGP community tags; advertising prefixes can never be controlled
- C. Use BGP community tags while advertising prefixes to AWS in the us-west-1 region
- D. Use Local preference BGP community tags while advertising prefixes to AWS in the us-west-1 region

Cevap C



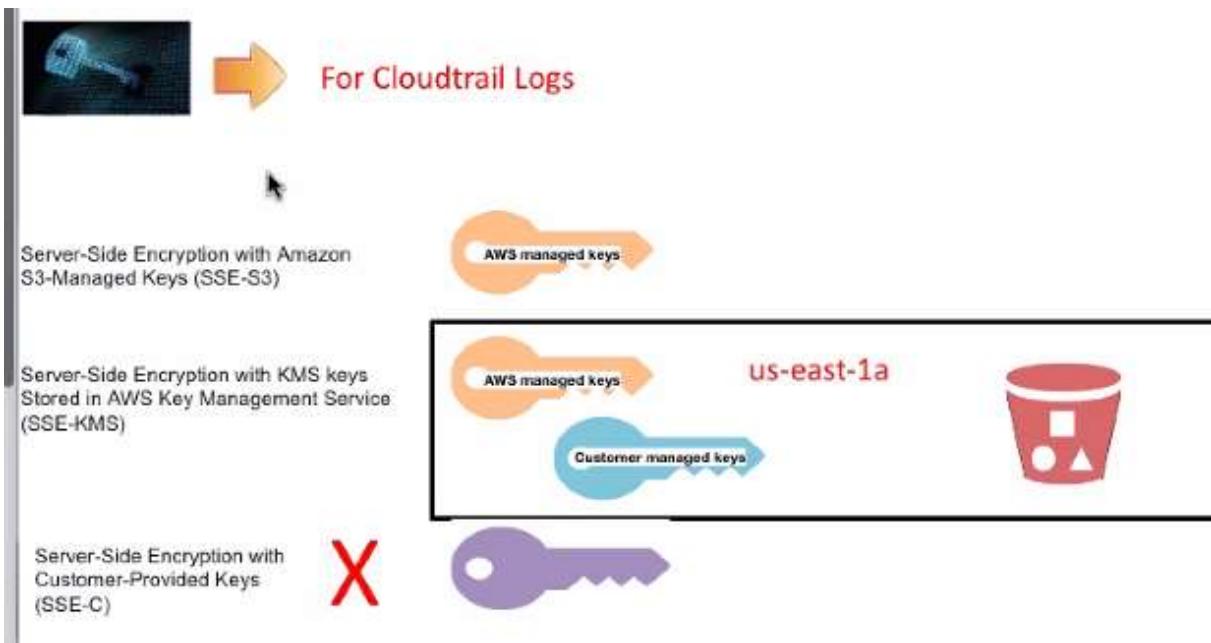
An engineering firm uses Amazon CloudTrail to record user activities across multiple accounts. Log files for this CloudTrail are stored in the Amazon S3 bucket in the us-east-1 region. Keys used for encrypting these logs should be managed by the Security team in this firm. Only specific users in the team should have permission to use this key for encrypting and decrypting log files. You have been assigned to work on this solution to suggest an efficient solution for additional security to log files from multiple accounts.

Which of the following solutions can you propose?

- A. Use Server-side encryption with Amazon S3-managed encryption keys (SSE-S3) for CloudTrail log files. Use different keys for encrypting and decrypting log files for multiple accounts across different regions
- B. Use Server-side encryption with AWS KMS-managed keys (SSE-KMS) for CloudTrail log files. Create a KMS key in the same region as the S3 bucket storing the log files. Use the same key for encrypting and decrypting log files for multiple accounts across different regions
- C. Use Server-side encryption with Amazon S3-managed encryption keys (SSE-S3) for CloudTrail log files. Use the same keys for encrypting and decrypting log files for multiple accounts across different regions
- D. Use Server-side encryption with AWS KMS-managed keys (SSE-KMS) for CloudTrail log files. Create a KMS key in a different region than the S3 bucket storing the log files. Use the different keys for encrypting and decrypting log files for multiple accounts across different regions

Cevap B

Key ler sadece security team tarafından kontrol edilsin



Domain: Design High-Performing Architectures

An online educational platform is developing a web application that millions of students will use to access educational content. They will be using the Amazon EC2 instance for their compute services and are looking for a high-performance shared storage solution that will be accessed parallelly. For each of the user sessions accessing a file system, it should be able to provide high throughput of more than 25 GB/s.

Which file system can be selected to meet performance requirements?

- A. FSx for NetApp ONTAP
- B. FSx for Lustre
- C. FSx for Windows File Server
- D. FSx for OpenZFS

Cevap B

Amazon FSx file systems



Amazon FSx for NetApp ONTAP

Fully managed shared storage built on NetApp's popular ONTAP file system.



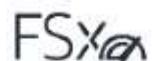
Amazon FSx for OpenZFS

Fully managed shared storage built on the popular OpenZFS file system.



Amazon FSx for Windows File Server

Fully managed shared storage built on Windows Server.



Amazon FSx for Lustre

Fully managed shared storage built on the world's most popular high-performance file system.

- Compatible with
Windows Active Directory

- HPC
- Compatible with S3

- Compatible with
FSx File Gateway

Question 48 of 65

[Exit Quiz](#)

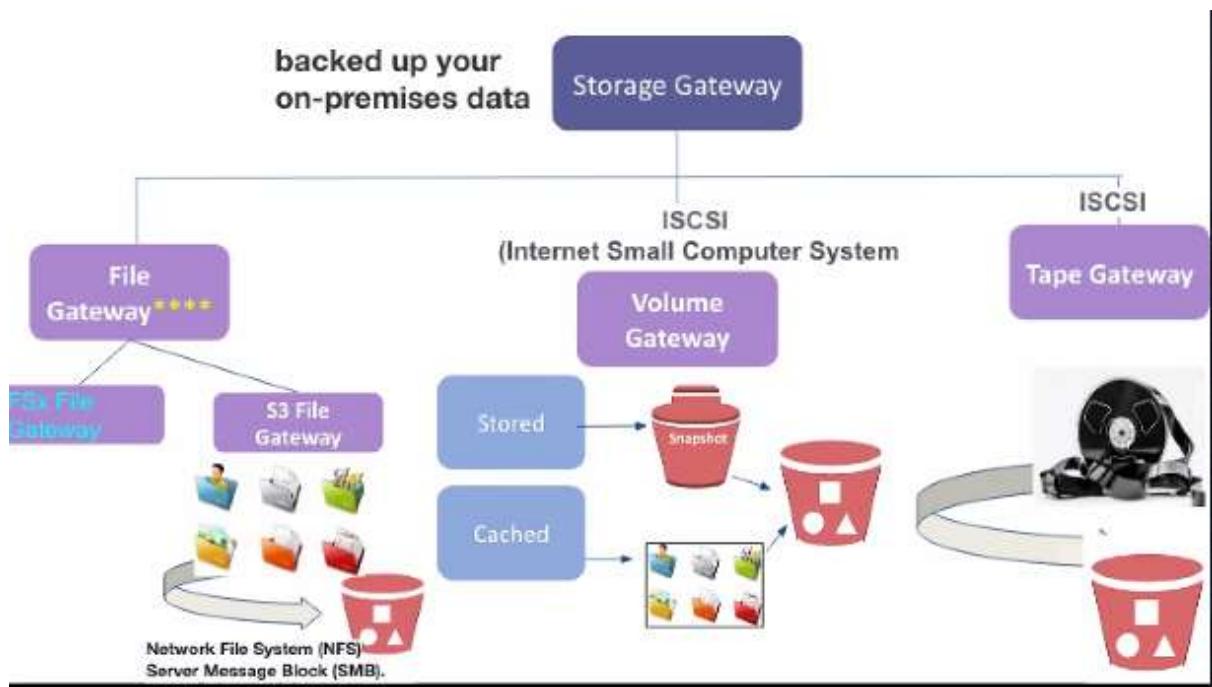
Domain: Design High-Performing Architectures

An R&D firm is looking to store research data in the AWS cloud. At on-premises locations, currently, they are storing this data in Windows servers and are looking for a storage option in AWS without any changes in the application. This data transfer to the cloud should be transparent to the end-users and should not have any performance impact. Files stored in AWS need to be accessed by SMB (Server Message Block) clients as well.

How can hybrid cloud storage be designed?

- A. Implement hybrid cloud storage using the Amazon S3 file gateway
- B. Implement hybrid cloud storage using cached mode volume gateway
- C. Implement hybrid cloud storage using Amazon FSx file gateway
- D. Implement hybrid cloud storage using a stored mode volume gateway

Cevap C Windows serverlar var



Question 49 of 65

[Exit Quiz](#)

Domain: Design Cost-Optimized Architectures

A large manufacturing company is looking to track IoT sensor data collected from thousands of equipment across multiple factory units. This is extremely high-volume traffic that needs to be collected in real time and should be efficiently visualized. The company is looking for a suitable database in the AWS cloud for storing these sensor data.

Which of the following cost-effective databases can be selected for this purpose?

- A. Send sensor data to Amazon RDS (Relational Database Service) using Amazon Kinesis and visualize data using Amazon QuickSight
- B. Send sensor data to Amazon Neptune using Amazon Kinesis and visualize data using Amazon QuickSight
- C. Send sensor data to Amazon DynamoDB using Amazon Kinesis and visualize data using Amazon QuickSight
- D. Send sensor data to Amazon Timestream using Amazon Kinesis and visualize data using Amazon QuickSight

Cevap D

Database type	Use cases	AWS service
Relational	Traditional applications, enterprise resource planning (ERP), customer relationship management (CRM), ecommerce	 Amazon Aurora  Amazon RDS  Amazon Redshift
Key-value	High-traffic web applications, ecommerce systems, gaming applications	 Amazon DynamoDB
In-memory	Caching, session management, gaming leaderboards, geospatial applications	 Amazon ElastiCache  Amazon MemoryDB for Redis
Document	Content management, catalogs, user profiles JSON workloads	 Amazon DocumentDB (with MongoDB compatibility)
Wide column	High-scale industrial apps for equipment maintenance, fleet management, and route optimization	 Amazon Keyspaces  Apache Cassandra
Graph	Fraud detection, social networking, recommendation engines	 Amazon Neptune
Time series	Internet of Things (IoT) applications, DevOps, industrial telemetry	 Amazon Timestream
Ledger	Systems of record, supply chain, registrations, banking transactions	 Amazon Ledger Database Services (QLDB)

Domain: Design High-Performing Architectures

A start-up firm is using a **JSON-based database** for content management. They are planning to rehost this database to AWS Cloud from on-premises. For this, they are looking for a suitable option to deploy this database, which can **handle millions of requests per second with low latency**. Databases should have a flexible schema that can store any type of user data from multiple sources and should effectively process similar data stored in different formats.

Which of the following databases can be selected to meet the requirements?

- A. Use Amazon DocumentDB (with MongoDB compatibility) in the AWS cloud to rehost the database from an on-premises location
- B. Use Amazon Neptune in the AWS cloud to rehost the database from an on-premises location
- C. Use Amazon Timestream in AWS cloud to rehost database from an on-premises location
- D. Use Amazon Keyspaces in AWS cloud to rehost database from an on-premises location

Cevap A

James is trying to launch a t2.micro EC2 instance having permissions as per the AmazonEC2FullAccess Policy. He is using Amazon Linux 2 AMI and proceeded without a key pair. For the root volume, he has selected General Purpose EBS Volume with 8 GB of storage and has encrypted it by manually specifying the encryption key name. After reviewing the configurations, he launches the instance.

The instance appears in the list, but the instance has changed its state from Pending to Shutting Down to Terminated immediately without being in Running State. Upon finding it, he checked it happened because of the error: Client.InternalError: Client error on launch. Which of the following options describe the root cause and the solution for the given problem in the best way?

- A. The EBS Volume is encrypted with a KMS key which James has no permission for decryption. The admin should give him permission to access KMS Key
- B. The mentioned policy applies to only certain instance types. James should try launching the EC2 instance with another type than t2.micro
- C. James has created the instance without a key pair. He should launch the instance again by using a key pair
- D. The EBS Volume limits in his account have been exceeded. He needs to submit a request to AWS to increase the limit

Cevap A

The screenshot shows the AWS CloudFormation 'Create New Stack' wizard at the 'File systems' step. It displays a table of file system configurations:

Volume 1 (AMI Root) (Custom)		
Storage type Info	Device name - required Info	Snapshot Info
EBS	/dev/xvda	snap-07bb85133b158fe99
Size (GiB) Info	Volume type Info	IOPS Info
8	gp2	100 / 3000
Delete on termination Info	Encrypted Info	KMS key Info
Yes	Encrypted	Select Q Specify a custom value... Select (default) aws/ebs Key ID: 050b1069-46f9-4999-8119-b665edf12e38 149cf4c-85e0-43d7-8477-3cd965cb65e1

Below the table, there is a note: "Free tier eligible customers can get up to 30 GB of EBS General Purpose". At the bottom, there is a "Add new volume" button.

Domain: Design Resilient Architectures

You are designing an architecture for one of your client's applications. You created a VPC within the us-east-1 region consisting of two public subnets and a private subnet. The first public subnet contains a web server and the private subnet contains the database server. All the other VPC Networking components are also set including Internet gateway, route tables, security groups, etc. Your team lead reviewed the architecture and informed you about the client's requirement of making the current architecture highly available with MINIMAL COST. He also specified that the client could accept a minor downtime, but the cost is the priority here.

Which options can make your current architecture highly available with minimum effort and costs?

- A. Span the architecture components to a second Availability zone by creating similar resources in it as AZ-1 and then distributing the traffic using an application load balancer
 - B. Create a standby instance in the second public subnet. Attach a secondary ENI to the instance in the first public subnet. In case of a failure, detach this ENI and attach it to the standby instance
 - C. Span the architecture components to a second region by creating similar resources in it as available in US-East-1 and then distributing the traffic using an application load balancer
 - D. Implement dynamic scaling in this architecture by using Cloudwatch, Elastic Load Balancing, and Auto Scaling together. This solution will bring automated scaling and high availability with zero downtime
-

Cevap B

An *elastic network interface* is a logical networking component in a VPC that represents a virtual network card. It can include the following attributes:

- A primary private IPv4 address from the IPv4 address range of your VPC
- One or more secondary private IPv4 addresses from the IPv4 address range of your VPC
- One Elastic IP address (IPv4) per private IPv4 address
- One public IPv4 address
- One or more IPv6 addresses
- One or more security groups
- A MAC address
- A source/destination check flag
- A description



When you move a network interface from one instance to another, **network traffic is redirected to the new instance**.

Question 53 of 65

[Exit Quiz](#)

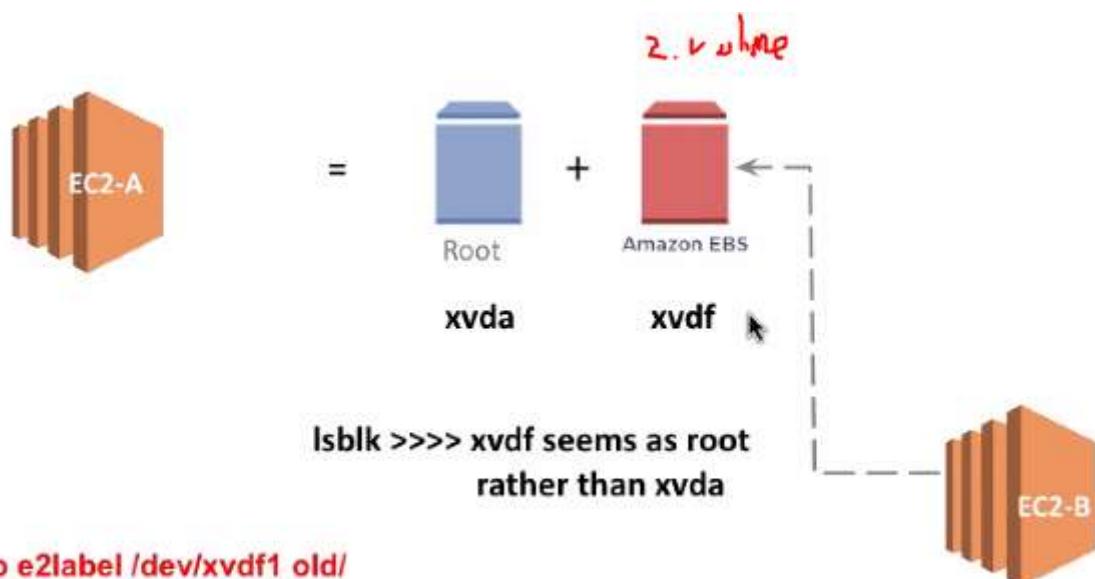
Domain: Design High-Performing Architectures

You have launched an EC2 instance with a **General Purpose EBS Volume of 16384 GB** as its root volume (`/dev/xvda`). After working on it for some time, you observed the need to have more storage with this instance. You **found a volume (`/dev/xvdfl`) available in your account**. Therefore this volume is attached as an additional volume with your EC2 instance.

One of your teammates informed you that **this volume (`/dev/xvdfl`) was created from the snapshot of the root volume of another instance**. Later on, upon rebooting the instance, you **connect to it and observe that the instance has booted from `/dev/xvdfl`, whereas it was supposed to boot from `/dev/xvda` (original root volume)**. What should you do now to **resolve the problem and let the instance boot from the correct volume?**

- A. Keep the additional volume in an Availability zone other than the one where the EC2 instance and its root volume reside, then attach it to the EC2 instance
- B. Change the label of the additional volume using the `e2label` command
- C. Unmount the additional volume and add more storage capacity to the existing General Purpose volume only
- D. Change the label of the additional volume using the `e2changelabel` command

Cevap B



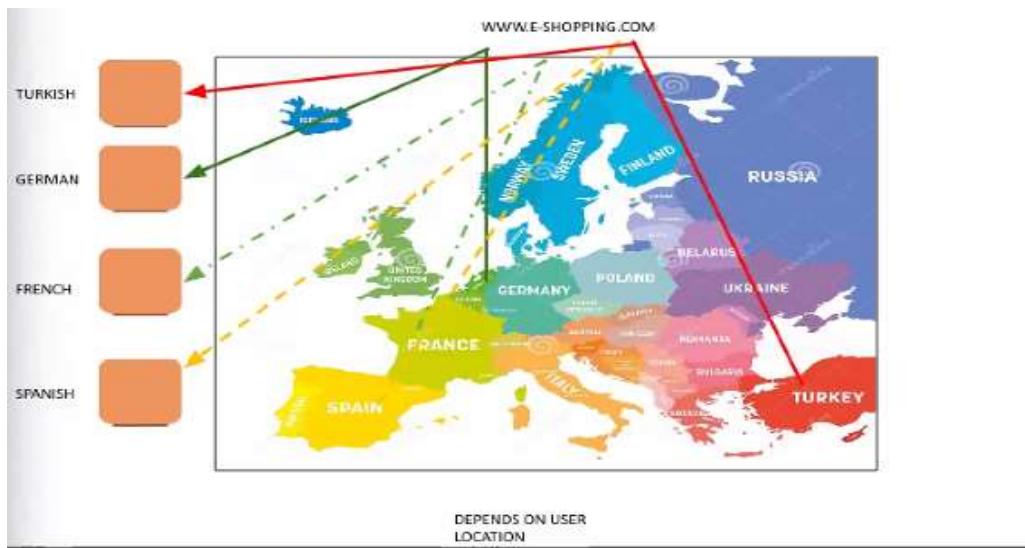
Domain: Design Resilient Architectures

You are working as a solutions architect in an E-Commerce based company with users from around the globe. There was feedback coming from various users of different countries to have the website content in their local languages. So, the company has now translated the website into multiple languages and is rolling out the feature soon for its users.

Now you need to send the traffic based on the location of the user. For example, if a request comes from Japan, it should be routed to the server in the ap-northeast-1 (Tokyo) region where the application is in the Japanese language. You can do so by specifying the IP address of that particular server while configuring the records in Route 53. Which one of the following routing policies should you use in Amazon Route 53 that will fulfill the given requirement?

- A. Weighted Routing Policy
- B. Geoproximity Routing Policy
- C. Geolocation Routing Policy
- D. Multivalue Answer Routing Policy

Cevap C



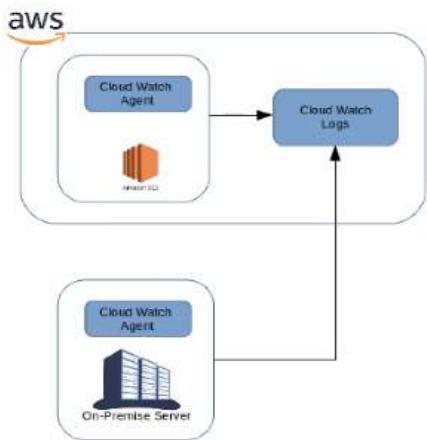
Domain: Design Resilient Architectures

You are working as a solutions architect in a team who is handling the infrastructure provisioning, scaling, load balancing and monitoring, etc for a workload. This workload is hosted on a fleet of Linux servers on-premises. You have been assigned a task to capture the application logs, percentage of total disk space that is used, the amount of memory currently in use, and the percentage of time that the CPU is idle from the Linux servers.

Using this data, your team is going to build an operational dashboard. This dashboard will help set high-resolution alarms to alert and take automated actions and understand the service and application health easily. Which one of the following options is the best solution to satisfy this requirement?

- A. Use Amazon Cloudtrail standard metrics and logs to collect the required metrics and application logs
- B. Use Amazon Cloudtrail unified agent to collect the required metrics and application logs
- C. Use Amazon Cloudwatch Standard metrics and logs to collect the required metrics and application logs
- D. Use Amazon Cloudwatch unified agent to collect the required metrics and application logs

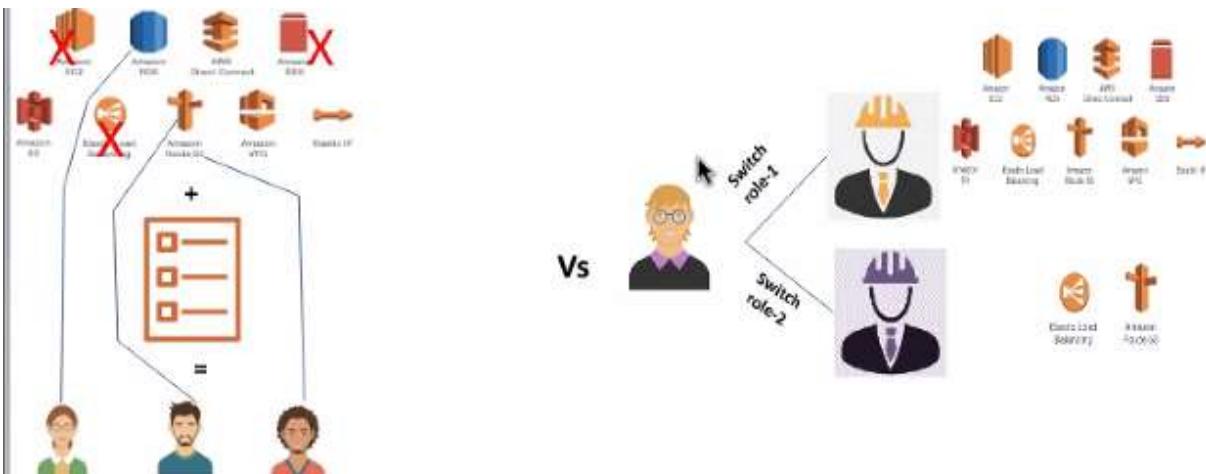
Cevap D



You are an IAM User in Account A and have created a Lambda function. One of the members of the Testing team, who is an IAM User in Account B, needs to perform some test cases and keep the code in his local system for record purposes. For the same, he needs to invoke your Lambda function in his account and should be able to download the code as well. Also, you need to make sure that the User of Account B can still perform his normal operations along with this Lambda invocation requirement. Which options will allow User B to perform the required actions in his account with minimal effort?

- A. Use a Resource-based policy with your Lambda function and grant permissions using `lambda:*`
- B. Create an IAM Role in Account A so that User B can assume this role and perform the required actions
- C. Use a Resource-based policy with your Lambda function and grant permissions using `lambda:InvokeFunction` and `lambda:GetFunction`
- D. Create an IAM User for User B in Account A so that he can log in and perform the required actions

Cevap C



Resource based Policy can not be attach all resources .(For example **except Lambda and Serverless App Repo** you can not attach any rbp for any compute service including EC2)

You need to switch the role to take action

You can use **only one** role at the same time

Question 57 of 65

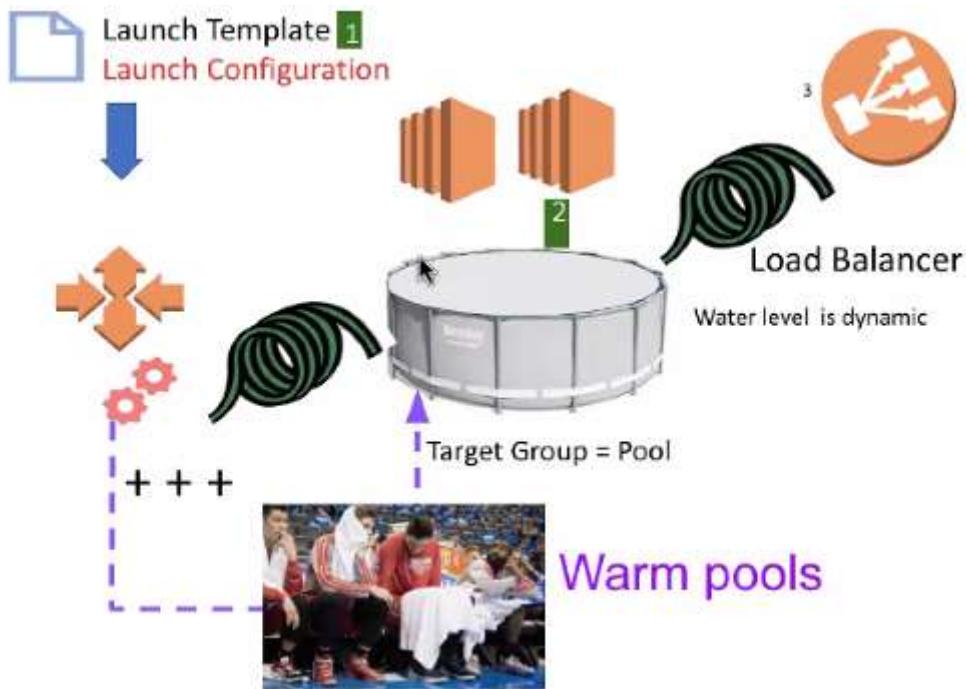
[Exit Quiz](#)

Domain: Design High-Performing Architectures

You have designed an architecture for a highly available web application with a dynamic workload. This architecture consists of a set of EC2 instances managed by an Auto Scaling Group and a load balancer for the even distribution of traffic. Your manager has reviewed the functioning of this architecture and is concerned about the latency of the scaling out process. Whenever there is a need to scale out, it takes a few seconds to a few minutes for Auto Scaling to launch the instance, register it with the load balancer, and start serving the requests. The Manager has asked you to look for a way to reduce this latency so that instances can start serving the requests as quickly as possible. Which of the following options can satisfy the given requirement?

- A. Use Warm Pool with Auto Scaling
- B. Use the Desired capacity parameter to scale out quickly
- C. Use Lifecycle hooks with Auto Scaling
- D. Use the Scheduled scaling type to scale out quickly

Cevap A



Question 58 of 65

Exit Quiz

Domain: Design High-Performing Architectures

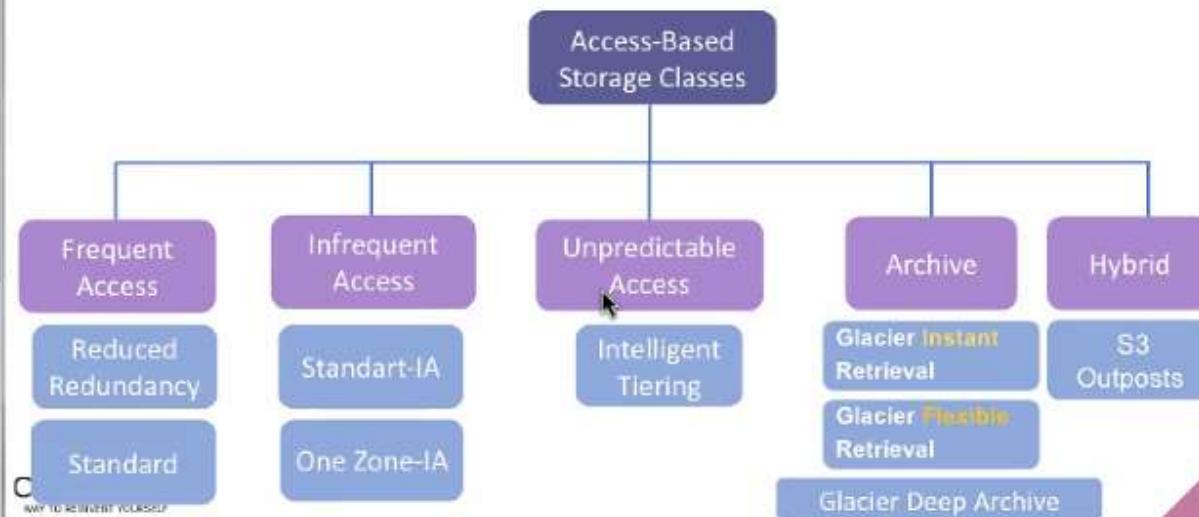
A Healthcare Industry has a requirement to store patient records on AWS securely and cost-effectively. This data isn't frequently accessed but needs to be kept as records for about six years as per the industry's policies. But there can be scenarios where they might need to query this data within minutes. They will need to access only the required data by performing filtering over the existing data. As a Solutions Architect, which of the following solutions will you recommend for the given scenario?

- A. Store the data in Amazon S3 Standard and Use S3 Select to query the data
- B. Store the data in Amazon S3 Glacier and Use Glacier Select to query the data
- C. Store the data in Amazon Redshift and perform SQL queries to retrieve the required data
- D. Store the data in Amazon S3 Standard Infrequent Access and Use S3 Select to query the data

Cevap B

Storage Classes

Summary of Storage Classes



Question 59 of 65

Exit Quiz

Domain: Design High-Performing Architectures

You have stored some of your documents as objects in the Amazon S3 Standard Storage Class. You have an urgent requirement of fetching just a subset of the data from the S3 Object. As a best practice, you want to reduce the latency of data retrieval and the amount of data that S3 will transfer to your location. You have tried using S3 Select to filter and query over the required subset of data with an SQL Expression of about 300 KB. But you have not been able to get the results yet. What can be the possible cause behind this failure?

- A. You can't perform the queries using S3 Select on the data stored in S3 Standard Class
- B. The maximum length of a SQL expression is limited to 200 KB
- C. Use Glacier Select instead of S3 Select to query the data successfully
- D. The maximum length of a SQL expression is limited to 256 KB

Cevap D

With Amazon S3 Select, you can use simple structured query language (SQL) statements to filter the contents of an Amazon S3 object and retrieve just the subset of data that you need. By using Amazon S3 Select to filter this data, you can reduce the amount of data that Amazon S3 transfers, which reduces the cost and latency to retrieve this data.

The following limits apply when using Amazon S3 Select:

- The maximum length of a SQL expression is 256 KB.
- The maximum length of a record in the input or result is 1 MB.
- Amazon S3 Select can only emit nested data using the JSON output format.
- You cannot specify the S3 Glacier Flexible Retrieval, S3 Glacier Deep Archive, or REDUCED_REDUNDANCY storage classes. For more information, about storage classes see [Storage Classes](#).

Domain: Design Resilient Architectures

You have built a serverless architecture composed of Lambda Functions exposed through API Gateway for one of your client's applications. For the database layer, you have used DynamoDB. Your team lead has reviewed the architecture and is concerned about the cost of numerous API calls being made to the backend (Lambda Functions) for so many similar requests. Also, the client is concerned about providing as low latency as possible for the application users' requests. You have to look for a solution where the latency and overall cost can be reduced for the current architecture without much effort.

- A. Cache the computed request's responses using the CloudFront CDN caching
- B. Use the API Gateway QuickResponse feature to reduce the latency and number of calls to the backend
- C. Enable API Gateway Caching to cache the computed request's responses
- D. Adjust API Gateway Throttling settings to reduce the latency and number of calls to the backend

Cevap C

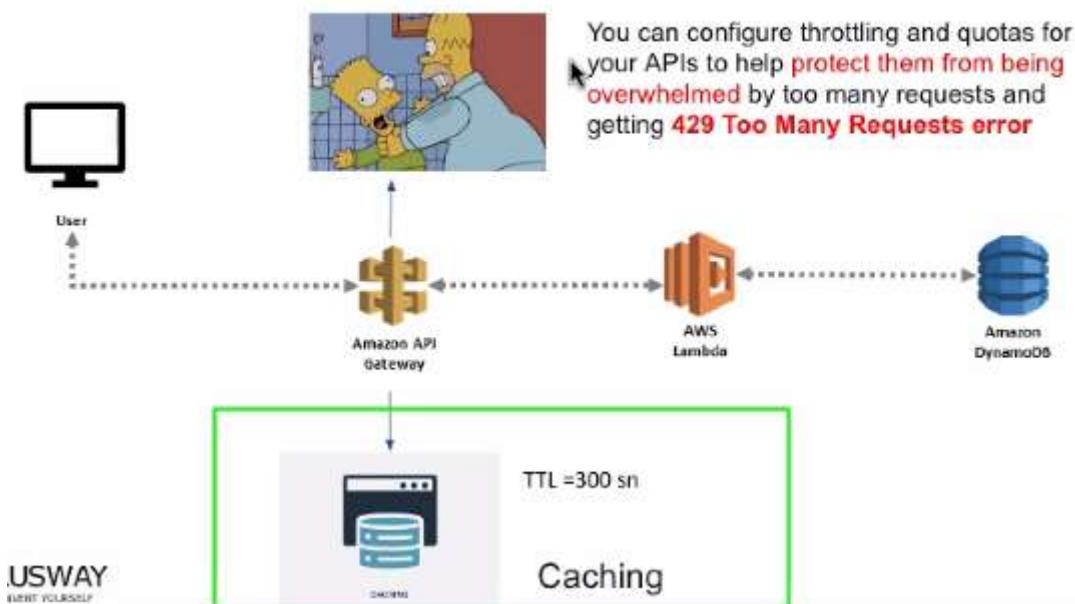
Explanation:

Correct Answer: C

- Option A is incorrect. Cloudfront CDN caching is mainly used for caching the responses to limited HTTP requests. API Gateway can cache the responses to requests of any type. Also, configuring Cloudfront CDN would be an added effort. You have to select a solution with minimal effort. Therefore, this option can be eliminated.
- Option B is incorrect because there is no feature like API Gateway QuickResponse. This option is just a distractor and thus, can be eliminated.
- Option C is CORRECT. When you enable caching on API Gateway, it caches the responses of the requests processed by the backend. When a similar request comes, it will serve it quickly from the API Gateway Cache itself instead of passing the call to the backend. This will reduce the number of calls made to the backend, eventually reducing the overall cost. Also, this will help in reducing the latency of the responses sent to the user.

Enabling API caching to enhance responsiveness

Throttle API requests



Domain: Design High-Performing Architectures

You have designed a gaming application with game servers hosted on EC2 Instances. For the leaderboards, you have used a DynamoDB table named Scores. Your manager has instructed you to build a feature to publish a Congratulations and Rewards Message with the winner's name on the app's Social Media network. This message should be posted each time the HighestScore Attribute is updated in the Scores DynamoDB table. Which of the following options can satisfy the requirement in the BEST way?

- A. DynamoDB Streams with AWS Lambda
- B. DynamoDB Global Tables with AWS Lambda
- C. DynamoDB Accelerator with AWS Lambda
- D. DynamoDB Transactions with AWS Lambda

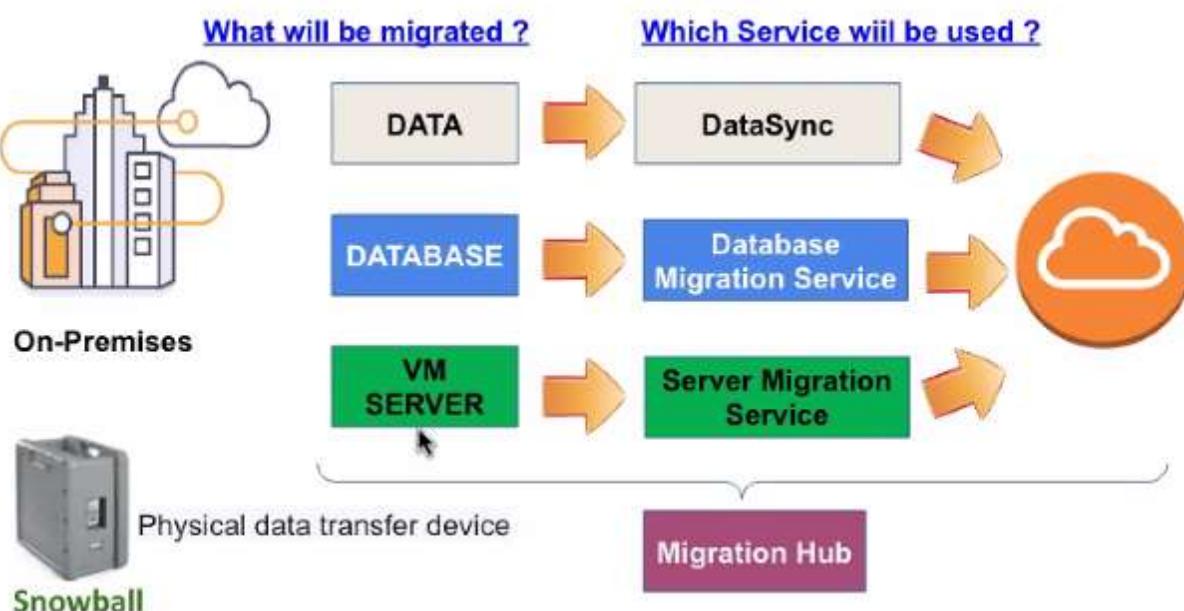
Cevap A



Domain: Design High-Performing Architectures

You are working in an organization who is using a Multi-Cloud environment. They have workloads deployed and data stored in Google cloud, Microsoft Azure, and AWS Cloud as well. Your manager has informed you about the decision to migrate one of the workloads from Google Cloud to AWS. For the same, he has asked you to migrate all the data stored in Google cloud storage buckets to Amazon S3 buckets without much effort. Which of the following services can help in accomplishing this task?

- A. AWS Migration Hub
- B. Amazon Storage Gateway
- C. AWS DataSync
- D. S3 Transfer Acceleration

Cevap C

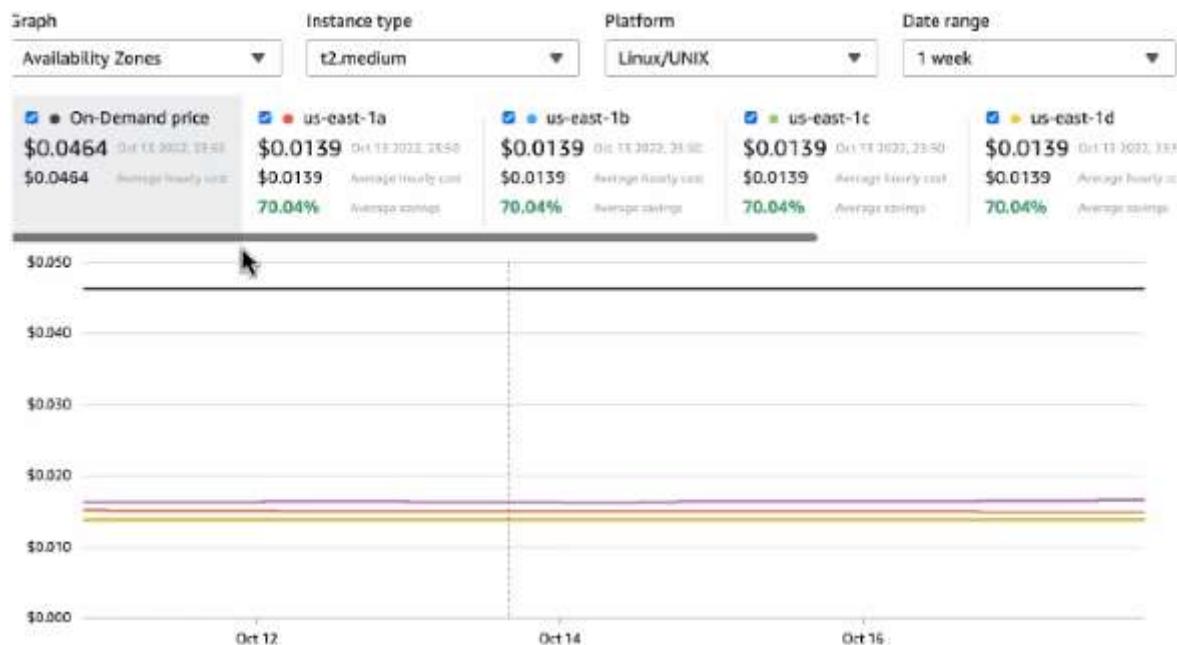
Domain: Design Cost-Optimized Architectures

Jackson handles the capacity provisioning for an E-Commerce application's workload. A sale is starting in the coming week, and the team has predicted a huge traffic spike during the sale days. Jackson has provisioned 5 t2.large Reserved Instances for the consistent portion of the workload. As per the predictions, there will be huge spikes in traffic which will be temporary.

Jackson is looking for a way to serve this portion of the workload in the MOST cost-effective way without much effort. Note that the application is already designed to recover from any instance failures. Which of the following options satisfies the given requirement in the best way?

- A. Use Savings Plans for the spiky traffic
- B. Use Spot Instances for the spiky traffic
- C. Use On-Demand Instances for the spiky traffic
- D. Use Dedicated Hosts for the spiky traffic

Cevap B



Domain: Design High-Performing Architectures

An IT Company is working on the disaster recovery strategy for one of their workloads deployed on AWS. They have begun with the recovery planning of the storage components. As of now, they have an EFS File System deployed in us-east-1 being used by hundreds of instances in this region. As a part of their recovery strategy, they also want the file system to be available in ap-south-1. For the same, they are looking for a way to bring a copy of this file system to ap-south-1 in the easiest way possible. Which of the following options can provide the BEST solution in this case?

- A. Use Amazon EFS Replication to accomplish the cross-region replication
- B. Bring the data in an S3 Bucket in us-east-1 and enable cross-region replication to copy the data to a bucket in ap-south-1
- C. Use AWS DataSync to accomplish the cross-region replication
- D. Use AWS Direct Connect to accomplish the cross-region replication

Cevap A

otto (fs-01f5416a63a3f3590)

Delete Attach

General Edit

Performance mode	Automatic backups
General Purpose	<input checked="" type="radio"/> Enabled
Throughput mode	Encrypted
Bursting	95c231cc-c4f7-4db7-9e1f-048dc850bb82 (aws/elasticfilesystem)
Lifecycle management	File system state

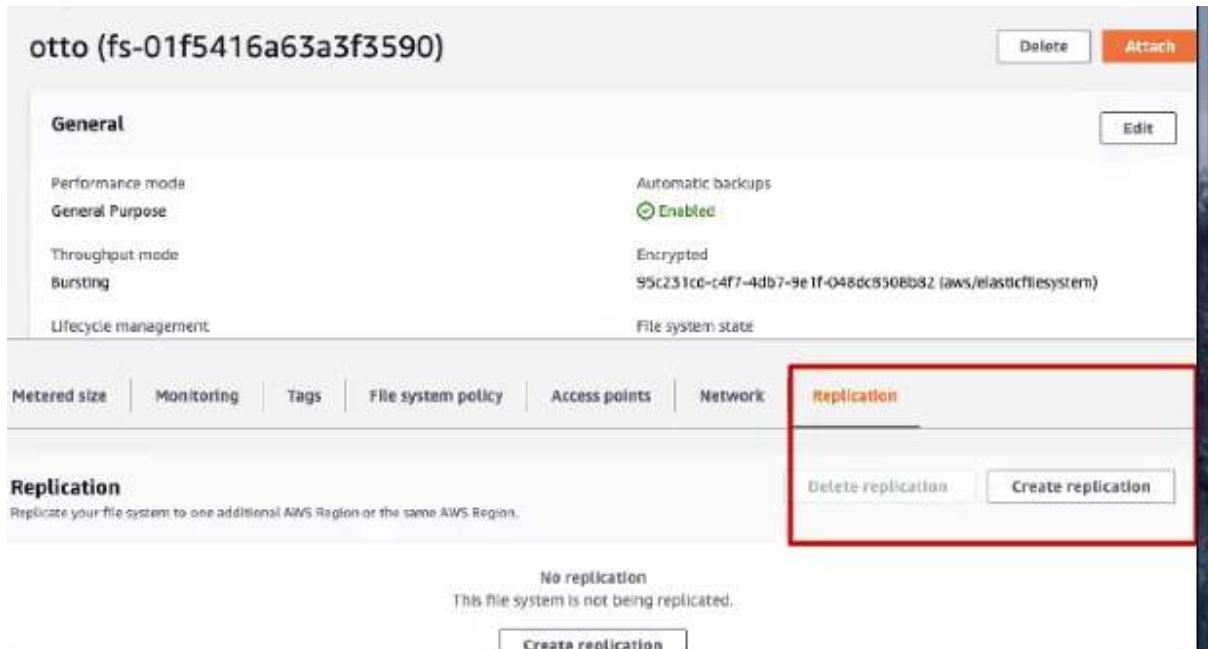
Metered size Monitoring Tags File system policy Access points Network Replication

Replication Delete replication Create replication

Replicate your file system to one additional AWS Region or the same AWS Region.

No replication
This file system is not being replicated.

Create replication



Domain: Design Secure Architectures

Your organization is looking for a standalone Active Directory solution on AWS with certain features and budget constraints. The first requirement is that the users in this directory should be able to use their directory credentials to log in to AWS Management Console and manage AWS resources. Also, there's a requirement to have daily automated snapshots. **The number of users will not exceed 450.** Which one of the following options can provide the required features at the lowest possible cost?

- A. AWS Directory Service for Microsoft AD
- B. Active Directory Connector
- C. Simple Active Directory
- D. Amazon Cognito

Cevap C

